

Lec 03: Secure Coding

CSED415: Computer Security
Spring 2024

Seulbae Kim



Grading (for newcomers)

- Midterm exam 25%
 - Final exam 30%
 - Lab assignments 25% (five labs, 5% each)
 - Research project 10%
 - + additional 5% as bonus for extraordinary teams
 - Extraordinary: work that's potentially publishable at great conferences
 - Participation 10%
- Up to **105%** including the bonus

Research project (for newcomers)

- We will have 4~5 teams
 - Subject to change considering the class size
- Select and work on **ANY** topic on computer security
 - e.g., VR/AR privacy (vision pro?) / breaking ChatGPT / crypto / ...

Summary of schedule (for newcomers)

- Week 2-3: Lab 1 // now
- Week 3: Team forming
- Week 4-5: Lab 2
- Week 6-7: Lab 3
- Week 8: Midterm exam
- Week 9: Project proposal
- Week 10-11: Lab 4
- Week 12-13: Lab 5
- Week 15: Project presentation
- Week 16: Final exam

Administrivia – Lab01 is out (due March 10th)

- Warm-up lab with guided self-learning (check PLMS for details)
- Familiarize yourselves with the setting and tools
- Dust off your x86 assembly knowledge from CSED211
 - If you never learned machine language, this is a great opportunity to learn
- Learn basic static/dynamic analysis techniques
- Use the Q&A board on PLMS for questions
- Start early!
- Start early!
- Start early!

Recap

- Week 1 – basic principles and concepts
 - Understanding computer security and challenges
 - CIA+AA: pillars of a secure system
 - Analyzing security through threat modeling
 - 13 fundamental principles for secure design
 - You should be able to articulate each principle and give examples
 - e.g., “Tell me about the ‘open design’ principle”

This week: practical aspects

- Secure coding
 - How to implement (in)secure systems?
 - or, how to write (in)secure code?
- Trusting trust
 - Is secure coding enough for computer security?

Defensive Programming

Bug

- Some bugs can be exploited by an attacker to compromise the entire computer system
 - This class of bugs is called “vulnerability” (ref: Lec 02)



If we write a bug-free code, then
the program has no vulnerability!

Defensive programming

- Process of designing and implementing a bug-free software so it continues to function even when under attack
- Such software is able to
 - Detect erroneous conditions resulting from attacks
 - Continue executing safely or fail gracefully
- Secure coding is a type of defensive programming, primarily focused on enhancing computer security

Insecure vs. secure

```
1 int func(char* input) {  
2     char buf[8];  
3     strcpy(buf, input);  
4     /* ... */  
5 }
```

Overflow! Insecure :(

```
1 int func(char* input) {  
2     char buf[8];  
3     strncpy(buf, input, 8);  
4     /* ... */  
5 }
```

Copies up to 8 bytes. Secure! :)

Really??

Memory view (64-bit architecture)

- Memory is a key-value storage
 - Key: address, Value: memory contents
 - One cell takes 64 bits (== 8 bytes)

<Address>	<Value>
0x7fffffffffe2e0	0x99 0x7c 0xcf 0xfd 0x74 0x4b 0xb6 0x5b
0x7fffffffffe2e8	0x47 0x12 0xe4 0x64 0x18 0x07 0x54 0xb3
0x7fffffffffe2f0	0x92 0x2e 0x7b 0x35 0x6c 0x0a 0xa4 0xf8
0x7fffffffffe2f8	0x72 0x76 0xc7 0x33 0x01 0x5d 0x75
0x7fffffffffe300	0x50 0x08 0x18 0xc7 0xe0 0x1e
...	...

garbage data

Memory view (64-bit architecture)

- `char buf[8]` takes one cell
 - Q) address of `buf[2]`?

<Address>	<Value>	
0x7fffffffffe2e0	0x99 0x7c 0xcf 0xfd 0x74 0x4b 0xb6 0x5b	} buf
0x7fffffffffe2e8	0x47 0x12 0xe4 0x64 0x18 0x97 0x54 0xb3	
0x7fffffffffe2f0	0x92 0x2e 0x7b 0x35 0x6c 0xb2 0xa4 0xf8	
0x7fffffffffe2f8	0x72 0x76 0xc7 0x33 0x54 0x41 0x5d 0x75	
0x7fffffffffe300	0x50 0x08 0x18 0x69 0xcc 0xcf 0xe0 0x1e	
...	...	

Memory view (64-bit architecture)

- `strcpy(buf, input);`
 - `input = "aaaabbbbccccddddeeeeffffggggghhhh\0"`
 - Q) Why is this potentially dangerous?

<Address>	<Value>	
0x7fffffffffe2e0	0x61 0x61 0x61 0x61 0x62 0x62 0x62 0x62	} buf
0x7fffffffffe2e8	0x63 0x63 0x63 0x63 0x64 0x64 0x64 0x64	
0x7fffffffffe2f0	0x65 0x65 0x65 0x65 0x66 0x66 0x66 0x66	} overflown!
0x7fffffffffe2f8	0x67 0x67 0x67 0x67 0x68 0x68 0x68 0x68	
0x7fffffffffe300	0x00 0x08 0x18 0x69 0xcc 0xcf 0xe0 0x1e	
...	...	

Memory view (64-bit architecture)

- `strcpy(buf, input, 8);`
 - `input = "aaaabbbbccccddddeeeeffffggggghhhh\0"`
 - Q) Why is this still considered unsafe?

<Address>	<Value>	
0x7fffffffffe2e0	0x61 0x61 0x61 0x61 0x62 0x62 0x62 0x62	} buf
0x7fffffffffe2e8	0x47 0x12 0xe4 0x64 0x18 0x97 0x54 0xb3	
0x7fffffffffe2f0	0x92 0x2e 0x7b 0x35 0x6c 0xb2 0xa4 0xf8	} No longer overflown!
0x7fffffffffe2f8	0x72 0x76 0xc7 0x33 0x54 0x41 0x5d 0x75	
0x7fffffffffe300	0x50 0x08 0x18 0x69 0xcc 0xcf 0xe0 0x1e	
...	...	

C programming 101

- Strings in C are char arrays
 - `char cnum[8] = "CSED415";` // need a space for a **NULL** terminator

C	S	E	D	4	1	5	\0
---	---	---	---	---	---	---	----

- C has no semantic notion of strings
 - A string at address **p**
== a sequence of characters from **p** to the first **NULL** terminator
 - Without '**\0**', C does not know where the string ends

Memory view (64-bit architecture)

- `strcpy(buf, input, 8);`
 - e.g., `puts(buf);` // leads to an undefined behavior (UB)

<Address>	<Value>	
0x7fffffffffe2e0	0x61 0x61 0x61 0x61 0x62 0x62 0x62 0x62	buf
0x7fffffffffe2e8	0x47 0x12 0xe4 0x64 0x18 0x97 0x54 0xb3	
0x7fffffffffe2f0	0x92 0x2e 0x7b 0x35 0x6c 0xb2 0xa4 0xf8	
0x7fffffffffe2f8	0x72 0x76 0xc7 0x33 0x54 0x41 0x5d 0x75	
0x7fffffffffe300	0x50 0x08 0x18 0x69 0xcc 0xcf 0xe0 0x1e	
...	...	Memory contents up to a NULL byte are printed

Test it yourself

```
1 #include <stdio.h>
2 #include <string.h>
3
4 void func(char* input) {
5     char buf[8];
6     strncpy(buf, input, 8);
7     puts(buf); // what does this print?
8 }
9
10 int main(void) {
11     func("aaaabbbbccccddddeeeeffffgggghhhh");
12     return 0;
13 }
```

One more step after strncpy

- Null-terminate the string

```
1 int func(char* input) {  
2     char buf[8];  
3     strcpy(buf, input);  
4     /* ... */  
5 }
```

Bad

```
1 int func(char* input) {  
2     char buf[8];  
3     strncpy(buf, input, 8);  
4     /* ... */  
5 }
```

Still bad

```
1 int func(char* input) {  
2     char buf[8];  
3     strncpy(buf, input, 8);  
4     buf[7] = 0; // Beware: not buf[8]  
5     /* ... */  
6 }
```

Safe now

Secure Coding Guidelines

SEI CERT C coding standard

- <https://wiki.sei.cmu.edu/confluence/display/c>
 - Documented by CMU
 - Conformance to this standard are necessary (but not sufficient) to ensure the safety, reliability, and security of software systems developed in C
 - A lot of rules – do not need to memorize all, but recommended to read through them at least once when you have time

#1: Declare objects with appropriate storage duration

- The lifetime of an object is the portion of program execution during which storage is guaranteed to be reserved for it. An object exists, has a constant address, and retains its last-stored value throughout its lifetime.
- If an object is referred to outside of its lifetime, **the behavior is undefined**. The value of a pointer becomes indeterminate when the object it points to reaches the end of its lifetime

Example #1-1-noncompliant

```
1 const char *p; // Static variable
2 void dont_do_this(void) {
3     const char c_str[] = "This will change"; // Automatic variable
4     p = c_str; /* Dangerous */ // lifetime mismatch
5 }
6
7 void innocuous(void) {
8     printf("%s\n", p); // out of scope → takes on an indeterminate value
9 }
```

Example #1-1-compliant

```
1 void this_is_OK(void) {  
2     const char c_str[] = "Everything OK";  
3     const char *p = c_str; // same storage duration  
4     /* ... */  
5 }  
6 /* p is inaccessible outside the scope of string c_str */
```

```
1 const char *p;  
2 void is_this_OK(void) {  
3     const char c_str[] = "Everything OK?";  
4     p = c_str;  
5     /* ... */  
6     p = NULL; // p is cleared before c_str is destroyed  
7 }
```


Example #1-2-noncompliant

```
1 char *init_array(void) {  
2     char array[10]; // Automatic storage duration  
3     /* Initialize array */  
4     return array;    // The caller can access the destroyed array  
5 }
```

Example #1-2-compliant

```
1 #include <stddef.h>
2 void init_array(char *array, size_t len) {
3     /* Initialize array */
4     return;
5 }
6
7 int main(void) {
8     char array[10]; // Keep object under same scope
9     init_array(array, sizeof(array) / sizeof(array[0]));
10    /* ... */
11    return 0;
12 }
```

#2: Do not call system()

- `system("cmd")` executes a command by invoking a shell
- `system("cmd")` can be exploited if
 - `cmd` is not sanitized or improperly sanitized
 - `cmd` is specified without a path name
 - relative path used in `cmd` can be modified by an attacker
 - executable program specified by `cmd` can be spoofed by an attacker

Example #2-1-noncompliant

```
1 void run_ls(const char *input) {  
2     char cmd[512];  
3     snprintf(cmd, 512, "ls %s", input);  
4     system(cmd);  
5 }
```

If input is `"/home/lab01"` then `system("ls /home/lab01");` is executed

If input is `"/home/lab01; useradd attacker"`
then two commands are invoked in sequence:
(1) `ls /home/lab01` and
(2) `useradd attacker`

#3: Do not depend on the order of evaluation for side effects

- Evaluation of an expression may produce side effects
- Sequence points are points that distinguish evaluations during execution
 - e.g., between the evaluations of the operands of &&
- Do not depend on the order of evaluation for side effects unless there is an intervening sequence point

Example #3-noncompliant and compliant

```
1 extern void c(int i, int j);
2 int glob;
3
4 int a(void) {
5     return glob + 10;
6 }
7
8 int b(void) {
9     glob = 42;
10    return glob;
11 }
12
13 void func(void) {
14     c(a(), b());
15 }    // order of argument
        evaluation is unspecified
```

```
1 extern void c(int i, int j);
2 int glob;
3
4 int a(void) {
5     return glob + 10;
6 }
7 int b(void) {
8     glob = 42;
9     return glob;
10 }
11
12 void func(void) {
13     int a_val = a();
14     int b_val = b();
15     c(a_val, b_val);
16 }
```

#4: Do not read uninitialized memory

- If an object that has automatic storage duration is not initialized explicitly, its value is indeterminate
 - Local, automatic variables are stored on the stack
 - Their initial values default to the current values of the stack
- Dynamic allocators have different behaviors
 - `calloc()`: Zero-initializes allocated memory
 - `malloc()`: Does not initialize allocated memory
 - `realloc()`: Copies contents from original ptr; may not init. all memory

Example #4-1-noncompliant

```
1 void set_flag(int number, int *sign_flag) {
2     if (NULL == sign_flag) return;
3
4     if (number > 0) {
5         *sign_flag = 1;
6     } else if (number < 0) {
7         *sign_flag = -1;
8     } // missing case: number == 0
9 }
10
11 int is_negative(int number) {
12     int sign;
13     set_flag(number, &sign); // sign is uninitialized
14     return sign < 0; // comparison exhibits undefined behavior
15 }
```


Example #4-2-noncompliant

- CVE-2008-0166
 - Debian devs got creative and decided to use uninitialized memory for random number generation
 - Why is this a bad idea?

```
1 void func(void) {  
2     struct timeval tv;  
3     unsigned long junk;  
4  
5     gettimeofday(&tv, NULL);  
6     srandom((getpid() << 16) ^ tv.tv_sec ^ tv.tv_usec ^ junk);  
7 }
```

// utilize random stack
contents as seed

Example #4-2-noncompliant

- CVE-2008-0166
 - Different compilers employ different optimization strategies
 - LLVM optimizes out uninitialized variables completely

```
1 void func(void) {  
2     struct timeval tv;  
3     unsigned long junk;  
4  
5     gettimeofday(&tv, NULL);  
6     srandom((getpid() << 16) ^ tv.tv_sec ^ tv.tv_usec ^ junk);  
7 }  
// removed by compiler
```


Rule #5: Do not dereference NULL pointers

- A NULL pointer represents a pointer that does not refer to a valid object or memory location

```
1 int *p = NULL;
2 /* ... */
3 *p = 415; // what happens?
```

<Address>
0x7fffffffffe2e0
...

<Value>			
0x99	0x7c	0xcf	0xfd
0x47	0x12	0xe4	0x64
0x92	0x2e	0x7b	0x35
0x72	0x76	0xc7	0x33
0x50	0x08	0x18	0x69
...			



Cannot write to a non-existent
memory location (Segmentation fault)

Example #5-noncompliant

- Q) Find the problem

```
1 #include <png.h> /* From libpng */
2 #include <string.h>
3                                     // if length == -1
4 void func(png_structp png_ptr, int length, const void *user_data) {
5     png_charp chunkdata;
6     chunkdata = (png_charp)png_malloc(png_ptr, length + 1);
7     /* ... */                                     // malloc of size 0 returns NULL
8     memcpy(chunkdata, user_data, length);
9     /* ... */ // chunkdata is NULL, and gets dereferenced
10 }
```

Rule #6: Ensure (un)signed int operations do not wrap

- signed and unsigned integers have valid ranges
 - int32_t: [-2147483648 to 2147483647]
 - uint32_t: [0 to 4294967295]

```
1 int x = 2147483647;  
2 int y = 1;  
3 printf("%d\n", x + y); // result?
```

// prints -2147483648 (silent wraparound)

Example #6-noncompliant

```
1 void func(unsigned int a, unsigned int b) {  
2     unsigned int sum;  
3     sum = a + b;  
4     char* buf = malloc(sum);  
5     /* ... */  
6 }
```

Example #6-compliant

```
1 void func(unsigned int a, unsigned int b) {
2     unsigned int sum;
3     if (UINT_MAX - a < b) { // Beware: if (a + b < UINT_MAX) doesn't work
4         /* Handle error */
5     } else {
6         sum = a + b;
7     }
8     char* buf = malloc(sum);
9     /* ... */
10 }
```

Rule #7: Ensure that integer conversions do not result in lost or misinterpreted data

- Implicit and explicit int conversions must be guaranteed not to result in lost or misinterpreted data

```
1 #include <stdio.h>
2
3 int main(void) {
4     int i = 0x100; // Uses two bytes
5     printf("%d\n", i);
6     char c = (char)i; /* explicit casting */
7     printf("%d\n", c); // char type can hold one byte
8     return 0;
9 }
```


Rule #8: Ensure that division and remainder operations do not result in divide-by-zero errors

- When the value of the second operand of the / or % operator is zero, i.e., divide-by-zero, the behavior is undefined (UB)

```
1 int main(void) {  
2     int x = 1 / 0; // Floating point exception  
3     return 0;  
4 }
```

Complicated cases exist in practice

Example #8-noncompliant

- CVE-2018-13097 in Linux kernel (F2FS file system)

```
1 // linux/fs/f2fs/f2fs.h
2 struct f2fs_sb_info {
3     struct super_block *sb; /* pointer to VFS super block */
4     struct proc_dir_entry *s_proc; /* proc entry */
5     /* ... */
6     block_t user_block_count; /* # of user blocks */
7     block_t total_valid_block_count; /* # of valid blocks */
8     /* ... */
9 }
10
11 // linux/fs/f2fs/segment.h
12 static inline int utilization(struct f2fs_sb_info *sbi)
13 {
14     return div_u64((u64)valid_user_blocks(sbi) * 100,
15                   sbi->user_block_count);
16 }
```

File system image metadata
(superblock)

Corrupted due to crash
(becomes zero)

Div-by-zero when mounting the corrupt image

Rule #9: Do not use floating-point variables as loop counters

- Computers cannot accurately represent all real numbers
- The precision differs depending on the CPU

```
1 #!/usr/bin/python3
2 print(0.1 + 0.2 == 0.3) # result?
```

```
>>> 0.1 + 0.2 == 0.3
False
>>> 0.1 + 0.2
0.30000000000000004
```

Example #9-1-noncompliant

```
1 void func(void) {  
2     for (float x = 0.1f; x <= 1.0f; x += 0.1f) {  
3         /* Loop may iterate 9 or 10 times */  
4     }  
5 }
```

Logically, the loop should be evaluated 10 times
(0.1, 0.2, 0.3, ..., 1.0)

The loop only iterates 9 times on an x86_64 processor
(our lab server)

Example #9-2-noncompliant

```
1 void func(void) {  
2     for (float x = 1000000001.0f; x <= 1000000010.0f; x += 1.0f) {  
3         /* Loop may not terminate */  
4     }  
5 }
```

$1000000001.0f + 1.0f = 1.000000001.0f$

(incremented by an amount that is too small given the precision)

Rule #10: Do not form or use out-of-bounds pointers or array subscripts

```
1 enum { TABLESIZE = 100 };
2
3 static int table[TABLESIZE];
4
5 int *get_ptr(int index) {
6     if (index < TABLESIZE) { // check if index is beyond array bounds
7         return table + index; // is the check sufficient?
8     }
9     return NULL;
10 }
```

Example #10-compliant

```
1 enum { TABLESIZE = 100 };
2
3 static int table[TABLESIZE];
4
5 int *f(int index) {
6     if (index >= 0 && index < TABLESIZE) {
7         return table + index;
8     }
9     return NULL;
10 }
```

do explicit check

```
1 #include <stddef.h>
2 enum { TABLESIZE = 100 };
3
4 static int table[TABLESIZE];
5
6 int *f(size_t index) {
7     if (index < TABLESIZE) {
8         return table + index;
9     }
10    return NULL;
11 }
```

use unsigned type


Rule #11: Do not access freed memory

- Freed memory == dangling pointer
 - `free(ptr);`
- Evaluating a dangling pointer is undefined behavior
 - Dereferencing it // `*ptr`
 - Using it as an operand // `*ptr + 20`
 - Type casting it // `(char*)ptr`
 - Using it as the right-hand side of an assignment // `*p = *ptr`

Example #11-noncompliant

- A famous example from Data Structure 101

```
1 #include <stdlib.h>
2
3 struct node {
4     int value;
5     struct node *next;
6 };
7
8 void free_list(struct node *head) {
9     for (struct node *p = head; p != NULL; p = p->next) {
10         free(p); // p is freed before p->next is executed.
11     }           // p->next reads freed memory
12 }
```



Example #11-compliant

```
1 #include <stdlib.h>
2
3 struct node {
4     int value;
5     struct node *next;
6 };
7
8 void free_list(struct node *head) {
9     struct node *q;
10    for (struct node *p = head; p != NULL; p = q) {
11        q = p->next; // Keep track of p->next before freeing p
12        free(p);
13    }
14 }
```

Rule #12: Detect and remove code that has no effect or is never executed

- Dead code and unreachable code can cause unexpected behavior
 - Usually optimized out during compilation by modern compilers, but not always!
- This rule is related to Lab01

Example #12-1-noncompliant

```
1 int func(int condition) {
2     char *s = NULL;
3     if (condition) {
4         s = (char *)malloc(10);
5         if (s == NULL) {
6             /* Handle Error */
7         }
8         /* Process s */
9         return 0; // function always returns here
10    }
11
12    if (s) {
13        /* This code is unreachable */
14    }
15    return 0;
16 }
```

Example #12-1-compliant

```
1 int func(int condition) {
2     char *s = NULL;
3     if (condition) {
4         s = (char *)malloc(10);
5         if (s == NULL) {
6             /* Handle Error */
7         }
8         /* Process s */
9         // return 0; // simple fix
10    }
11
12    if (s) {
13        /* now reachable */ // do cleanup, e.g., free(s);
14    }
15    return 0;
16 }
```

Example #12-2-noncompliant

```
1 int s_loop(char *s) {
2     size_t i;
3     size_t len = strlen(s); // strlen returns the number of chars preceding '\0'
4     for (i = 0; i < len; ++i) {
5         /* Code that doesn't change s, i, or len */
6         if (s[i] == '\0') { // this condition can never be satisfied
7             /* This code is never reached */
8         }
9     }
10    return 0;
11 }
```

Leads to Denial of Service (DoS)

And many more...

- Do check out the guideline!
 - <https://wiki.sei.cmu.edu/confluence/display/c>

Open question

- Why don't we write an analyzer that checks for all rules in the secure coding standard?

Coming up next: Trusting trust

- Is code-level analysis enough to build secure systems?

Questions?