

# Lec 05: x86 Assembly

CSED415: Computer Security  
Spring 2024

Seulbae Kim



# Administrivia

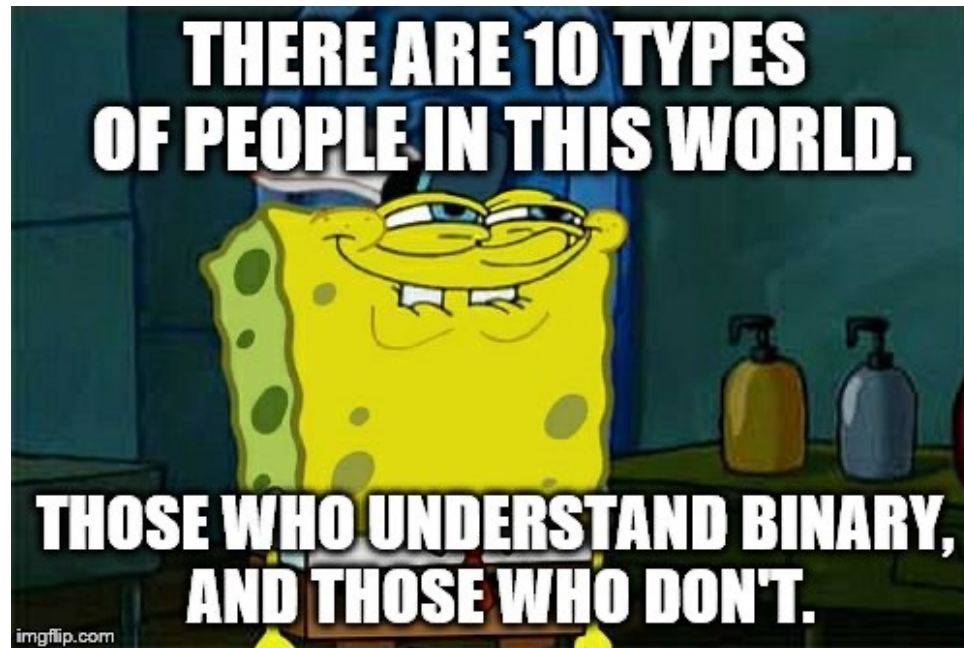
---

- Lab 01 is due this Sunday
  - Focus on today's lecture if you are struggling!
- Team forming is also due this Sunday
  - 2 students per team, six teams
  - Check “Find your Teammates” board on PLMS
- Jeongjin (TA) is taking a leave of absence
  - Please do not send him emails
  - Use PLMS Q&A board for questions

# Recap

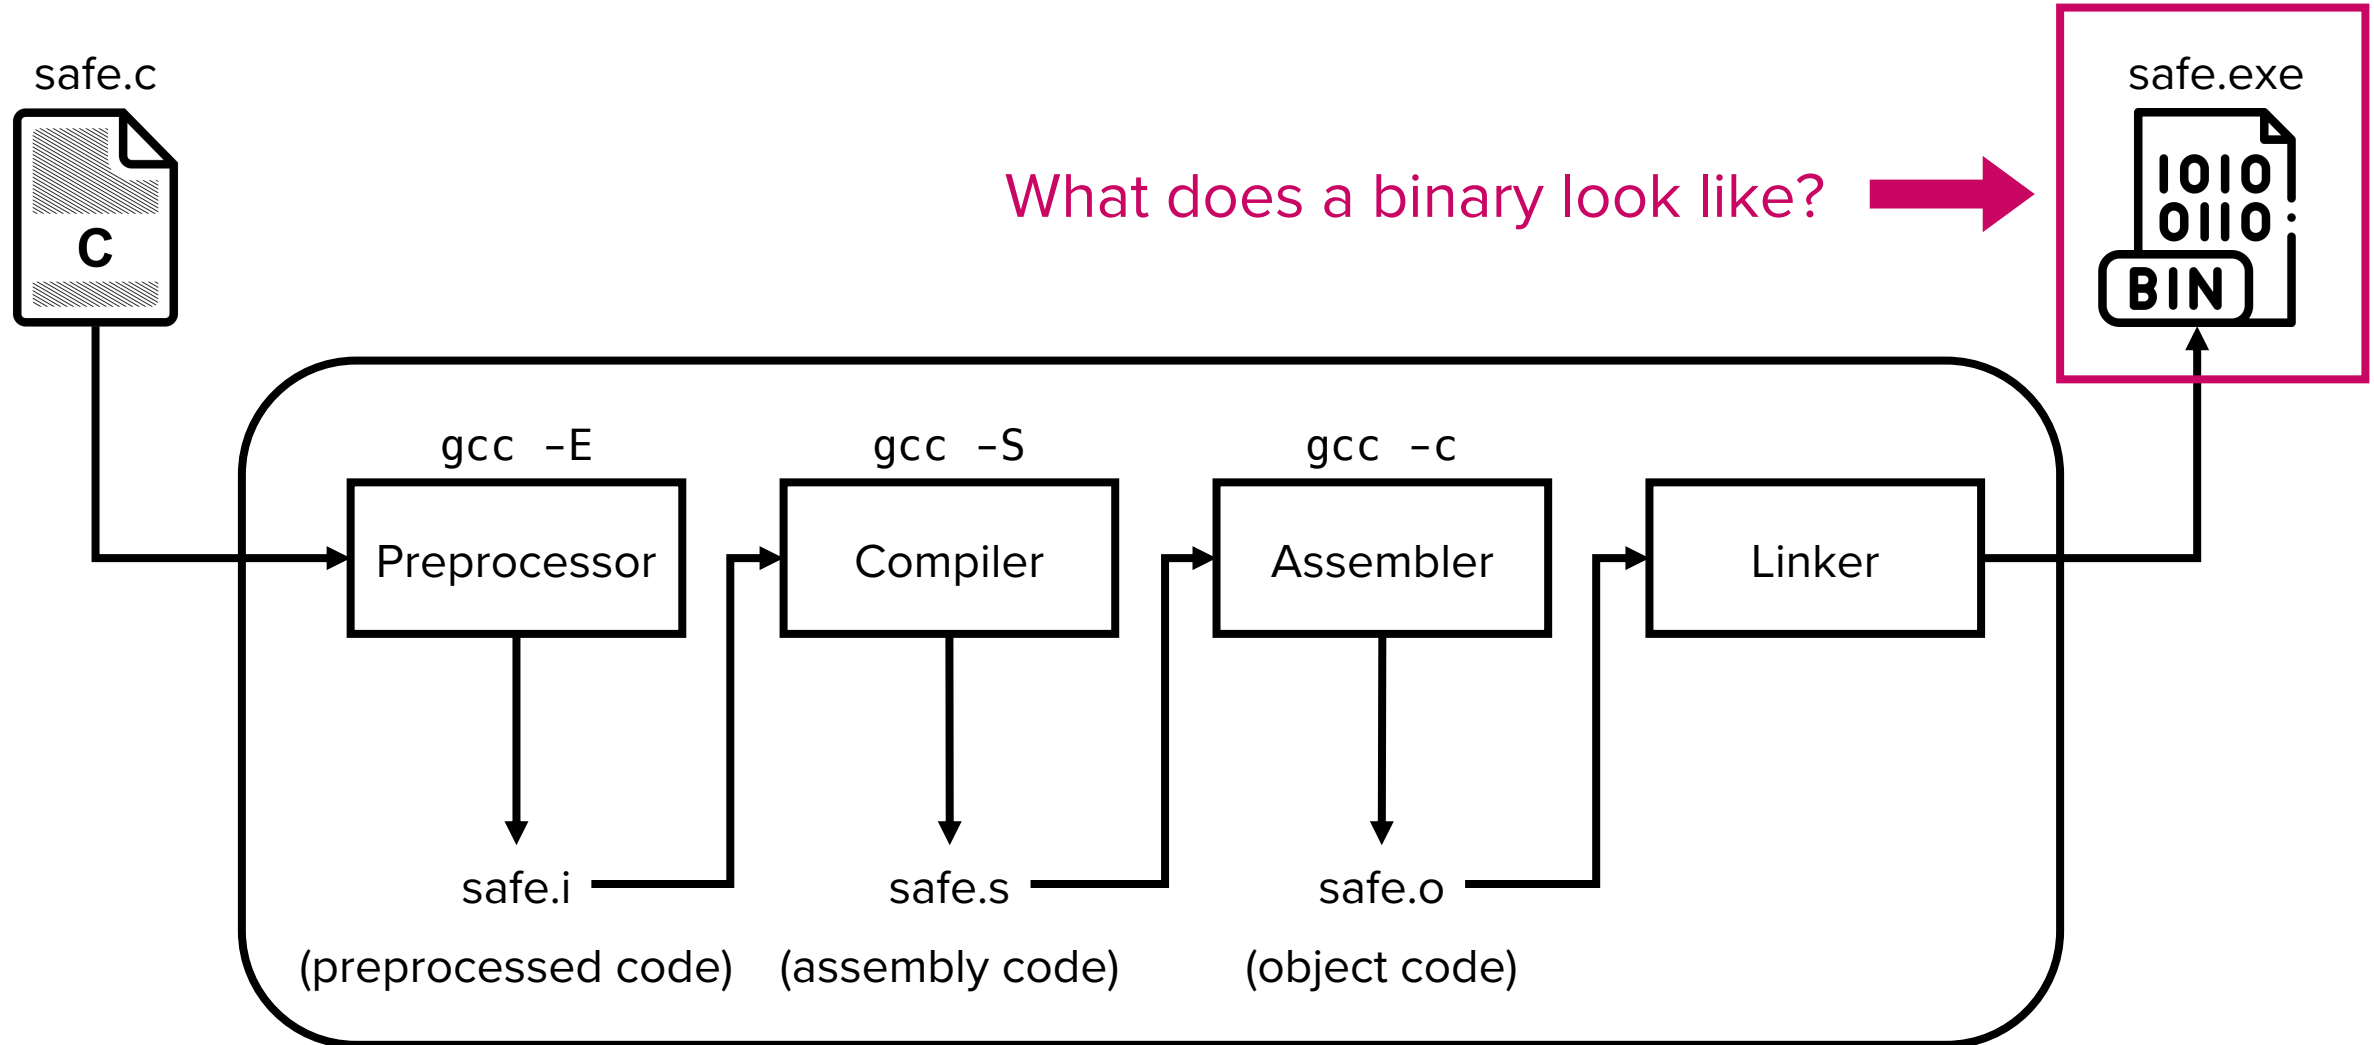
- Secure coding is necessary to build secure computer systems
- However, code-level mitigation may not be enough
  - We cannot trust anything that we did not create ourselves
  - e.g., Thompson compiler
- Binary analysis, like it or not, is required
  - We want to find out whether a system is secure or not
  - Analyze exactly what gets executed on the CPU

BIN

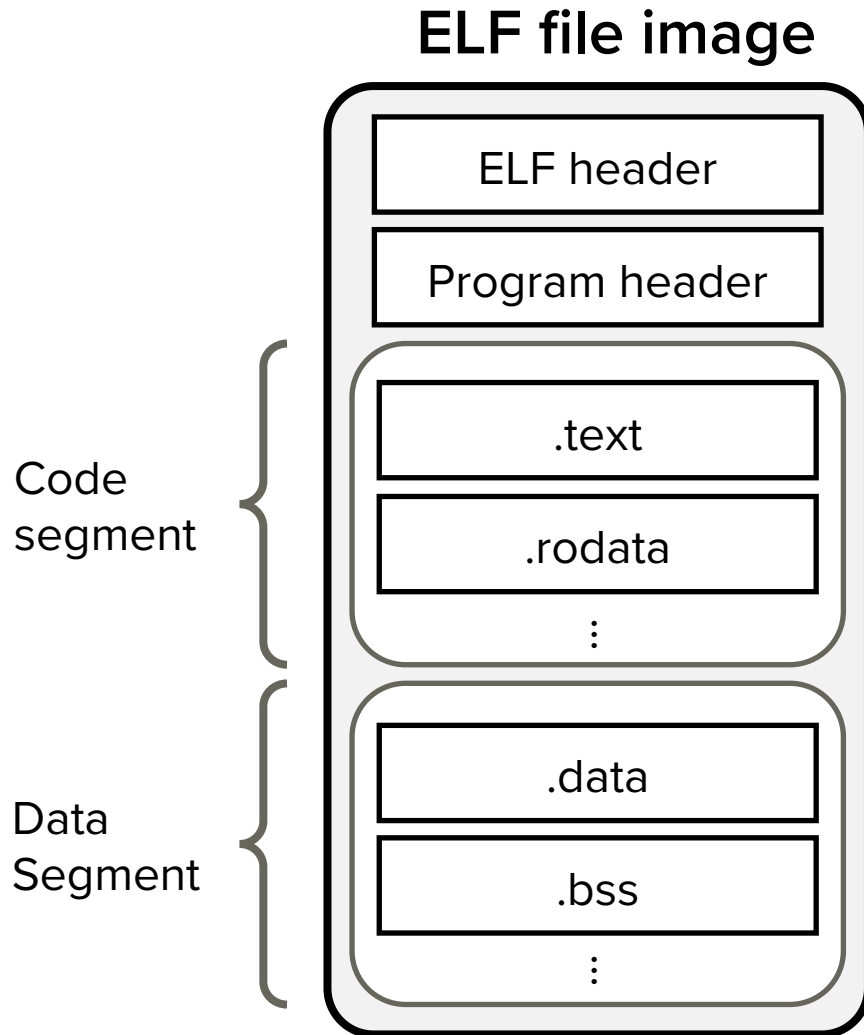


ARY

# Recap: Compiler 101

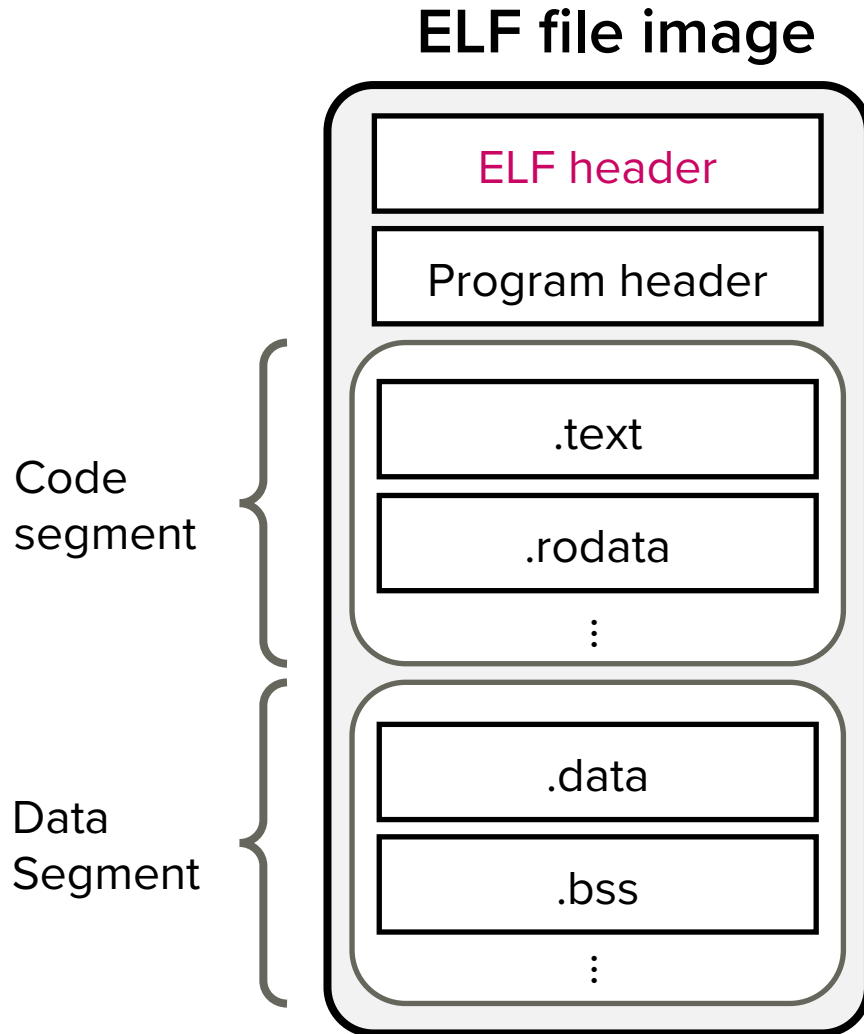


# Executable binary (ELF-formatted file)



- ELF (Executable and Linkable Format)
  - ELF header
  - Program header
  - Segments
    - Contains Sections

# Executable binary

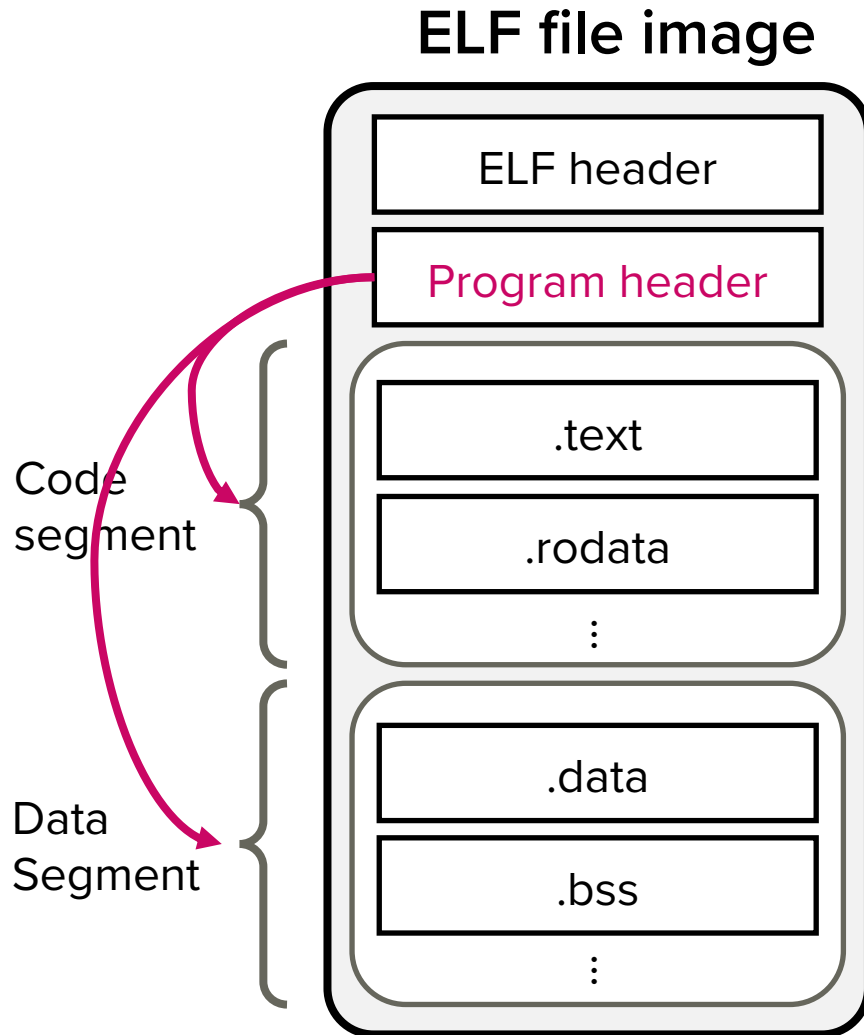


- ELF header

- Always at the beginning of an ELF file
- Contains magic number, architecture info, entry point address, ...

```
lab01@csed415:~$ readelf -h ./target
ELF Header:
  Magic:   7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:                                ELF32
  Data:                                      2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  EXEC (Executable file)
  Machine:                               Intel 80386
  Version:                               0x1
  Entry point address:                   0x8049110
  Start of program headers:              52 (bytes into file)
  Start of section headers:             14328 (bytes into file)
  Flags:                                  0x0
  Size of this header:                   52 (bytes)
  Size of program headers:               32 (bytes)
  Number of program headers:              11
  Size of section headers:               40 (bytes)
  Number of section headers:              29
  Section header string table index:     28
```

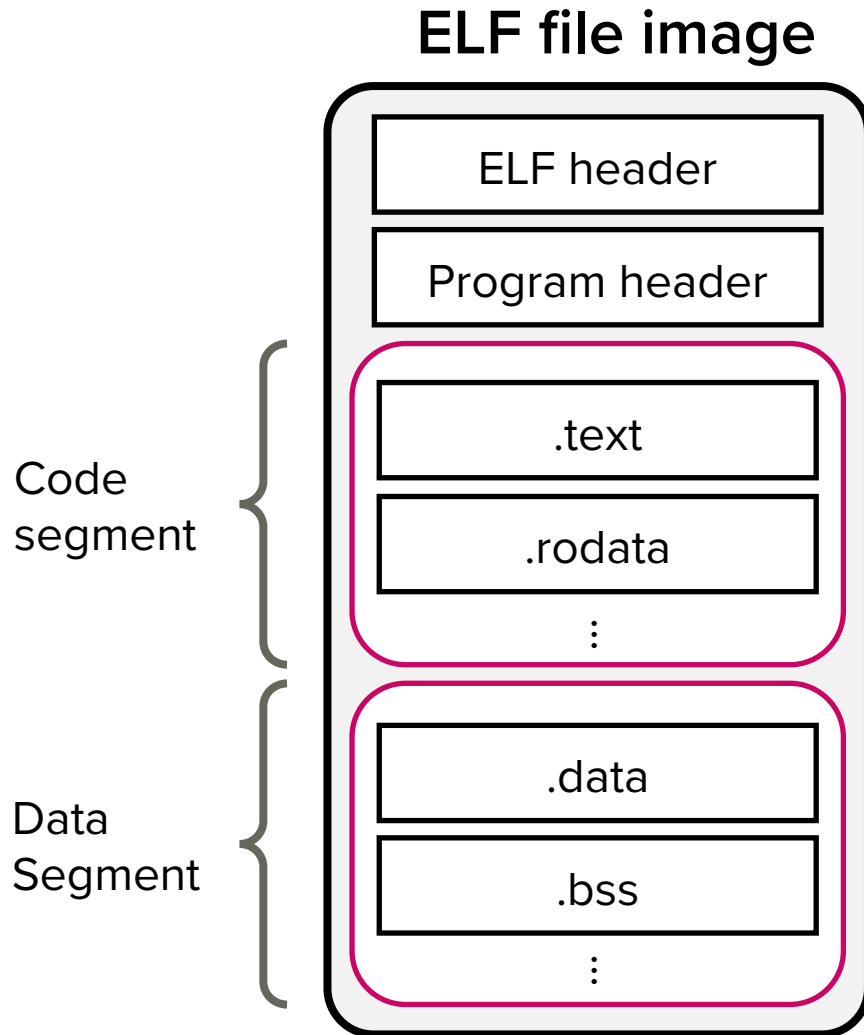
# Executable binary



- Program header
  - Contains segment information
    - Offset and size of each segment (code segment is size bytes from offset)
    - Virtual memory address each segment should be loaded
    - Permission of each segment



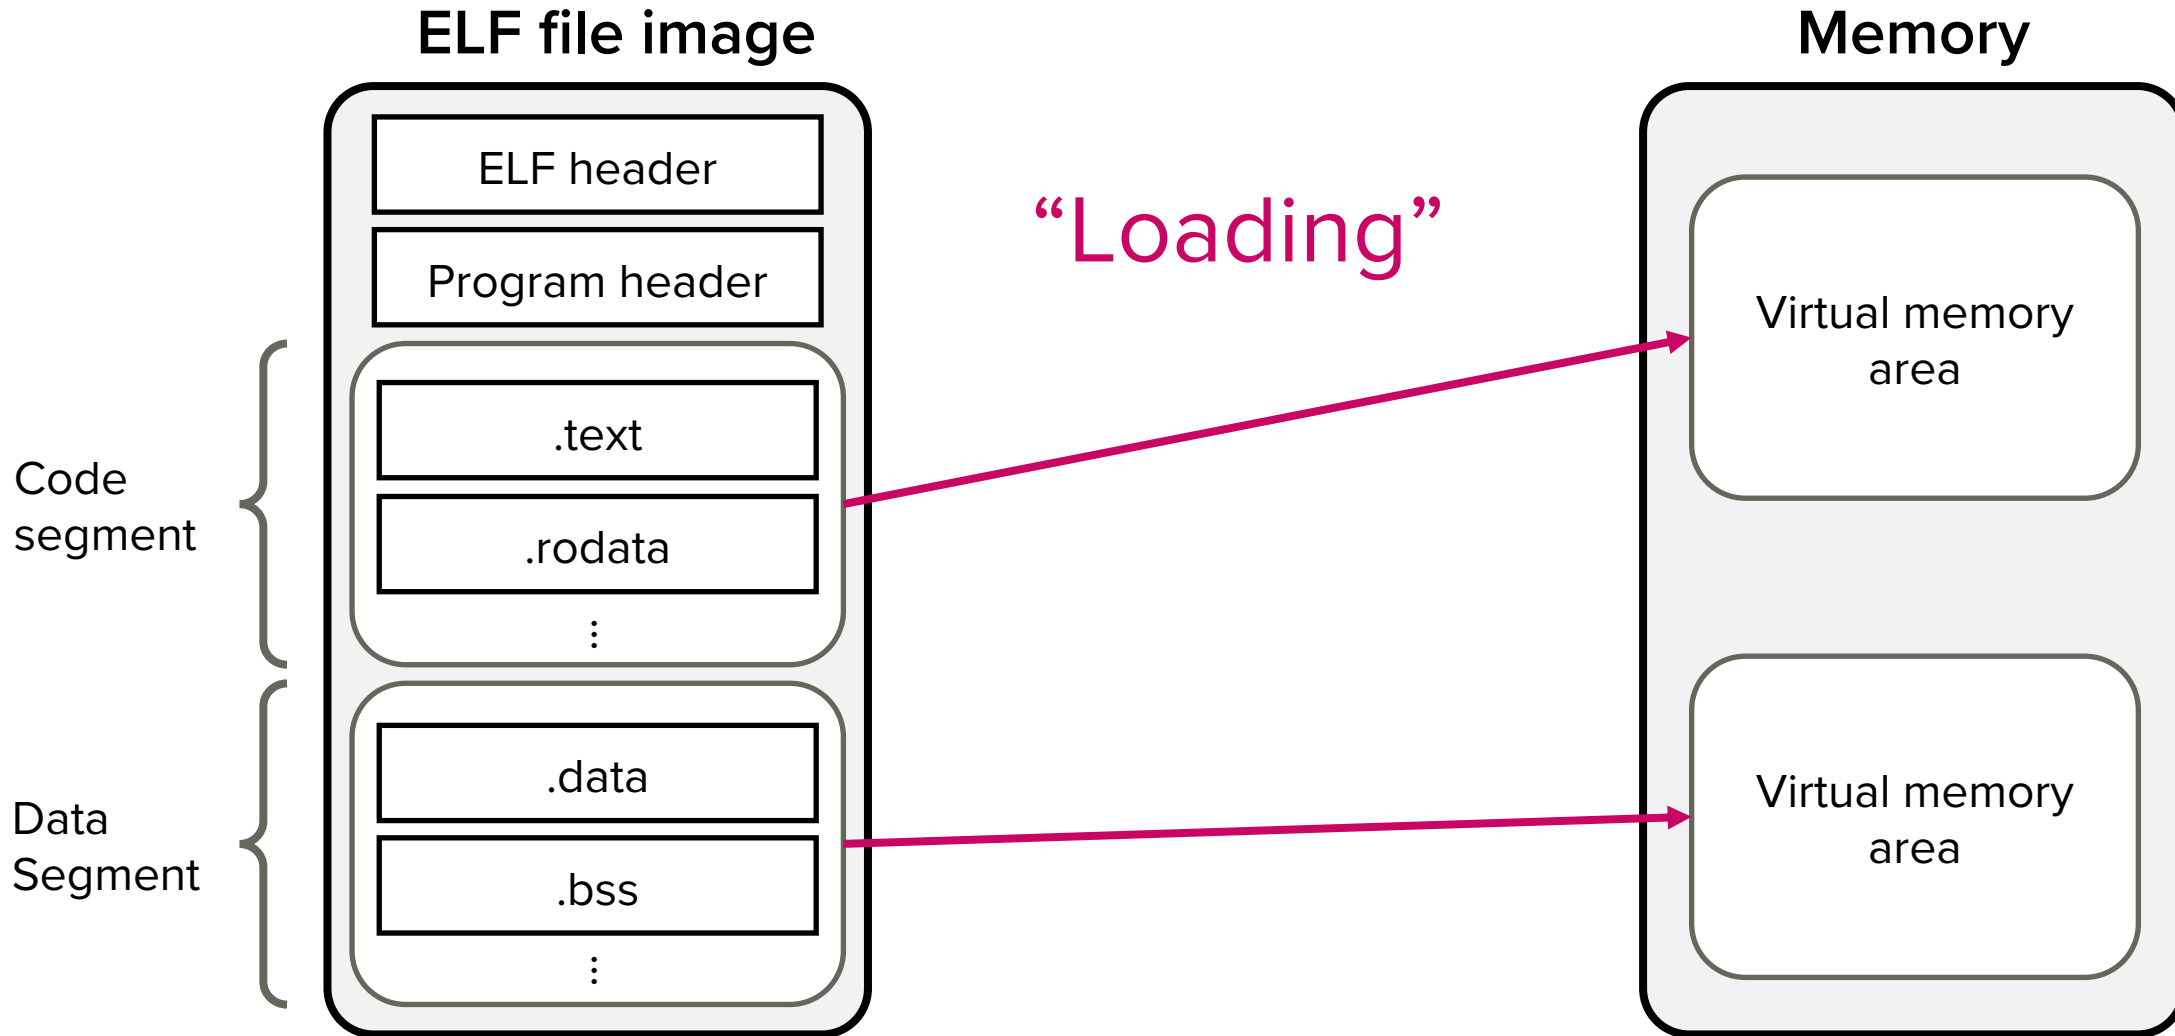
# Executable binary



- Segments

- Logical groups of sections
- Each segment is loaded onto one or more virtual memory areas at runtime
- Code segment (text segment)
  - Includes program instructions and read only data (e.g., string literals)
- Data segment
  - Includes static variables, uninitialized variables, ...

# Loader turns a binary image into a process



# ELF loader workflow

---

1. Read the ELF Header, which is always located at the very beginning of an ELF file
2. Read the program headers. These specify where in the file the program segments are located, and where they need to be loaded into memory
3. Parse the program headers to determine the number of program segments that must be loaded

# ELF loader workflow

## 4. Load each of the loadable segments

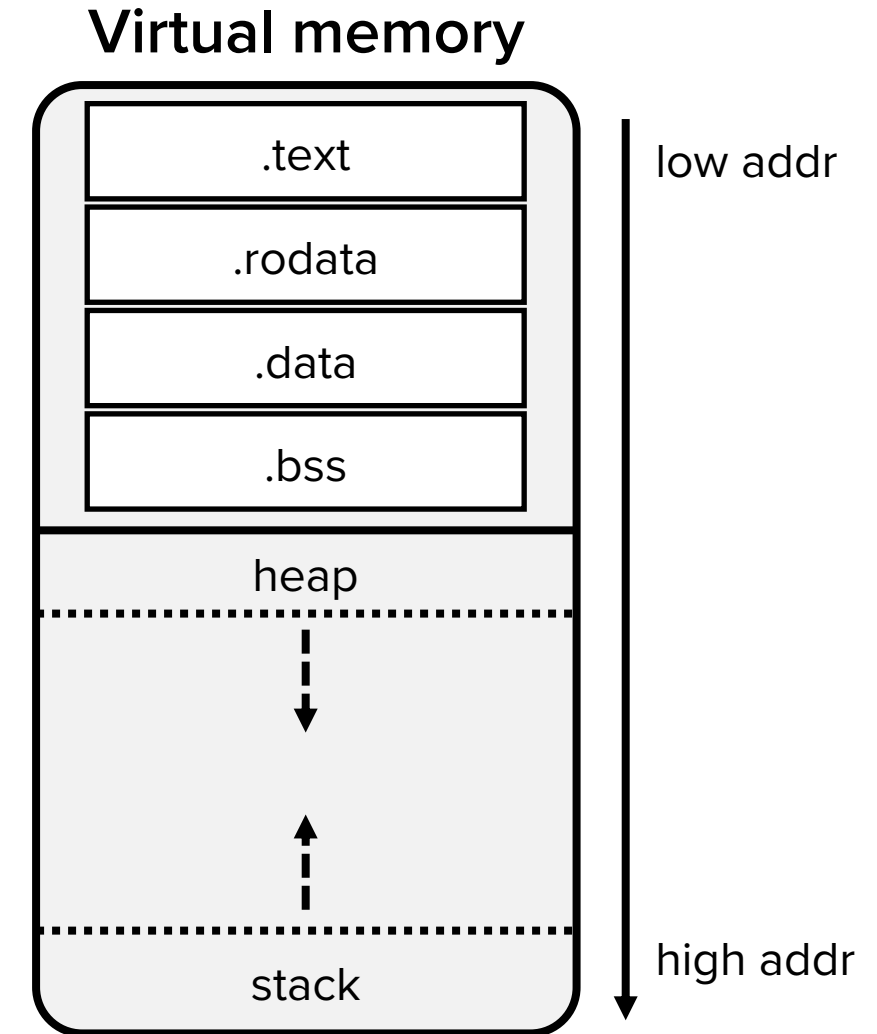
- a. Allocate virtual memory for each segment, at the address specified by the `p_vaddr` member in the program header
- b. Copy the segment data from the ELF file offset specified by `p_offset` to the virtual memory address specified by `p_vaddr`

## 5. Read the executable's entry point from the ELF header

## 6. Jump to the executable's entry point in the newly loaded memory

# A process in memory

- ELF segments + heap + stack
  - Heap
    - Dynamically allocated memory (e.g., using malloc)
    - Grows towards higher (larger) addresses
  - Stack
    - Contains runtime call stack
      - More on this later!
    - Grows towards lower (smaller) addresses



# Exercise: Examine Lab 01's target in GDB

```
$ ssh lab01@141.223.181.22
```

```
$ gdb ./target
```

```
pwndbg> break main
```

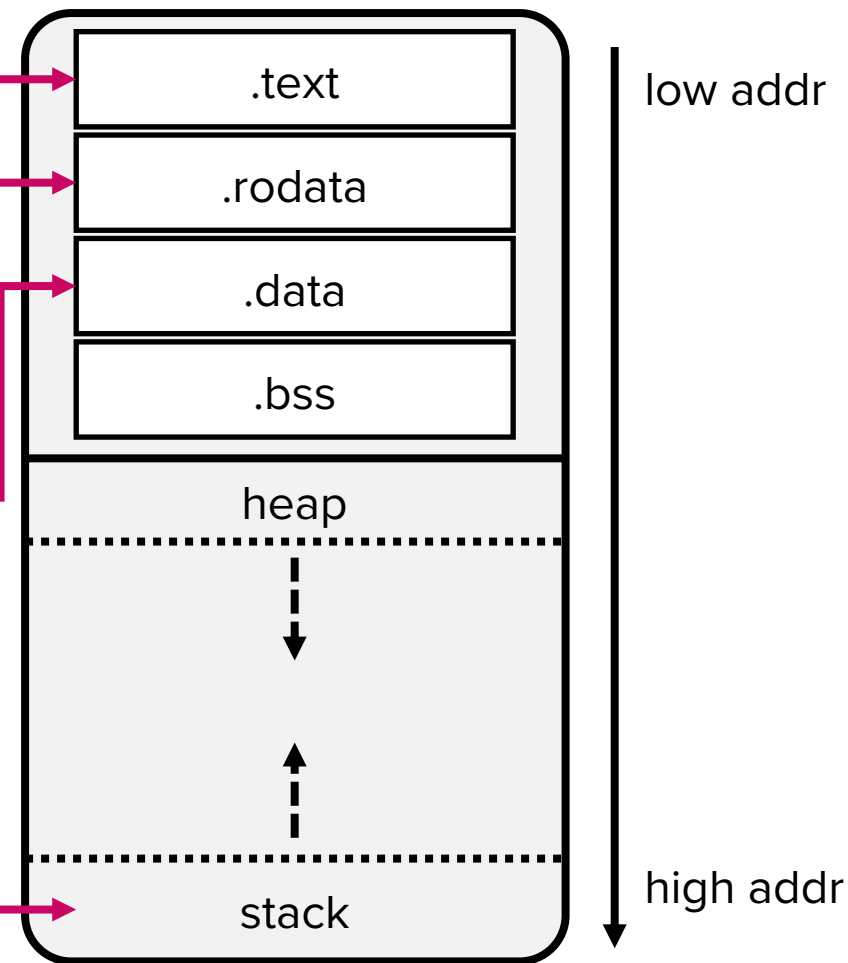
```
pwndbg> r
```

```
pwndbg> vmmmap
```

```
pwndbg> vmmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA

Start      End Perm  Size Offset File
0x8048000 0x8049000 r--p    1000    0 /home/lab01/target
0x8049000 0x804a000 r-xp    1000  1000 /home/lab01/target
0x804a000 0x804b000 r--p    1000  2000 /home/lab01/target
0x804b000 0x804c000 r--p    1000  2000 /home/lab01/target
0x804c000 0x804d000 rw-p    1000  3000 /home/lab01/target
0xf7cc1000 0xf7ce1000 r--p  20000    0 /lib/i386-linux-gnu/libc.so.6
0xf7ce1000 0xf7e63000 r-xp 182000 20000 /lib/i386-linux-gnu/libc.so.6
0xf7e63000 0xf7ee8000 r--p   85000 1a2000 /lib/i386-linux-gnu/libc.so.6
0xf7ee8000 0xf7ee9000 ---p    1000 227000 /lib/i386-linux-gnu/libc.so.6
0xf7ee9000 0xf7eeb000 r--p    2000 227000 /lib/i386-linux-gnu/libc.so.6
0xf7eeb000 0xf7eec000 rw-p    1000 229000 /lib/i386-linux-gnu/libc.so.6
0xf7eec000 0xf7ef6000 rw-p    a000    0 [anon_f7eec]
0xf7eff000 0xf7f01000 rw-p    2000    0 [anon_f7eff]
0xf7f01000 0xf7f05000 r--p    4000    0 [vvar]
0xf7f05000 0xf7f07000 r-xp    2000    0 [vdso]
0xf7f07000 0xf7f08000 r--p    1000    0 /lib/i386-linux-gnu/ld-linux.so.2
0xf7f08000 0xf7f2d000 r-xp   25000  1000 /lib/i386-linux-gnu/ld-linux.so.2
0xf7f2d000 0xf7f3c000 r--p    f000 26000 /lib/i386-linux-gnu/ld-linux.so.2
0xf7f3c000 0xf7f3e000 r--p    2000 34000 /lib/i386-linux-gnu/ld-linux.so.2
0xf7f3e000 0xf7f3f000 rw-p    1000 36000 /lib/i386-linux-gnu/ld-linux.so.2
0xffeca000 0xffeeb000 rw-p   21000    0 [stack]
```

## Virtual memory



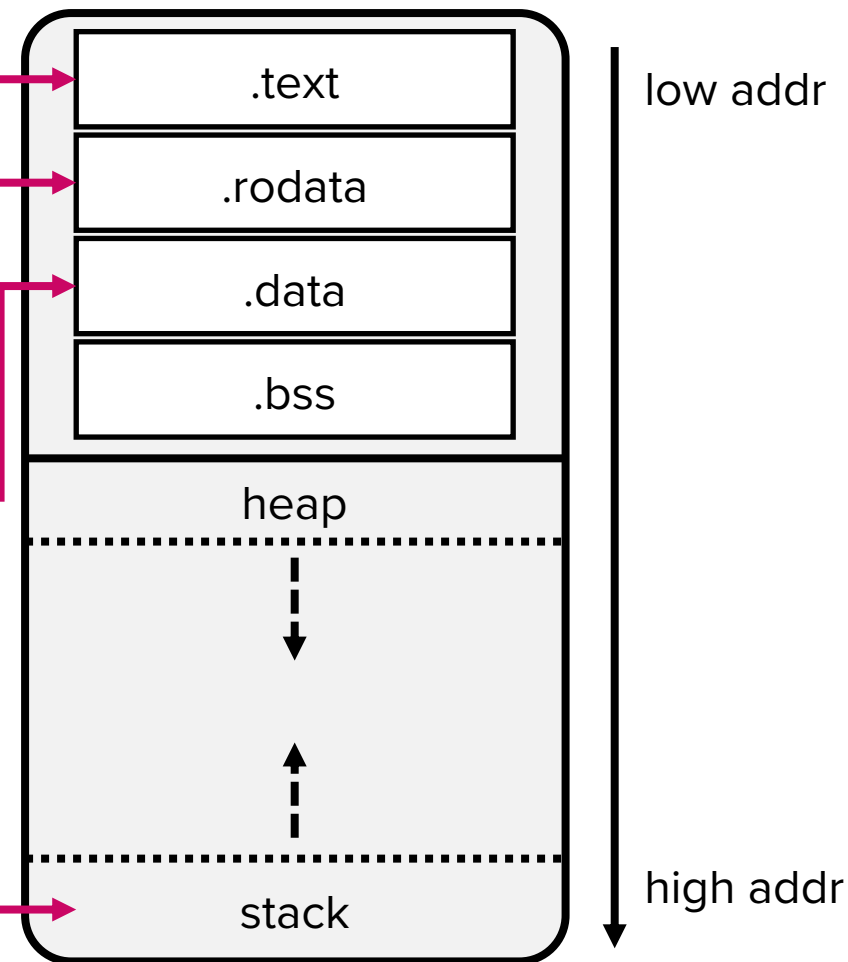
# Exercise: Examine Lab 01's target in GDB

```
pwndbg> elfsections // check section info
pwndbg> x/20i 0x8049110 // exam text sec. instructions
pwndbg> backtrace // backtrace from main (start_)
pwndbg> x/10s 0x804a000 // exam rodata sec. strings
pwndbg> x/100wx 0x804c000 // exam data/bss sec.
```

```
pwndbg> vmmmap
```

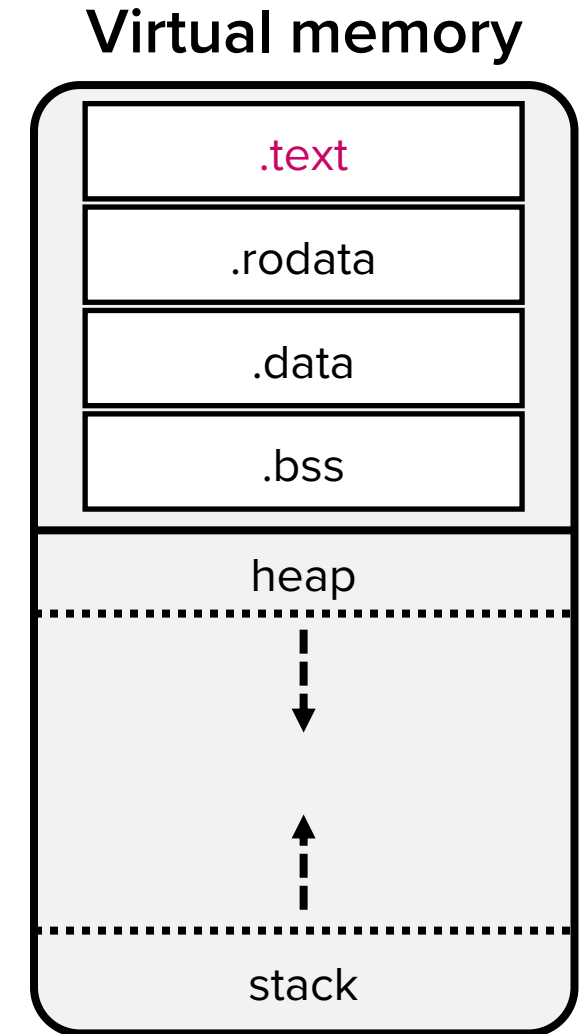
LEGEND: STACK   HEAP   CODE   DATA   RWX   RODATA						
Start	End	Perm	Size	Offset	File	
0x8048000	0x8049000	r--p	1000	0	/home/lab01/target	
0x8049000	0x804a000	r-xp	1000	1000	/home/lab01/target	
0x804a000	0x804b000	r--p	1000	2000	/home/lab01/target	
0x804b000	0x804c000	r--p	1000	2000	/home/lab01/target	
0x804c000	0x804d000	rw-p	1000	3000	/home/lab01/target	
0xf7cc1000	0xf7ce1000	r--p	20000	0	/lib/i386-linux-gnu/libc.so.6	
0xf7ce1000	0xf7e63000	r-xp	182000	20000	/lib/i386-linux-gnu/libc.so.6	
0xf7e63000	0xf7ee8000	r--p	85000	1a2000	/lib/i386-linux-gnu/libc.so.6	
0xf7ee8000	0xf7ee9000	---p	1000	227000	/lib/i386-linux-gnu/libc.so.6	
0xf7ee9000	0xf7eeb000	r--p	2000	227000	/lib/i386-linux-gnu/libc.so.6	
0xf7eeb000	0xf7eec000	rw-p	1000	229000	/lib/i386-linux-gnu/libc.so.6	
0xf7eec000	0xf7ef6000	rw-p	a000	0	[anon_f7eec]	
0xf7eff000	0xf7f01000	rw-p	2000	0	[anon_f7eff]	
0xf7f01000	0xf7f05000	r--p	4000	0	[vvar]	
0xf7f05000	0xf7f07000	r-xp	2000	0	[vdso]	
0xf7f07000	0xf7f08000	r--p	1000	0	/lib/i386-linux-gnu/ld-linux.so.2	
0xf7f08000	0xf7f2d000	r-xp	25000	1000	/lib/i386-linux-gnu/ld-linux.so.2	
0xf7f2d000	0xf7f3c000	r--p	f000	26000	/lib/i386-linux-gnu/ld-linux.so.2	
0xf7f3c000	0xf7f3e000	r--p	2000	34000	/lib/i386-linux-gnu/ld-linux.so.2	
0xf7f3e000	0xf7f3f000	rw-p	1000	36000	/lib/i386-linux-gnu/ld-linux.so.2	
0xffeca000	0xffeeb000	rw-p	21000	0	[stack]	

## Virtual memory



# A process contains instructions and data

- Q1) Which component of a computer executes instructions? → CPU
- Q2) How does that component keep track of which instruction to execute next?  
→ Using instruction pointer (special register)
- Q3) How does that component keep track of program execution states?  
→ Using general purpose registers





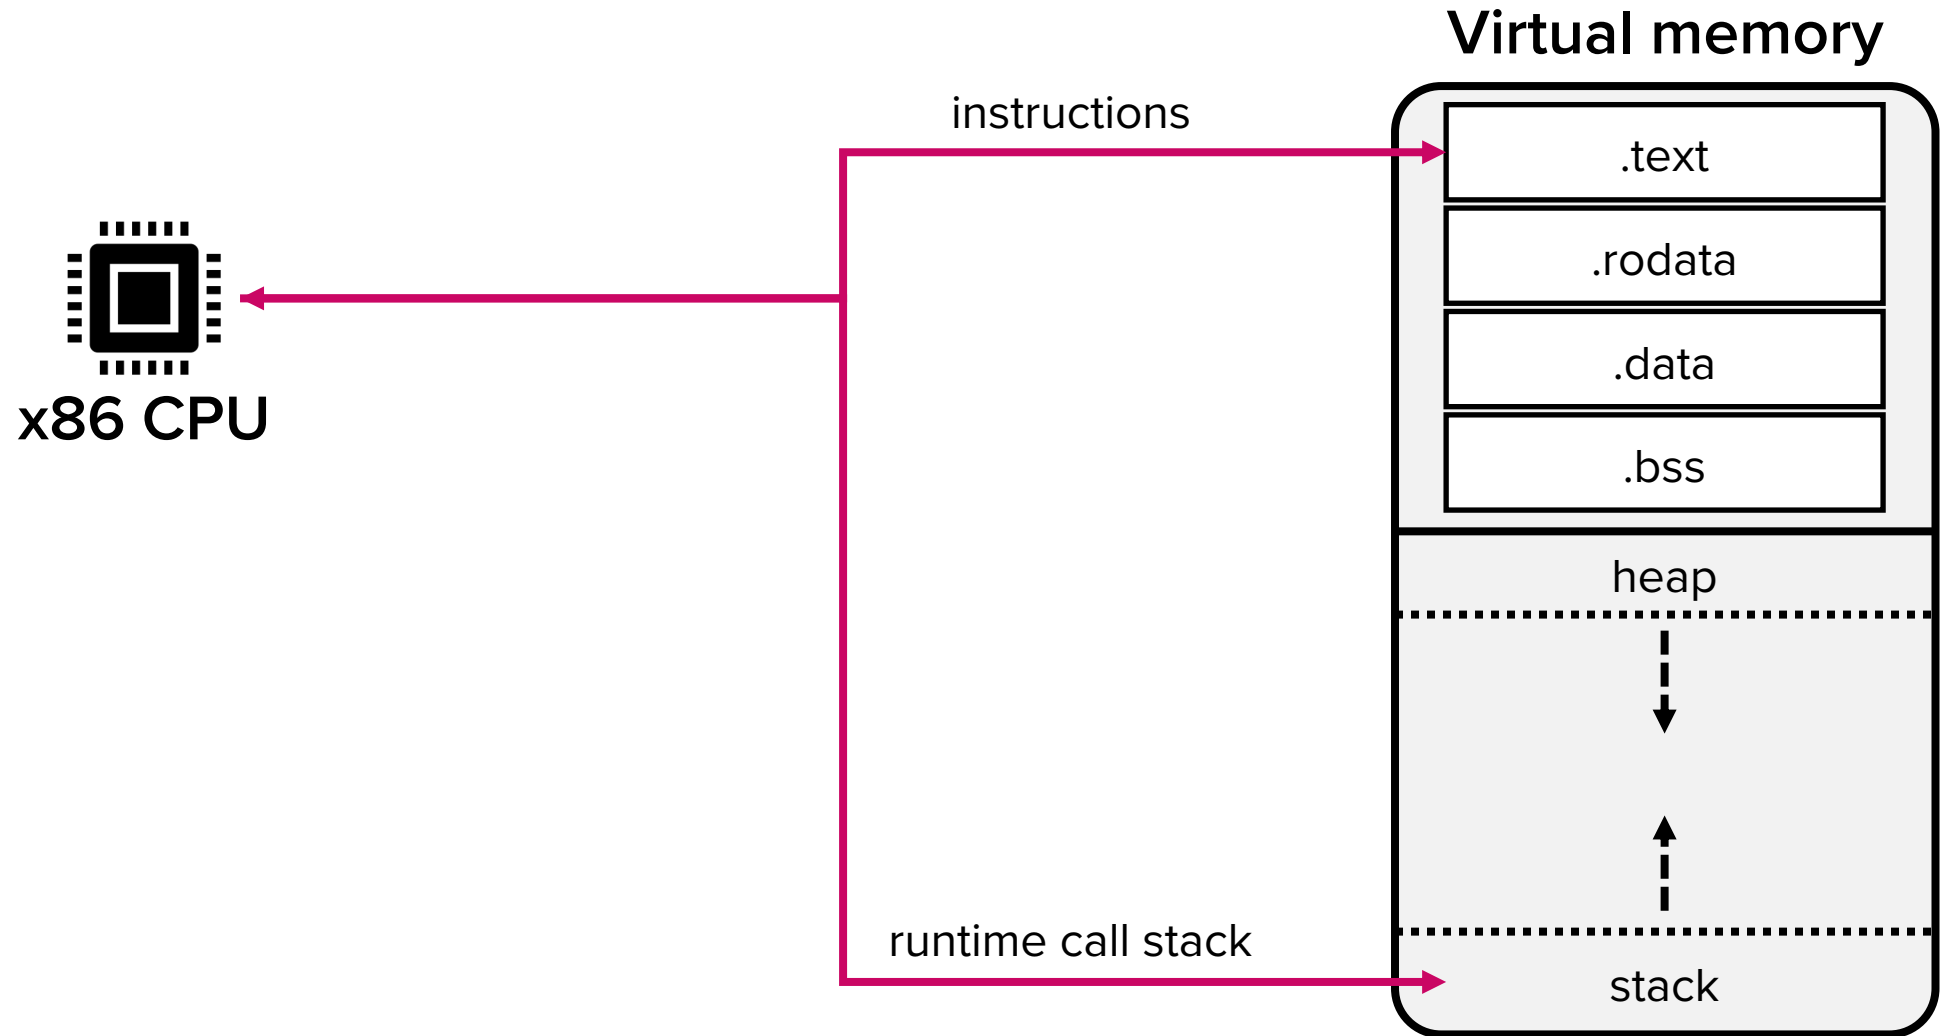
# x86 Registers

# x86 architecture

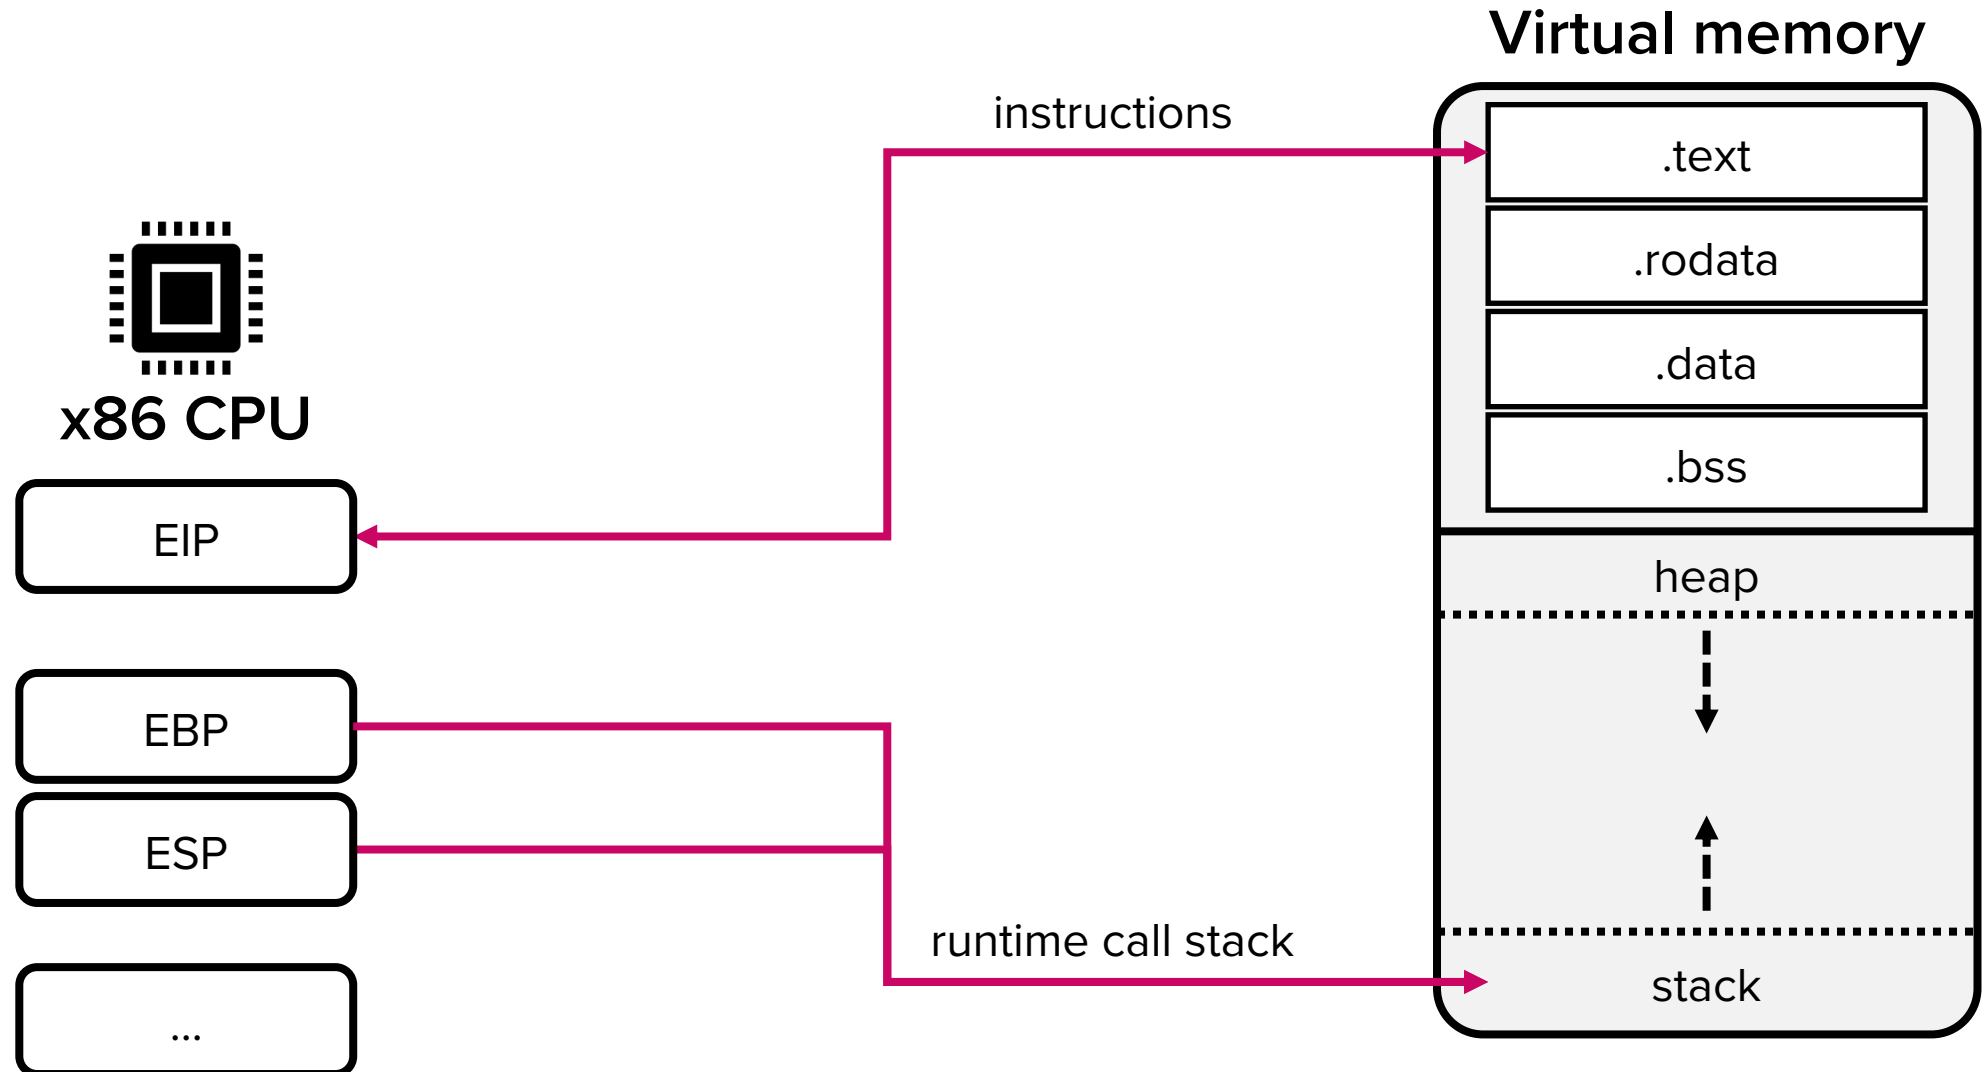
---

- Developed by Intel in 1985
- 32-bit address space
  - Theoretically  $2^{32} = 4\text{GB}$  of virtual memory can be used
- One of the most common architectures

# CPU executes a program



# CPU executes a program using registers

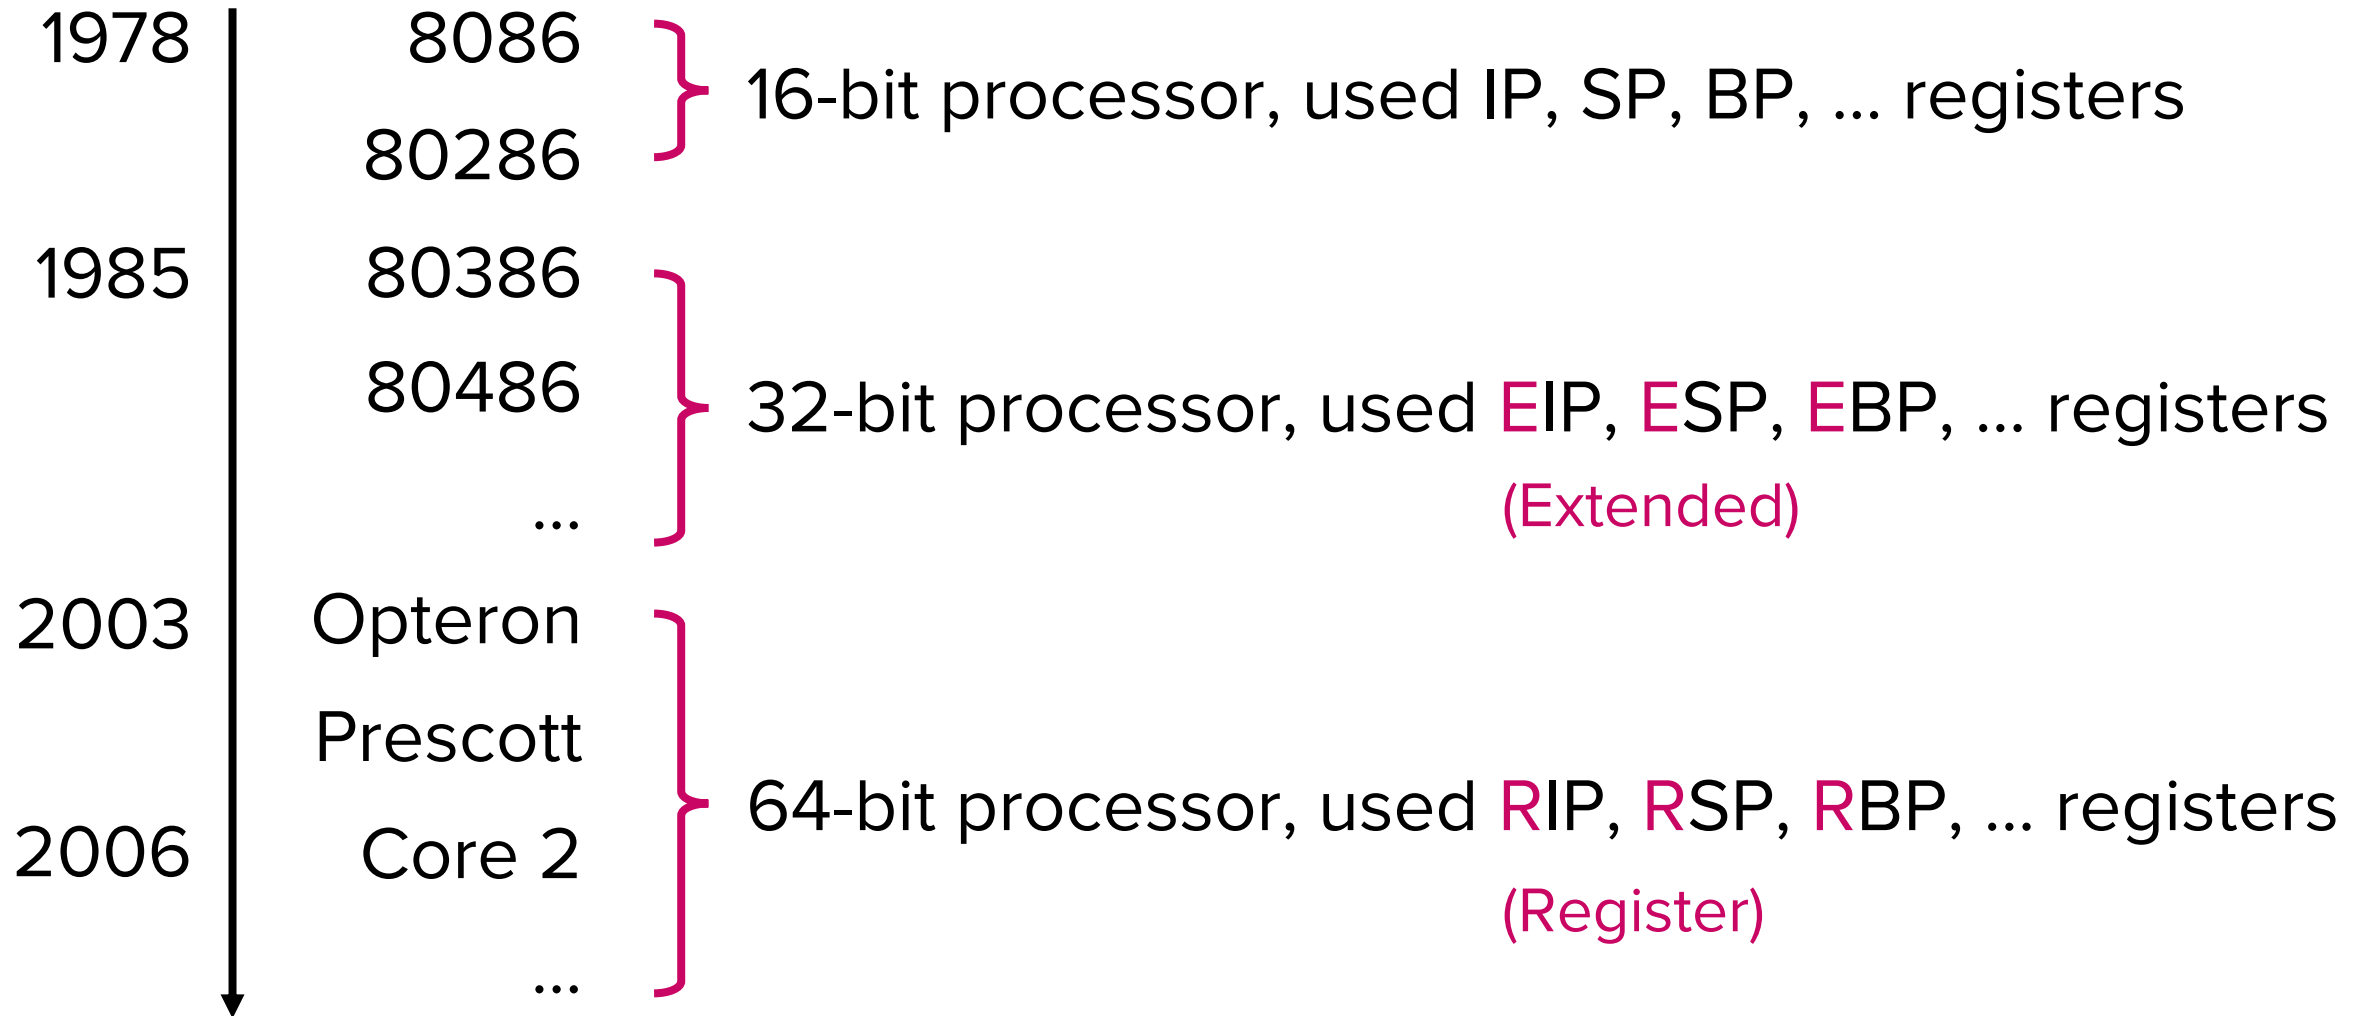


# x86 registers

---

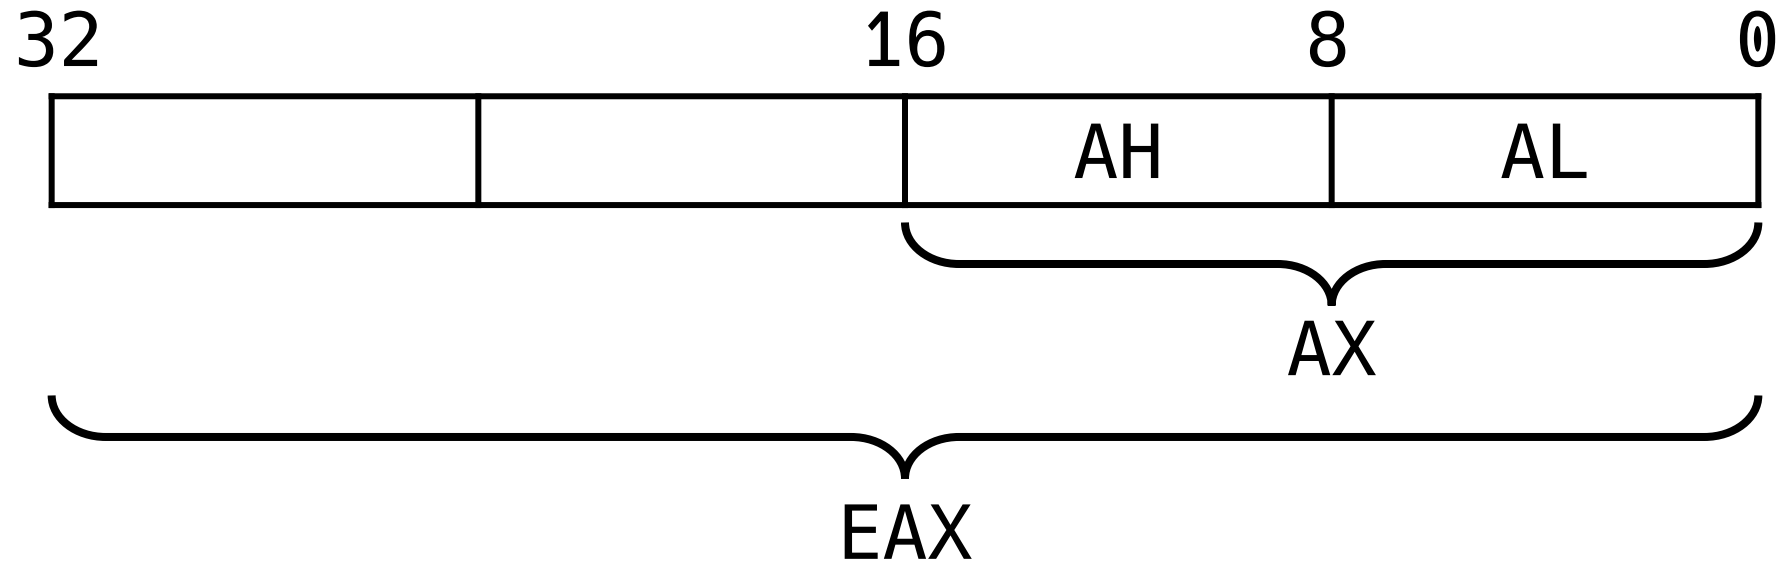
- General purpose registers (GPR)
  - EAX, EBX, ECX, EDX
- Pointers
  - ESI, EDI
- Stack pointers
  - ESP, EBP
- Special purpose registers
  - EIP, EFLAGS

# History of Intel/AMD processors



# x86 registers are extended from 16-bit registers

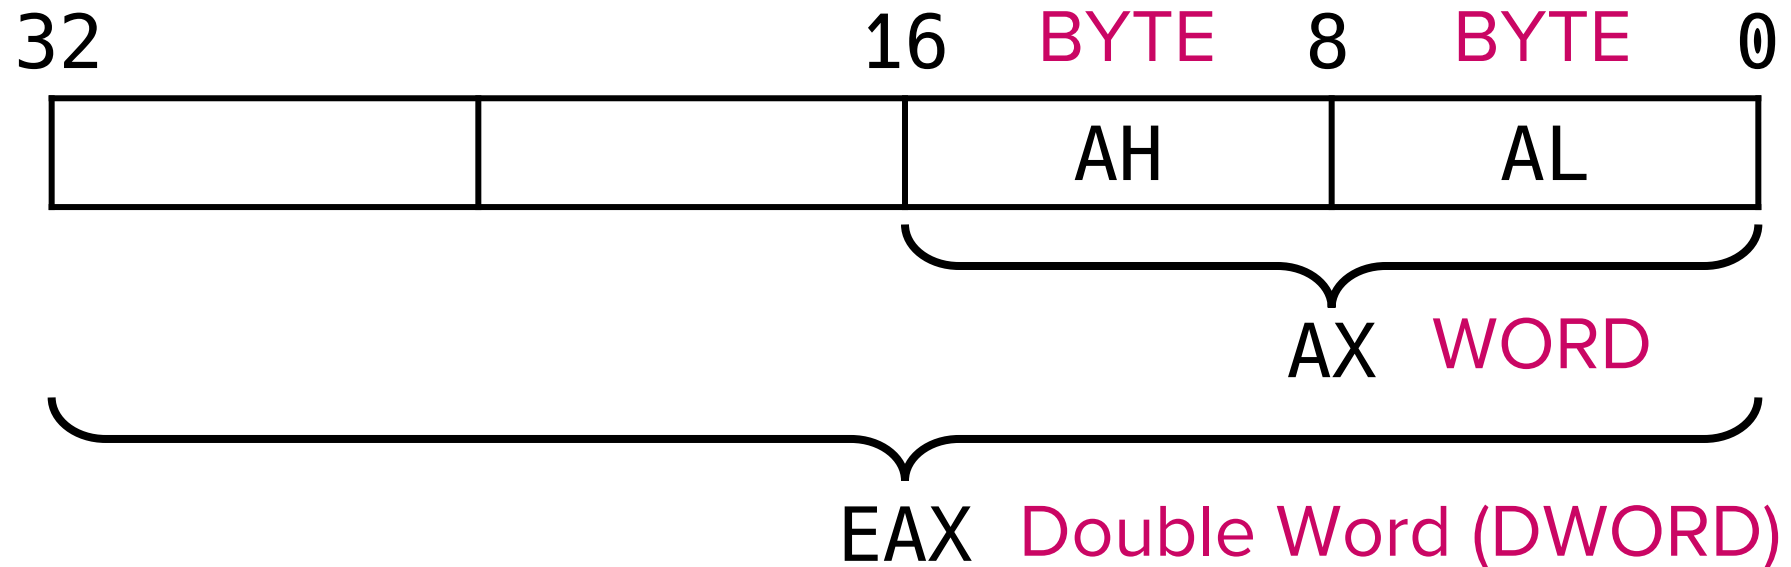
- AL, AH, AX, and EAX are accessible



(Same applies to other general purpose registers)

Q) Can you draw RAX for x86\_64?

# x86 size convention: Word is 16 bits



RAX: Quad Word (QWORD)



# x86 Assembly

# Basic format of x86 instructions

- x86 instructions have 0, 1, or 2 operands
  - 0 operand: `ret`
  - 1 operand: `inc eax`  
operand 1
  - 2 operands: `mov eax, ebx`  
operand 1  
operand 2

# Basic semantics

- Operand 1 (usually) stores the result

```
mov    eax, ebx    // eax = ebx
```

```
sub    esp, 0xc    // esp = esp - 0x8
```

```
inc    eax         // eax = eax + 1
```

# Operand types

- Operand can be register, memory, or constant

```
mov  eax, [ebx]
```

register    memory pointed to by ebx

```
sub  esp, 0xc
```

constant

```
mov  al, BYTE ptr [ecx]
```

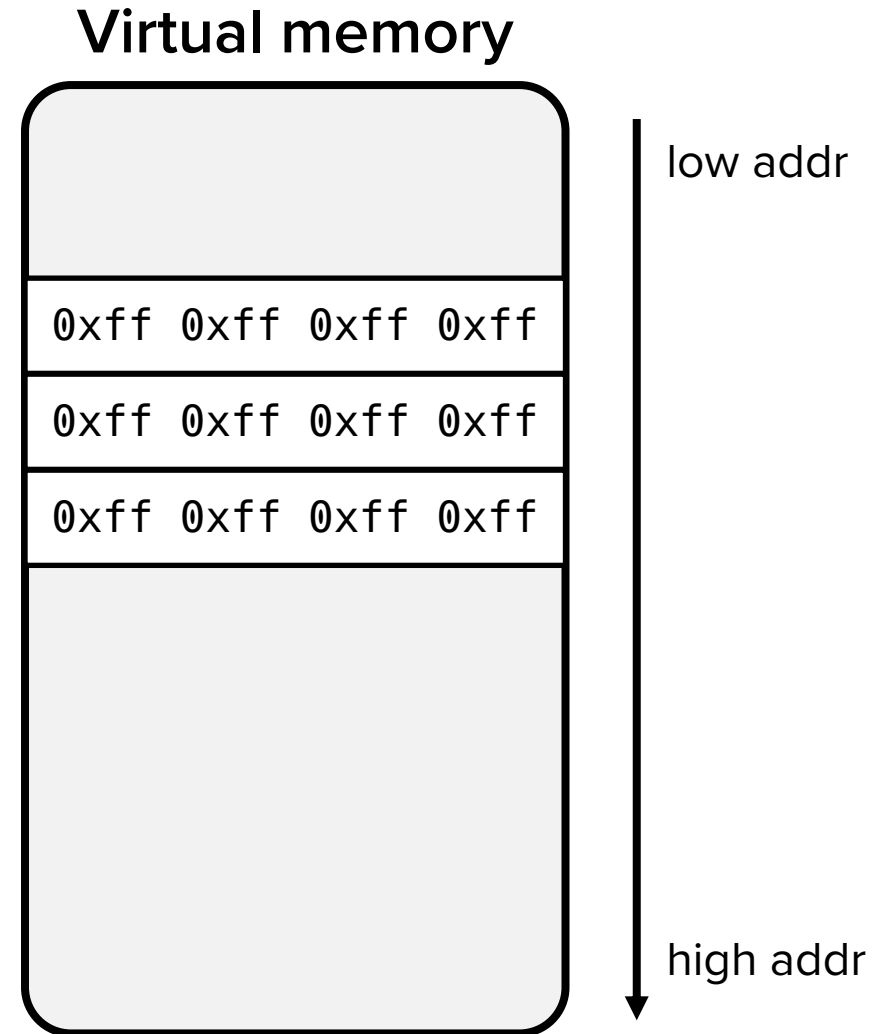
???

memory pointed to by ecx

# Size directives – motivation

```
mov [ebx], 2
```

ebx = 0x804c0a0 → 0x804c0a0  
0x804c0a4  
0x804c0a8



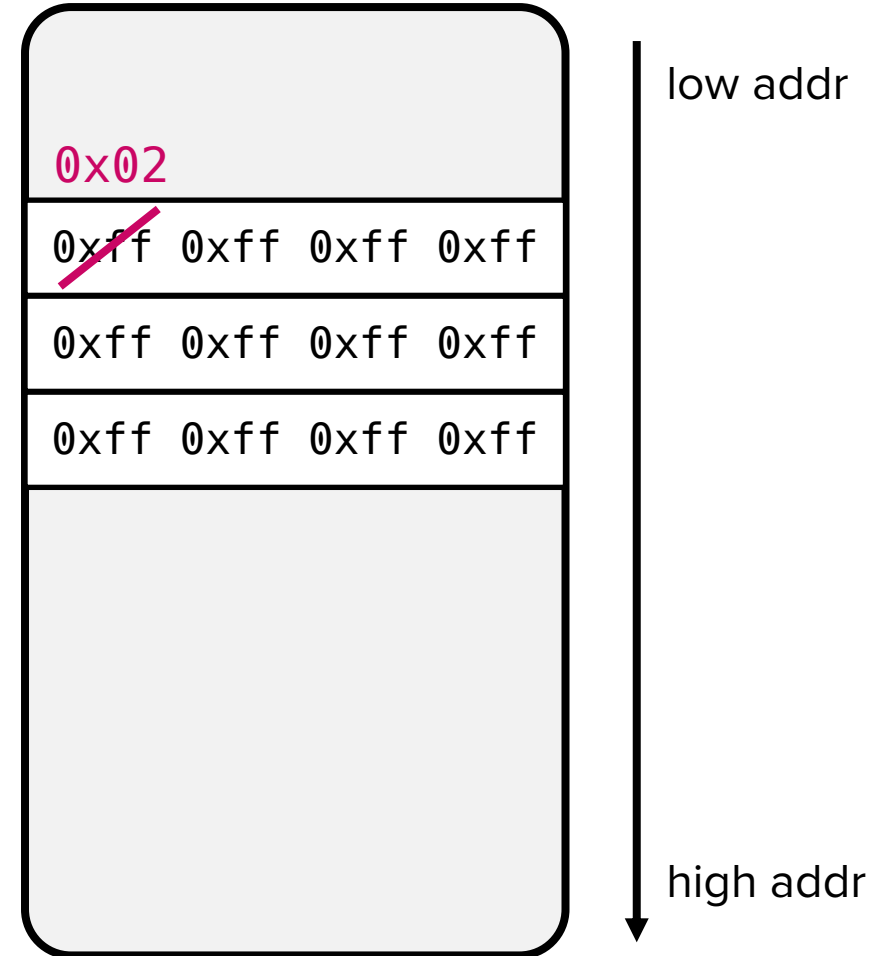
Q) Correct behavior?

# Size directives – motivation

```
mov [ebx], 2
```

ebx = 0x804c0a0 → 0x804c0a0  
0x804c0a4  
0x804c0a8

Virtual memory



Q) Correct behavior?

1. Overwrite a byte from **0x804c0a0** as 2
2. Overwrite 4 bytes from **0x804c0a0** as 2

# Size directives – motivation

```
mov [ebx], 2
```

ebx = 0x804c0a0 → 0x804c0a0  
0x804c0a4  
0x804c0a8

Virtual memory



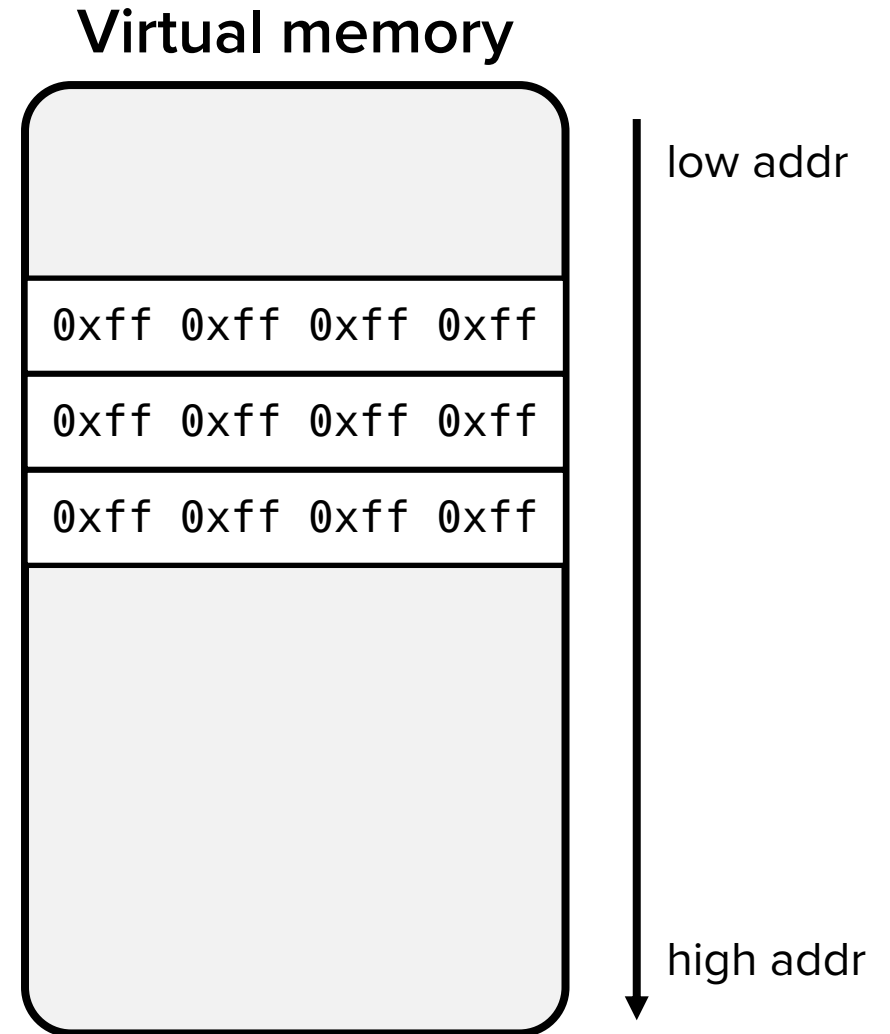
Q) Correct behavior?

1. Overwrite a byte from 0x804c0a0 as 2
2. Overwrite 4 bytes from 0x804c0a0 as 2

# Size directives – motivation

```
mov [ebx], 2
```

ebx = 0x804c0a0 → 0x804c0a0  
0x804c0a4  
0x804c0a8



Q) Correct behavior?

1. Overwrite a byte from **0x804c0a0** as 2
2. Overwrite 4 bytes from **0x804c0a0** as 2

A) Both are incorrect. The instruction leads to an error due to ambiguity.



# Size directives: BYTE/WORD/DWORD PTR

```
mov [ebx], 2
```

ERROR



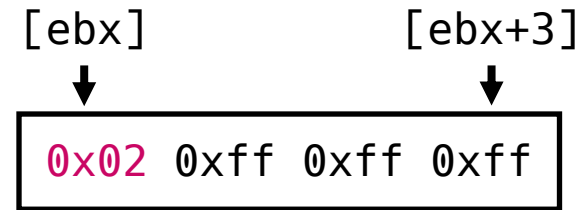
```
mov BYTE PTR [ebx], 2
```

or

```
mov WORD PTR [ebx], 2
```

or

```
mov DWORD PTR [ebx], 2
```



# Endianness

```
mov [ebx], 0xdeadbeef
```

ebx = 0x804c0a0 → 0x804c0a0

0x804c0a1

0x804c0a2

0x804c0a3

(1)

(2)

vs.

low addr

high addr

Q) Which is correct?

# Endianness: How a sequence of bytes is stored in memory

```
mov [ebx], 0xdeadbeef
```

MSB (most significant bit) ↑

LSB (least SB) ←

ebx = 0x804c0a0 → 0x804c0a0

0x804c0a1

0x804c0a2

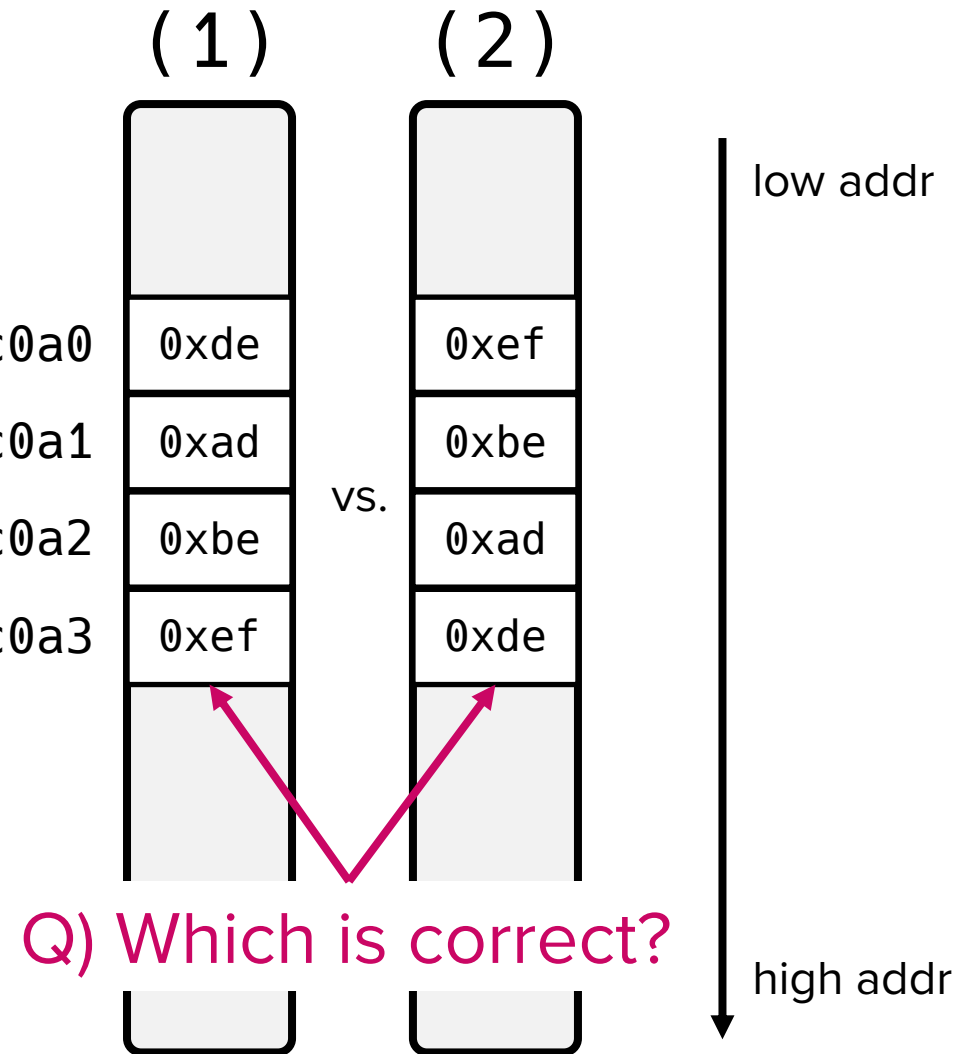
0x804c0a3

(1) is correct if Big Endian

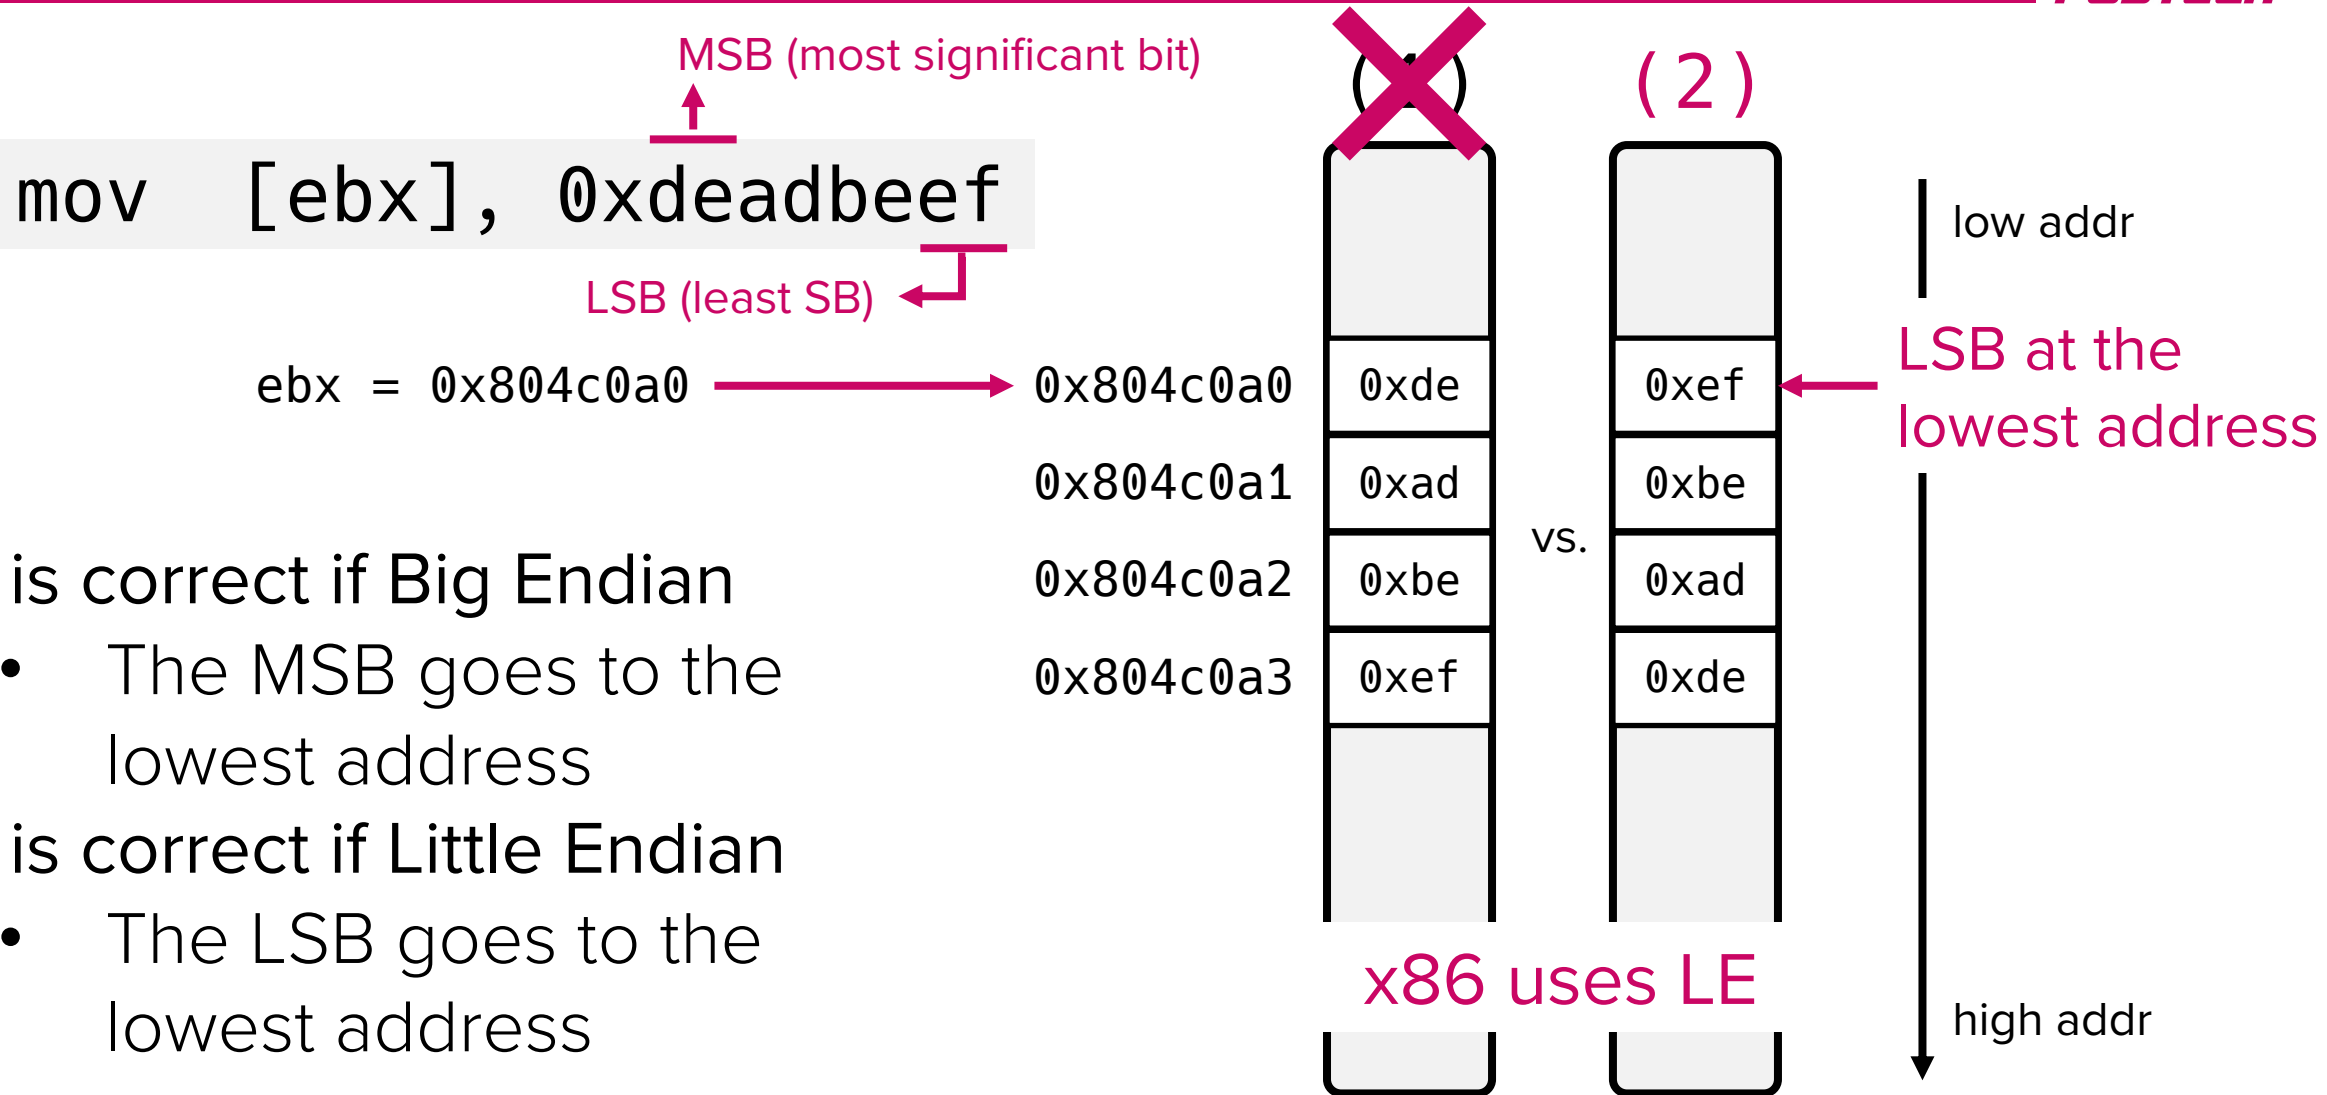
- The MSB goes to the lowest address

(2) is correct if Little Endian

- The LSB goes to the lowest address



# Endianness: x86 uses Little Endian (LE)



# mov instruction syntax w/ exercise

---

- Legal

1. `mov reg, reg`
2. `mov reg, mem`
3. `mov mem, reg`
4. `mov reg, const`
5. `mov mem, const`

- Illegal

- `mov mem, mem`
- `mov const, const`

# mov instruction syntax w/ exercise

- Legal

1. mov reg, reg
2. mov reg, mem
3. mov mem, reg
4. mov reg, const
5. mov mem, const

- Illegal

- mov mem, mem
- mov const, const

```
mov eax, ebx
```

Legal (R to R)

```
mov al, bl
```

Legal (R to R)

```
mov [eax], ebx
```

Legal (R to M)

```
mov eax, [ebx]
```

Legal (M to R)

```
mov eax, [ebx + edx * 4]
```

Legal (M to R)

```
mov al, BYTE PTR [esi]
```

Legal (M to R)

```
mov eax, 16
```

Legal (C to R)

```
mov [ebx], 16
```

Legal (C to M)

```
mov [eax], [ebx]
```

ILLEGAL (M to M)

# lea: Load effective address

```
lea  eax, [ebx]
```

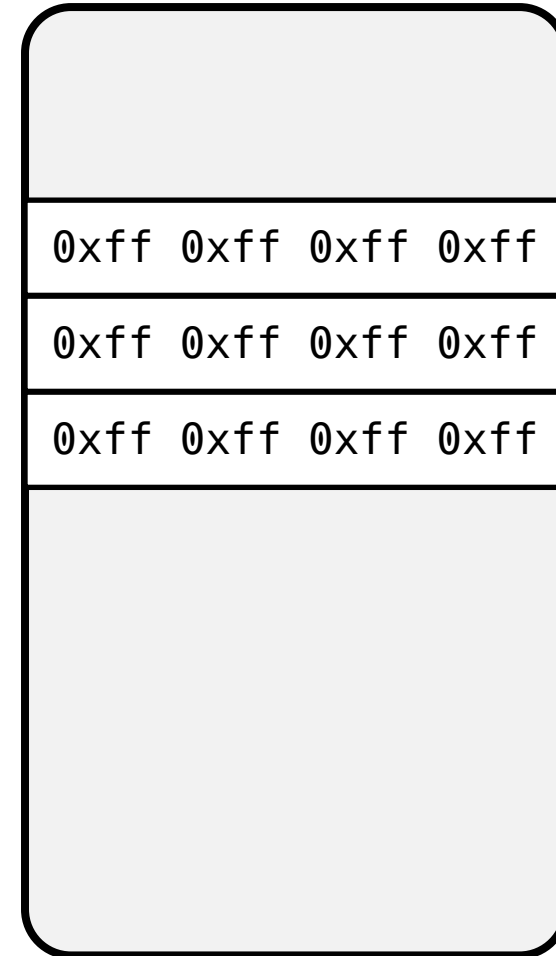
ebx = 0x804c0a0 → 0x804c0a0

0x804c0a4

0x804c0a8

Q) eax = ?

Virtual memory



low addr

high addr

# lea: Load effective address

```
lea  eax, [ebx]
```

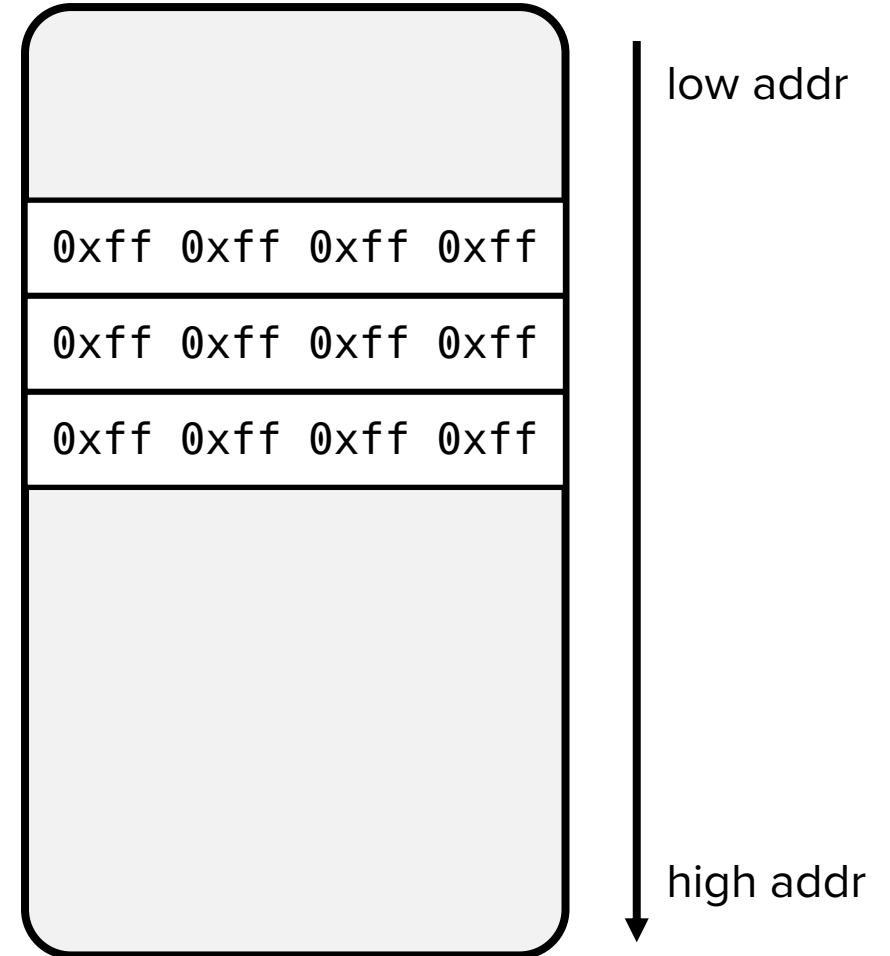
ebx = 0x804c0a0 → 0x804c0a0

0x804c0a4

0x804c0a8

A) **eax = 0x0804c0a0**  
The address is loaded!

Virtual memory





# mov vs. lea

```
mov    esp, [ebp-0x8]
```

// esp = value that address 0xffffd1450 points to  
i.e., esp = \*(ebp - 0x8);

```
lea    esp, [ebp-0x8]
```

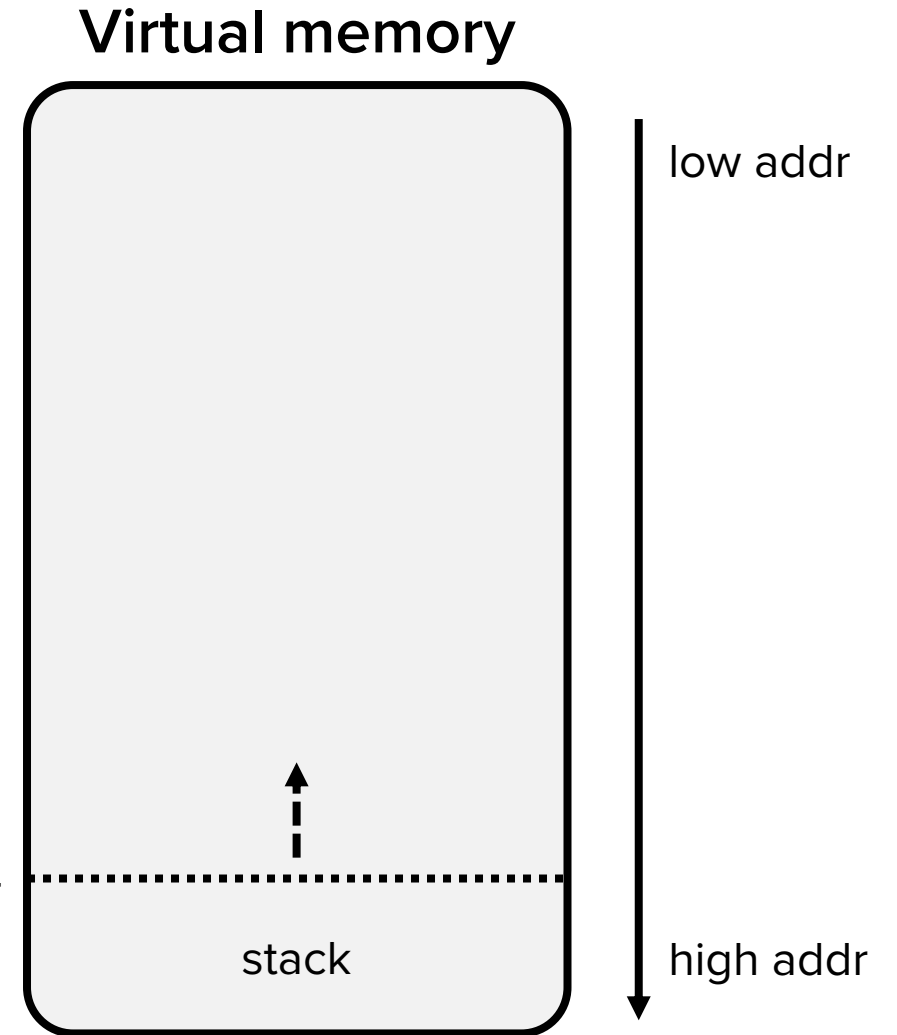
// esp = 0xffffd1450  
i.e., esp = (ebp - 0x8);

ebp = 0xffffd1458

# Stack operations

- `esp` points to the top of the stack

`esp = 0xffffd1404`  `0xffffd1404`



# Stack operations: push

- push enlarges the stack

```
push    eax
```

Push the value  
eax holds

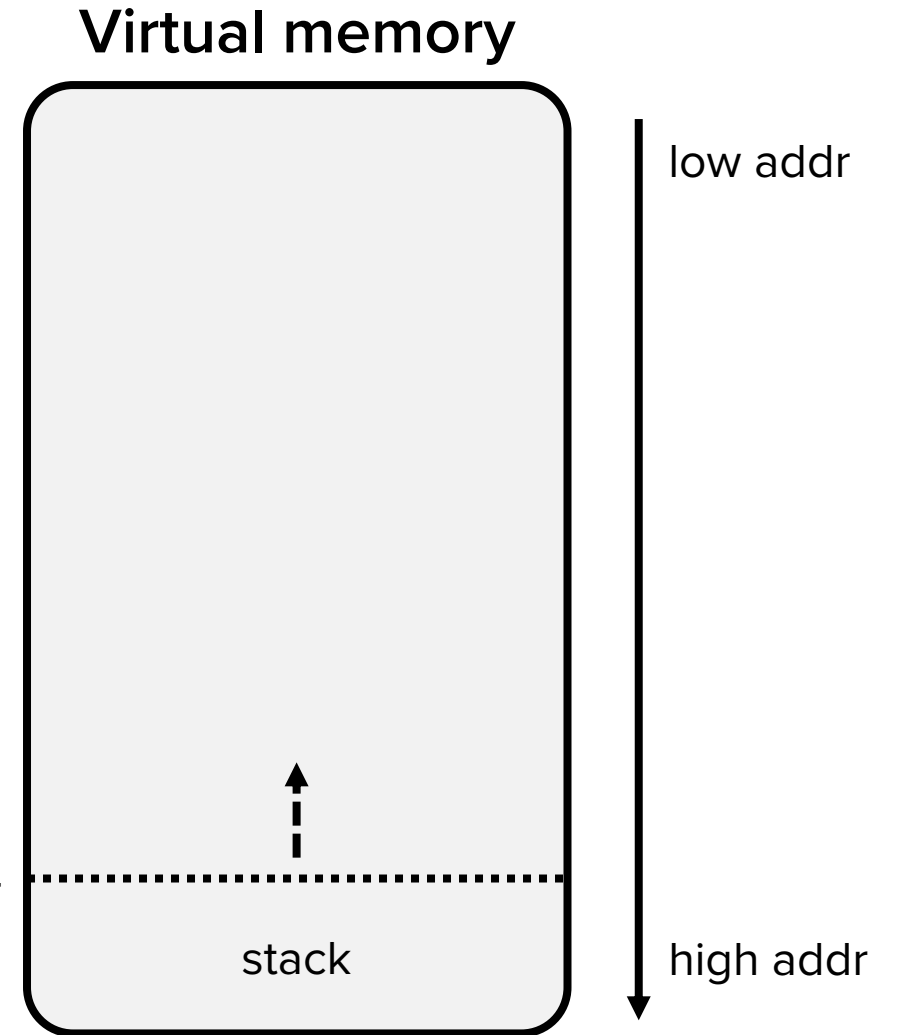
```
push    [eax]
```

Push the value at the  
address eax points to

```
push    0xdeadbeef
```

Push constant value

`esp` = 0xffffd1404 → 0xffffd1404



# Stack operations: push

- push enlarges the stack

```
push    eax
```

Push the value  
eax holds

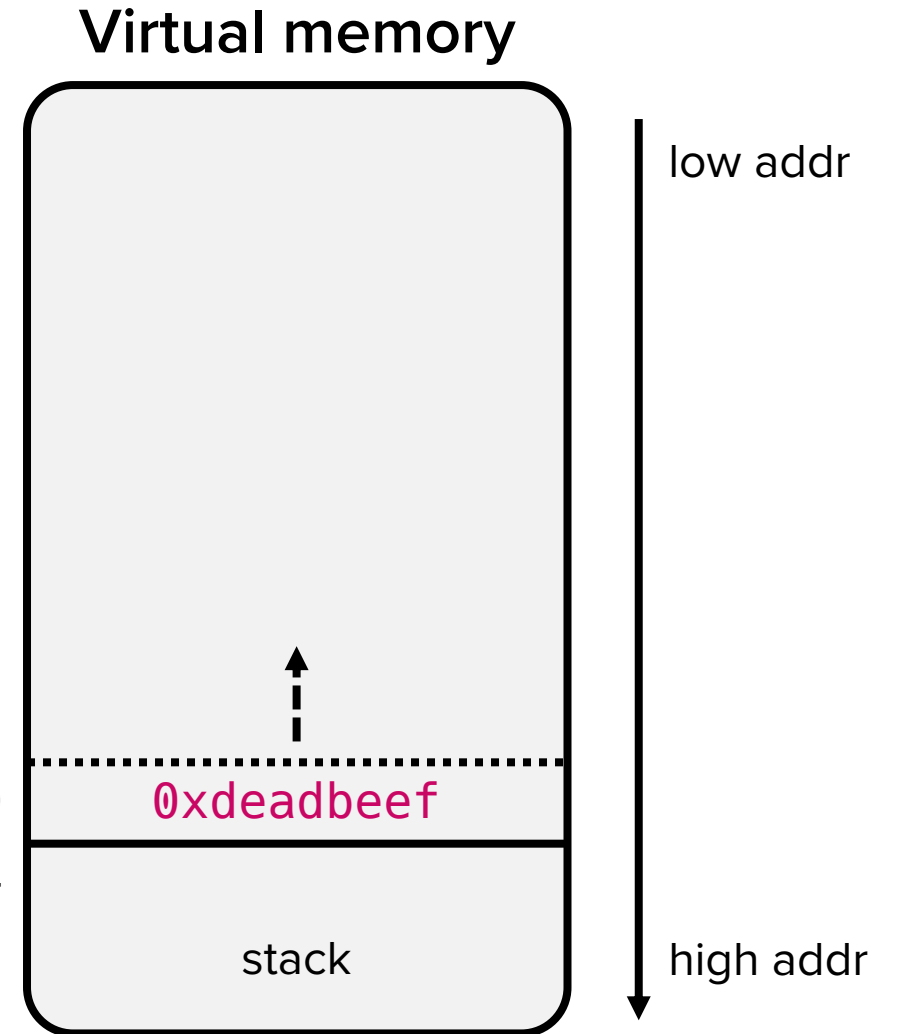
```
push    [eax]
```

Push the value at the  
address eax points to

```
push    0xdeadbeef
```

Push constant value

`esp = 0xffffd1400` → `0xffffd1400`  
`0xffffd1404`

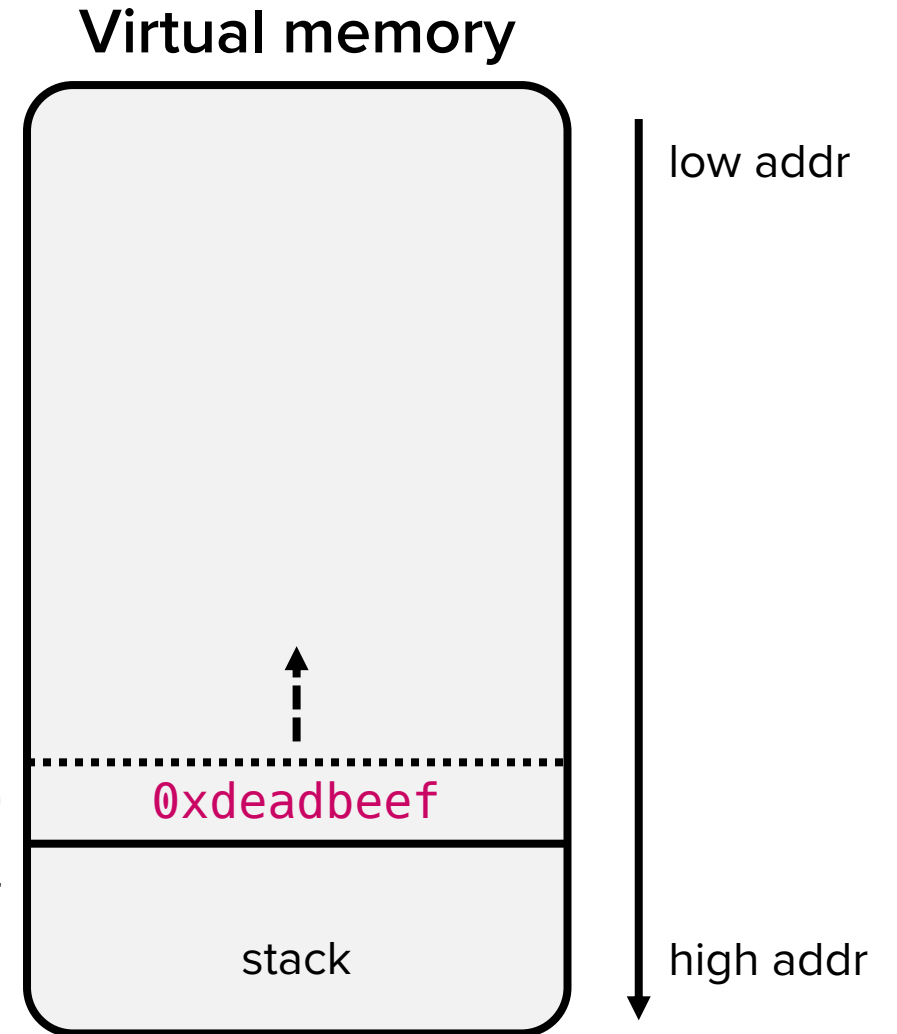


# Stack operations: push

- push equivalence

```
push    eax    ==    sub    esp, 4  
                        mov    [esp], eax
```

`esp = 0xffffd1400`  `0xffffd1400`  
`0xffffd1404`



# Stack operations: pop

- pop shrinks the stack

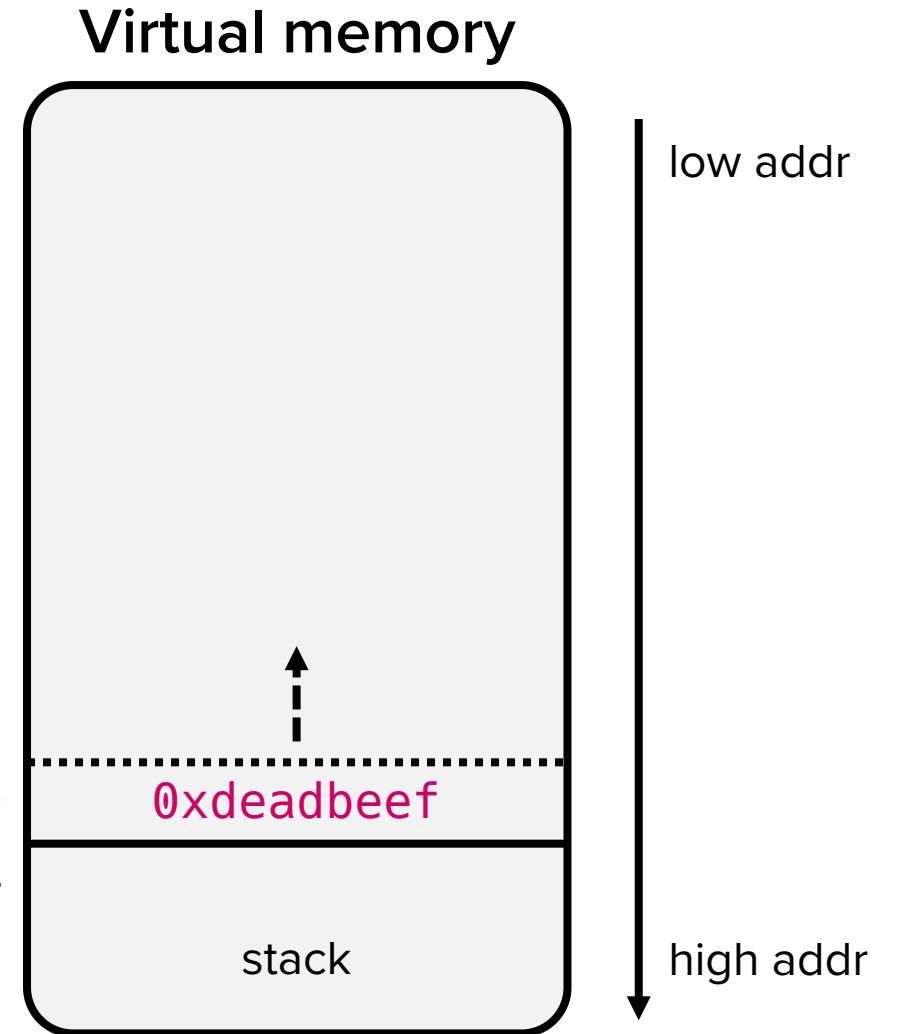
```
pop    eax
```

Pop the stack top  
into eax

```
pop    [eax]
```

Pop the stack top  
into the memory  
pointed to by eax

`esp = 0xffffd1400` → `0xffffd1400`  
`0xffffd1404`



# Stack operations: pop

- pop shrinks the stack

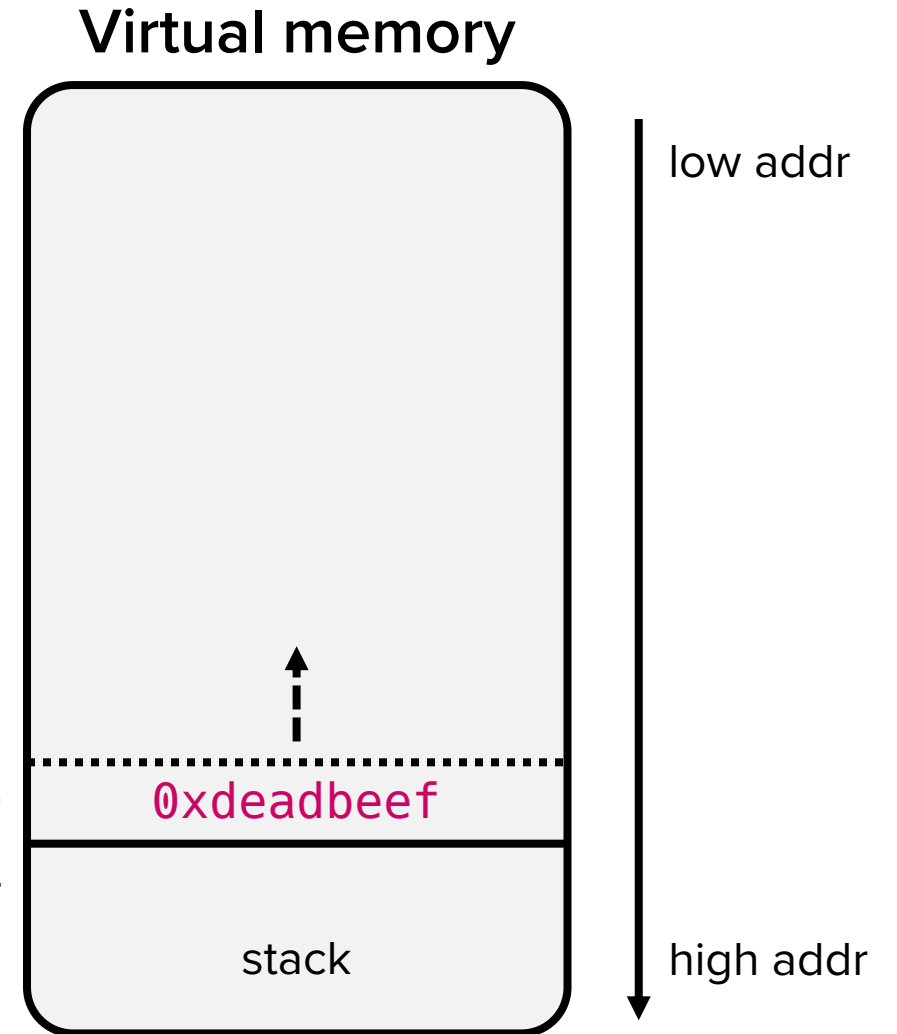
```
pop    eax
```

Pop the stack top  
into eax

```
pop    [eax]
```

Pop the stack top  
into the memory  
pointed to by eax

`esp = 0xffffd1404` → `0xffffd1400`  
`eax = 0xdeadbeef`

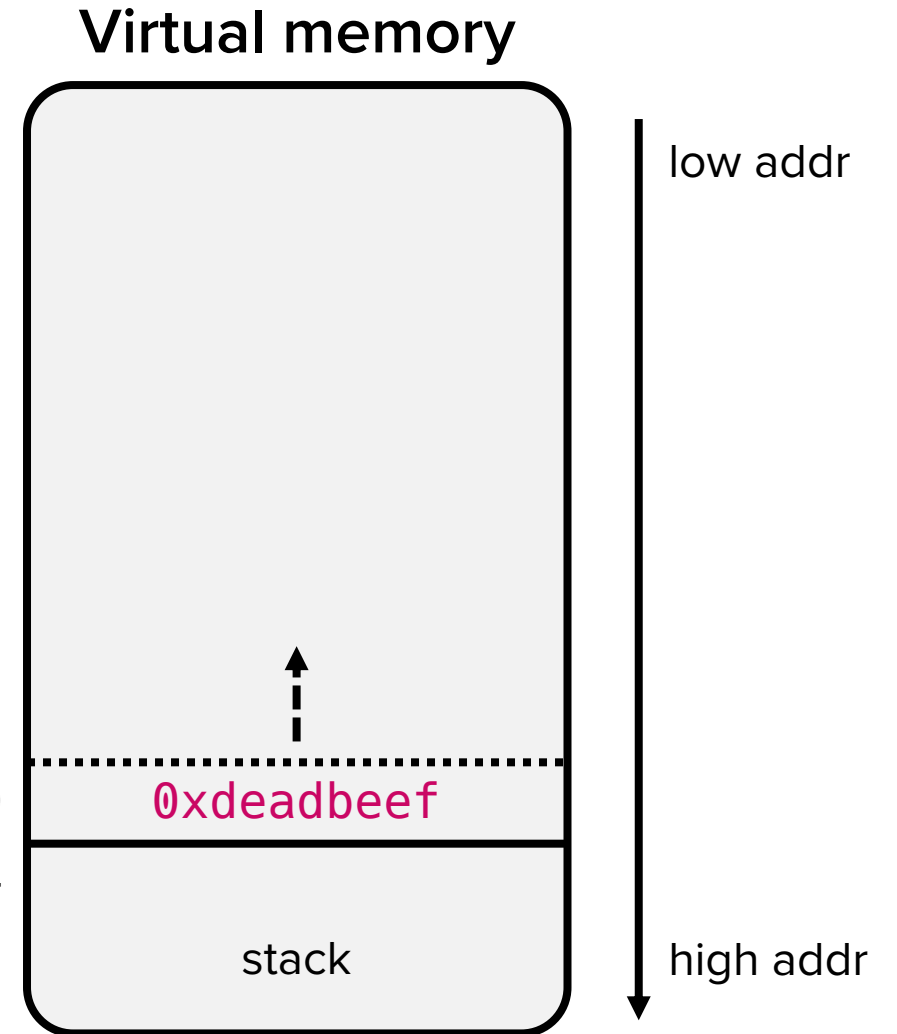


# Stack operations: pop

- pop equivalence

```
pop    eax    ==    mov    eax, [esp]
                        add    esp, 4
```

`esp = 0xffffd1404` → `0xffffd1400`  
`eax = 0xdeadbeef`

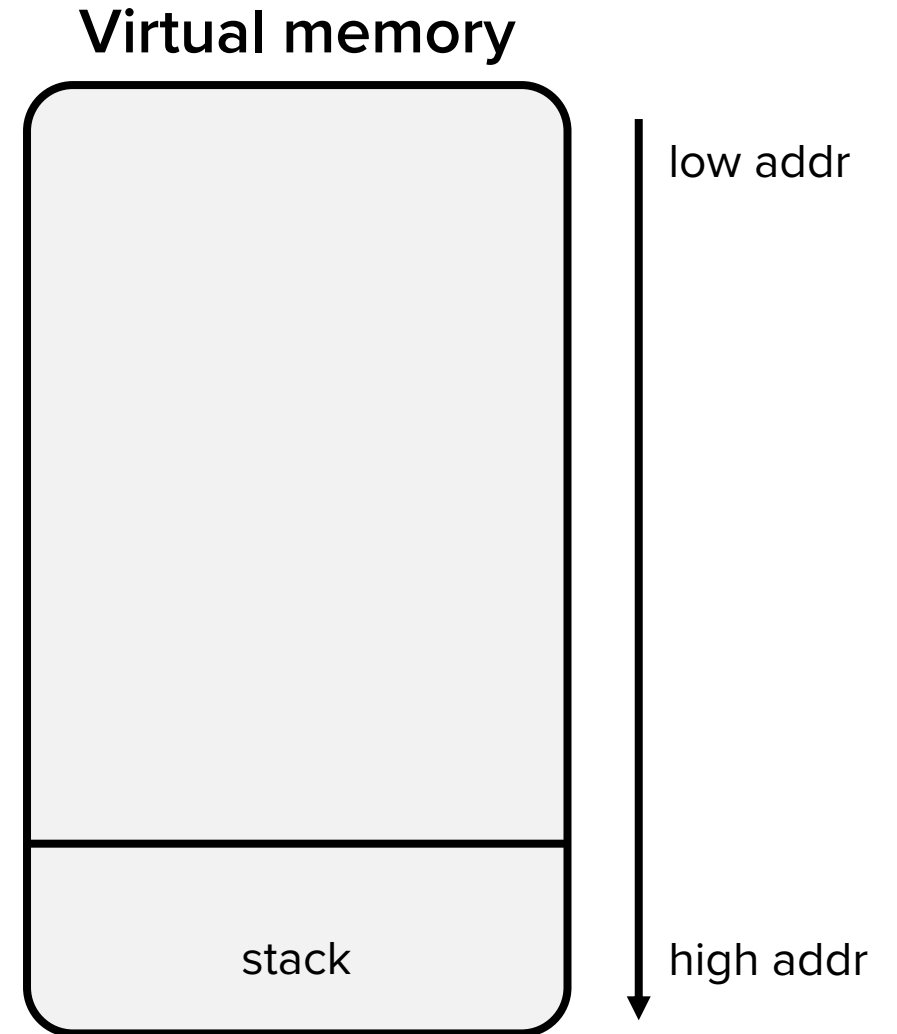




# Subroutine instructions: `call`

- `call <label>`
  - **push** the address to return to
  - perform unconditional **jmp** to `<label>`
  - e.g., in `main()` of `lab01`'s target

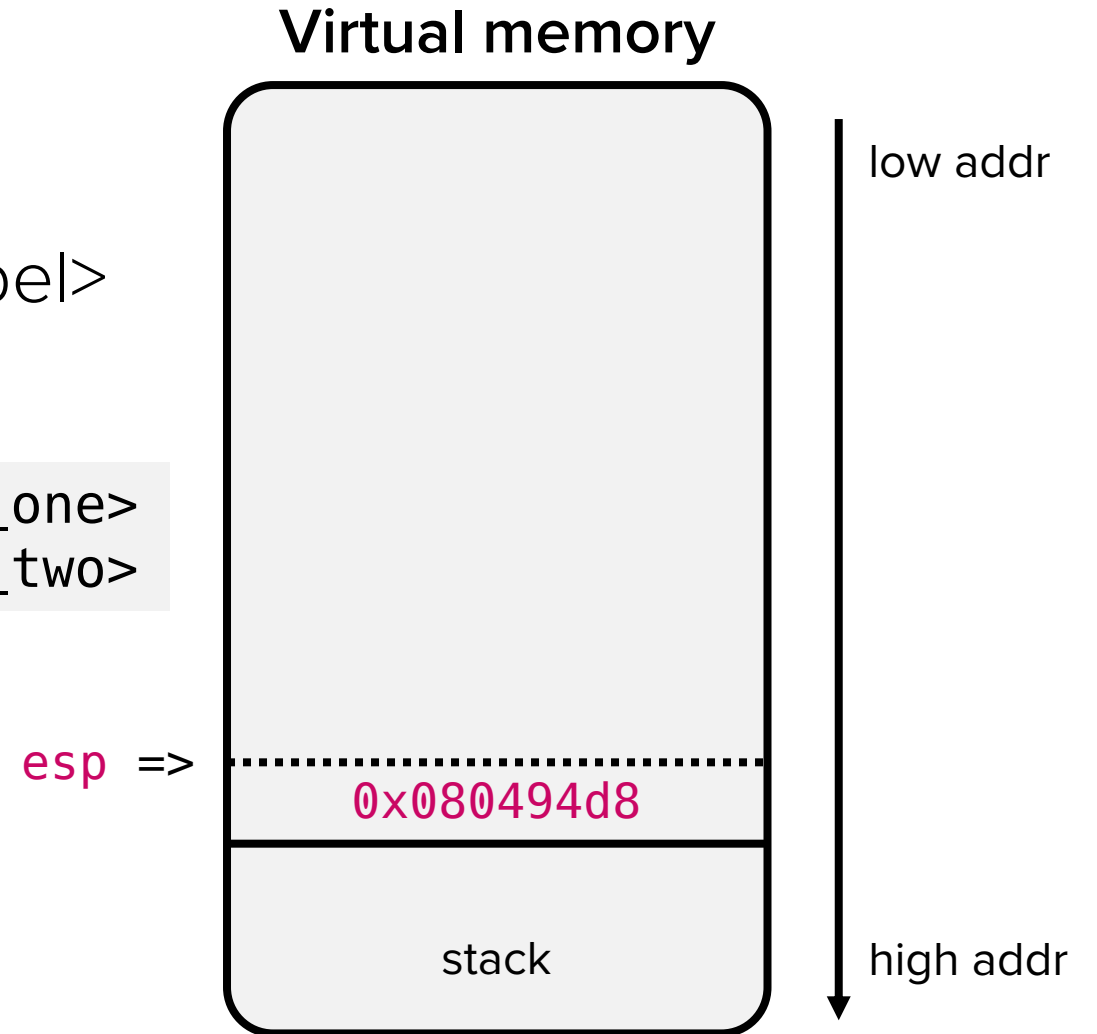
`eip => 0x080494d3: call 0x8049385 <phase_one>`  
`0x080494d8: call 0x8049352 <phase_two>`



# Subroutine instructions: `call`

- `call <label>`
  - **push** the address to return to
  - perform unconditional **jmp** to `<label>`
  - e.g., in `main()` of `lab01`'s target

`eip => 0x080494d3: call 0x8049385 <phase_one>`  
`0x080494d8: call 0x8049352 <phase_two>`

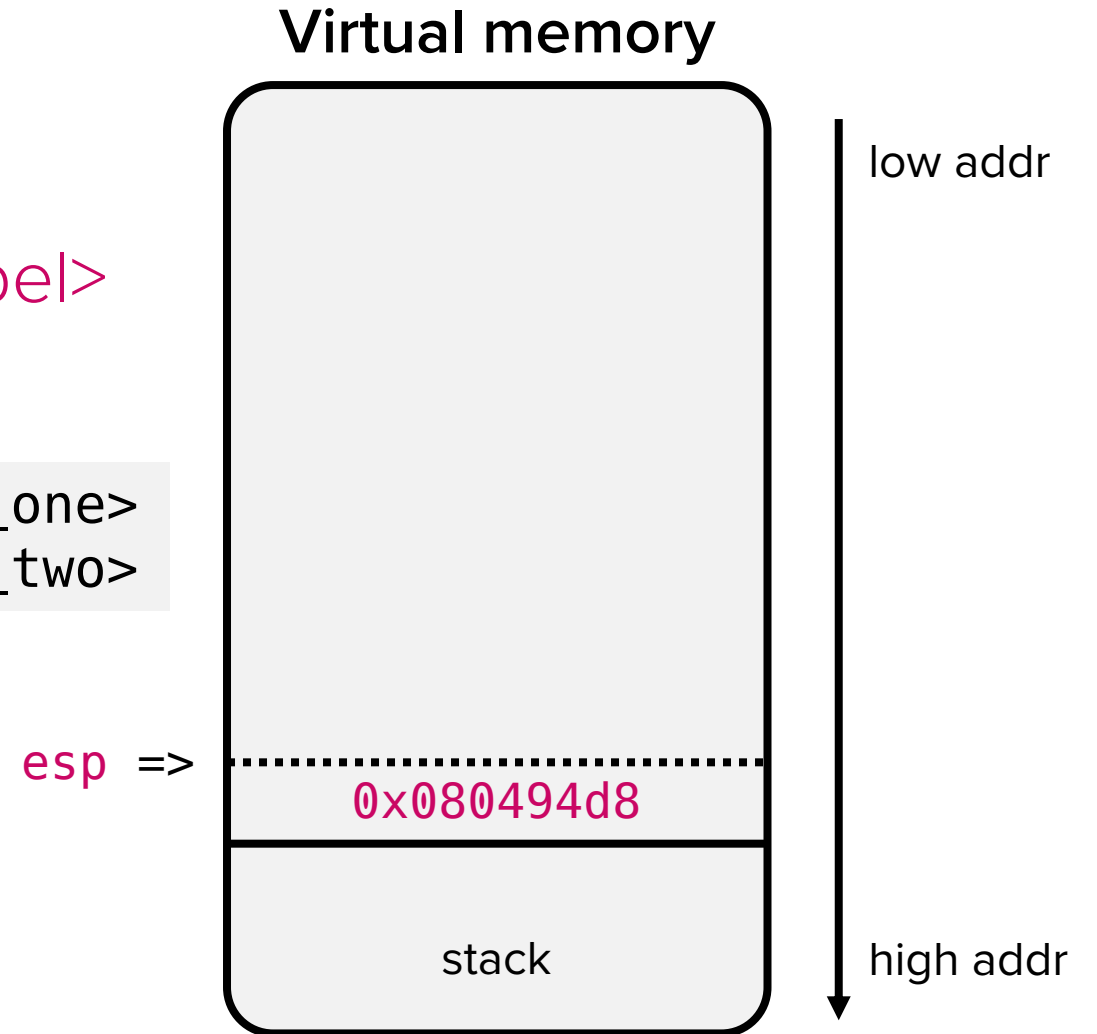


# Subroutine instructions: `call`

- `call <label>`
  - **push** the address to return to
  - perform unconditional **jmp** to `<label>`
  - e.g., in `main()` of `lab01`'s target

```
0x080494d3: call 0x8049385 <phase_one>
0x080494d8: call 0x8049352 <phase_two>
```

```
eip => <phase_one>
0x08049385: push ebp
0x08049386: mov  ebp, esp
...
```



# Subroutine instructions: ret

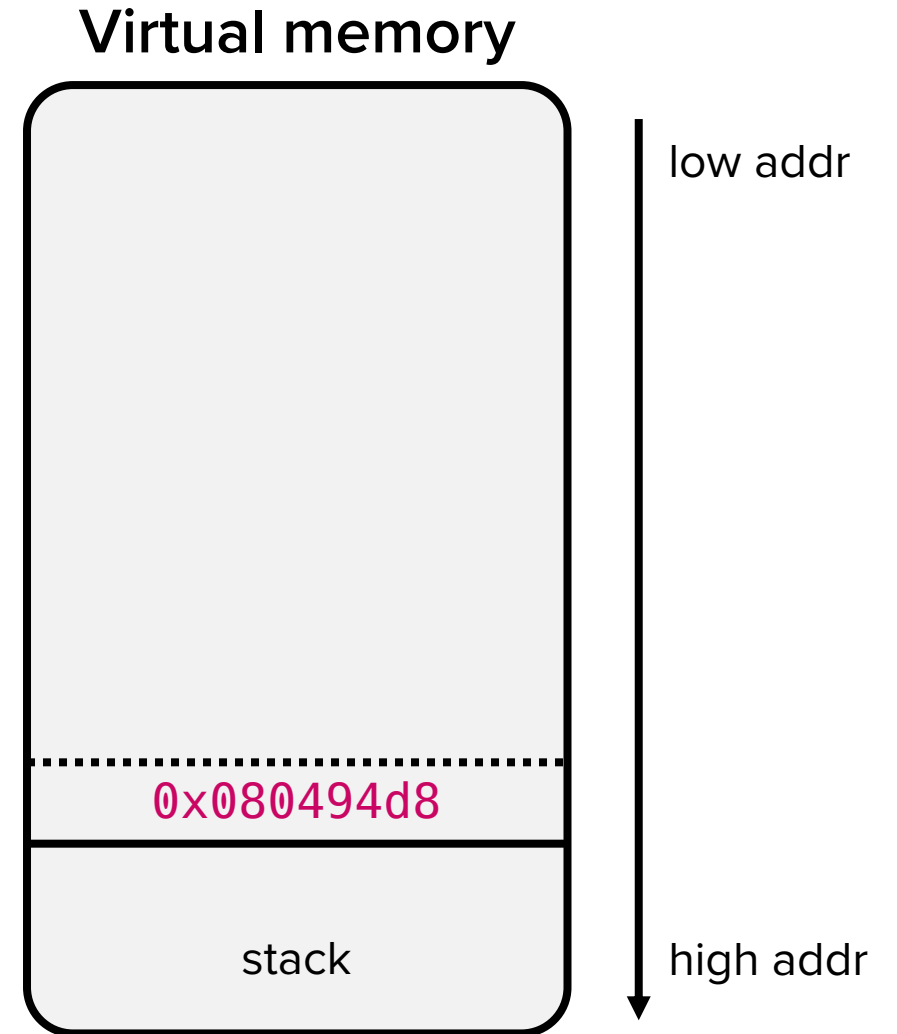
- **ret**
  - **pop** stack top into eip
  - e.g., in main() of lab01's target

```
0x080494d3: call 0x8049385 <phase_one>
0x080494d8: call 0x8049352 <phase_two>
```

```
<phase_one>
0x08049385: push ebp
0x08049386: mov  ebp, esp
...
0x08049482: ret
```

eip =>

esp =>



# Subroutine instructions: ret

- **ret**

- **pop** stack top into eip
- e.g., in main() of lab01's target

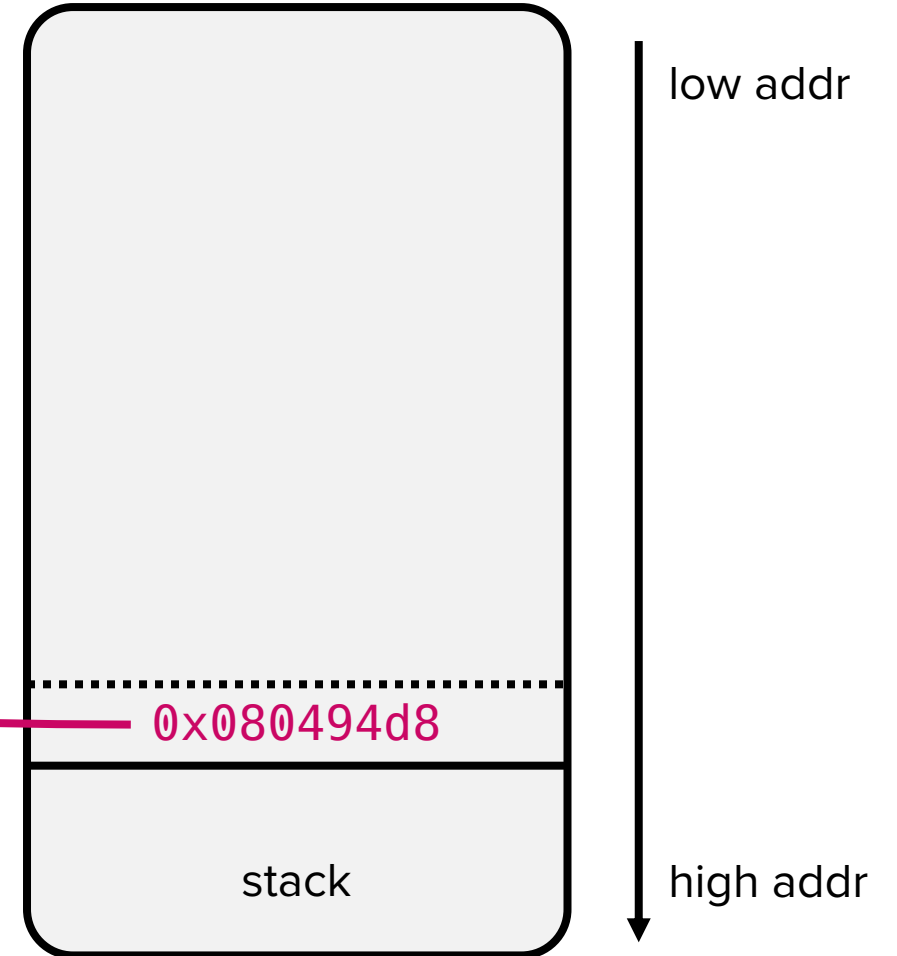
```
0x080494d3: call 0x08049385 <phase_one>  
0x080494d8: call 0x08049352 <phase_two>
```

eip =>

```
<phase_one>  
0x08049385: push ebp  
0x08049386: mov  ebp, esp  
...  
0x08049482: ret
```

esp =>

## Virtual memory



# Branches in assembly

- Implemented using `cmp` + conditional jump

```
cmp    eax, 0xdeadbeef
je     addr
```

```
if (eax == 0xdeadbeef) {
    goto addr;
}
```

```
cmp    eax, 0xdeadbeef
jne     addr
```

```
if (eax != 0xdeadbeef) {
    goto addr;
}
```

```
cmp    eax, 0xdeadbeef
jge     addr
```

```
if (eax >= 0xdeadbeef) {
    goto addr;
}
```

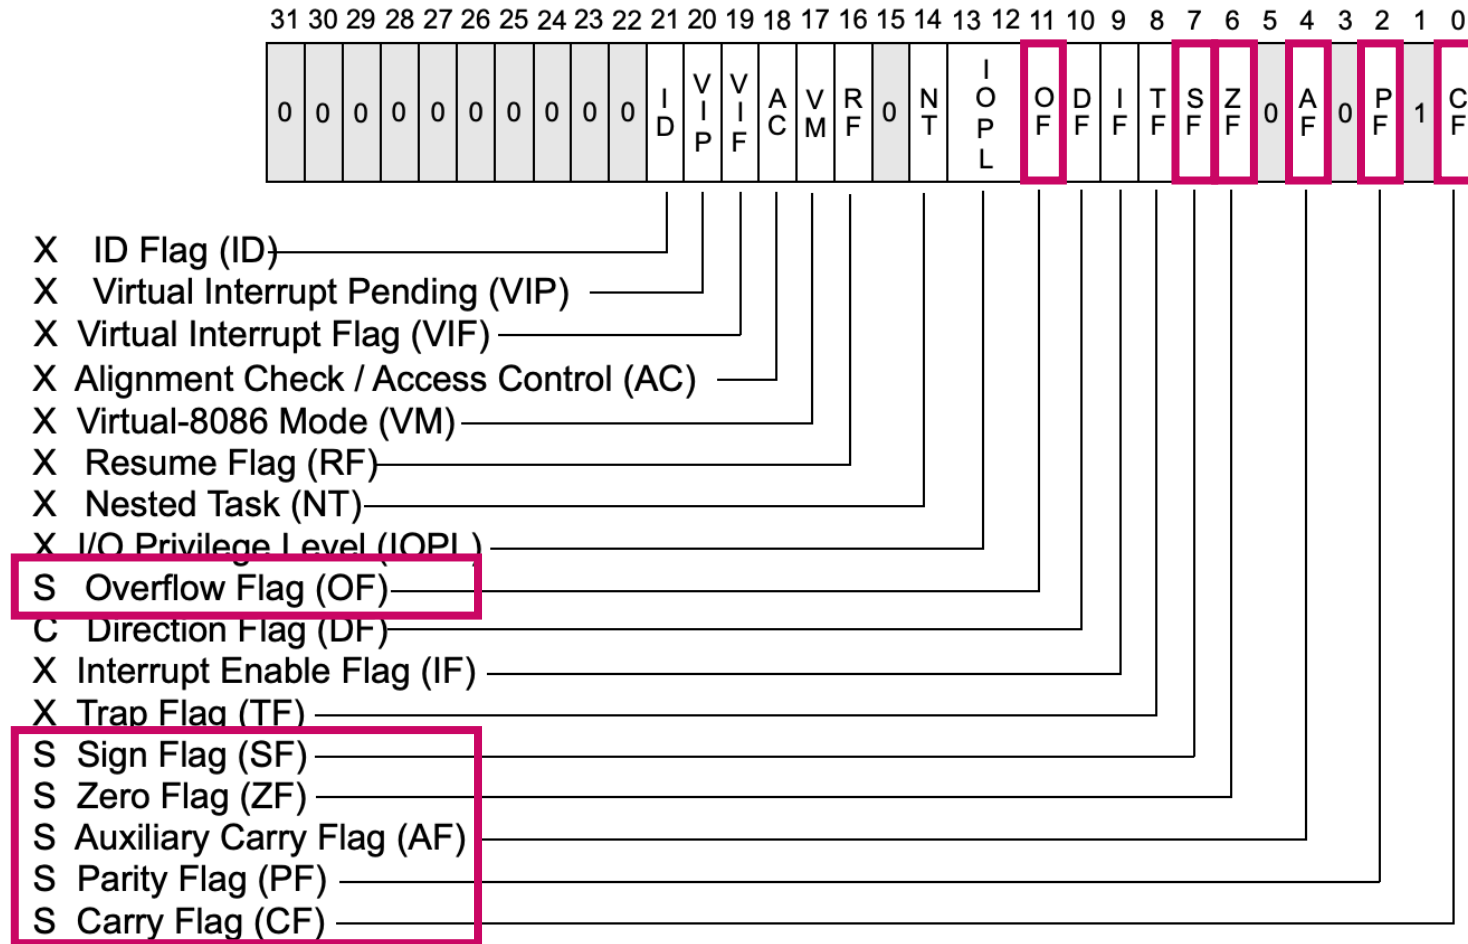
Q) Where does `cmp` store the result??

# x86 registers

---

- General purpose registers (GPR)
  - EAX, EBX, ECX, EDX
- Pointers
  - ESI, EDI
- Stack pointers
  - ESP, EBP
- Special purpose registers
  - EIP, EFLAGS ← This guy!

# EFLAGS register stores status / control / system flags



S Indicates a Status Flag  
C Indicates a Control Flag  
X Indicates a System Flag



# Conditional jumps and their conditions

Instruction Mnemonic	Condition (Flag States)	Description
<b>Unsigned Conditional Jumps</b>		
JA/JNBE	(CF or ZF) = 0	Above/not below or equal
JAЕ/JNB	CF = 0	Above or equal/not below
JB/JNAE	CF = 1	Below/not above or equal
JBE/JNA	(CF or ZF) = 1	Below or equal/not above
JC	CF = 1	Carry
JE/JZ	ZF = 1	Equal/zero
JNC	CF = 0	Not carry
JNE/JNZ	ZF = 0	Not equal/not zero
JNP/JPO	PF = 0	Not parity/parity odd
JP/JPE	PF = 1	Parity/parity even
JCXZ	CX = 0	Register CX is zero
JECXZ	ECX = 0	Register ECX is zero
<b>Signed Conditional Jumps</b>		
JG/JNLE	((SF xor OF) or ZF) = 0	Greater/not less or equal
JGE/JNL	(SF xor OF) = 0	Greater or equal/not less
JL/JNGE	(SF xor OF) = 1	Less/not greater or equal
JLE/JNG	((SF xor OF) or ZF) = 1	Less or equal/not greater
JNO	OF = 0	Not overflow
JNS	SF = 0	Not sign (non-negative)
JO	OF = 1	Overflow
JS	SF = 1	Sign (negative)

# Summary: With assembly, we can

- Move data around
  - Load/store data from/to data and registers
  - Compute pointer addresses
  - Use stack as a buffer (push, pop)
- Perform arithmetic operations
  - add, sub, and, xor, ...
- Call and return from functions
  - Push return address on call, pop that on ret
- Control execution flow
  - cmp followed by conditional jump family

**Turing complete!**

Given enough time and memory, can solve any computational problem

# Exercise: reading lab01's target

---

- Focus:
  - mov operation
  - push and pop
  - Function call and return
  - cmp and jumps

# Summary

---

- Do not get intimidated by assembly code!
  - They ARE human-readable

# Coming up next

---

- Writing malicious assembly code
- Exploiting buffer overflow to alter program's execution flow

# Questions?