

Lec 19: Access Control

CSED415: Computer Security
Spring 2024

Seulbae Kim



Administrivia


- Lab 04 is due this weekend!
 - Questions?

Recap

- User authentication
 - Enforces coarse-grained control for the entire system
 - Makes a binary decision: Grant or deny access

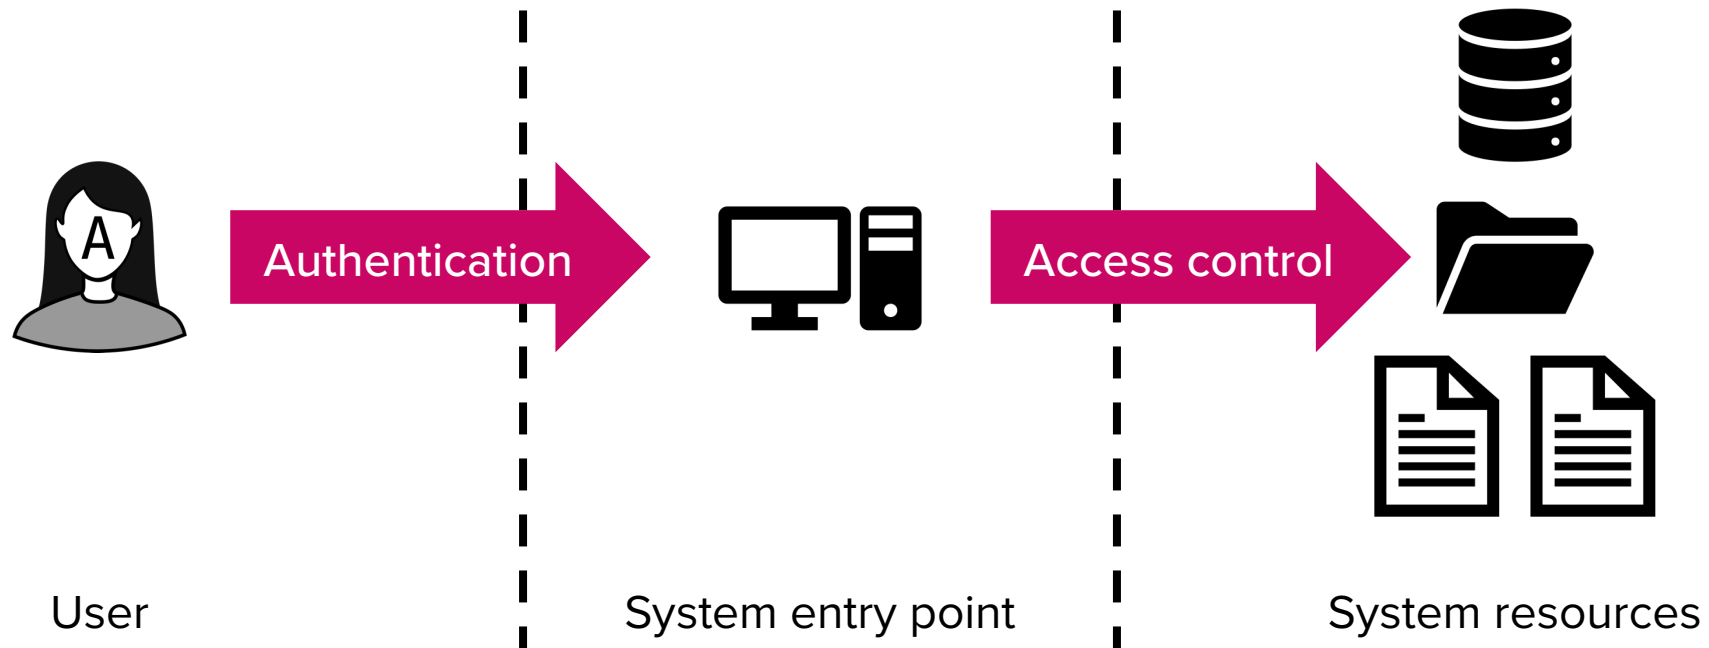
A nightclub example

- Authentication
 - ID check at the gate
 - Access control
 - Over 18 – only allowed in
 - Over 21 – allowed to buy and drink alcohol
 - On artist list – allowed to enter backstage and perform on stage
 - On VIP list – allowed to access VIP area
- Focus: What you are **allowed** to do



Access control

- Definition: Process of granting or denying specific requests to obtain and use information or resources



Models for access control

- Entities
 - **Subject:** An entity capable of accessing objects
 - Owner: Creator of a resource
 - Ownership is exclusive; a resource cannot be co-owned
 - Group: Named group of users who can exercise access rights
 - World (or others): Users who are not the owner nor group for a resource
 - **Object:** A resource to which access is controlled
 - Files, processes, memory, ...

Models for access control

- Access rights: Describe how a subject may access an object
 - File access rights:
 - Read: View information in a file
 - Write: Add, modify, or delete data in a file (includes read access)
 - Execute: Execute specified file
 - Directory access rights:
 - Delete: Delete files in a directory
 - Create: Create new file in a directory
 - Search: List files in a directory

Models for access control

- Two types of policy: DAC and MAC
 - DAC: Discretionary Access Control
 - Controls access based on the identity of the requestor and associated rules
 - Owner of the file determines access rights (hence “discretionary”)
 - MAC: Mandatory Access Control
 - Controls access based on security labels and clearances
 - System determines access rights (hence “mandatory”)

Discretionary Access Control (DAC)

Elementary forms of access control

- Authentication == Access control
 - Only available for single-subject, single-object environment
 - e.g., a safe
 - Allow access to authenticated subjects

Elementary forms of access control

- Blacklists and whitelists
 - Only available for multiple-subject, single-object environments
 - e.g., Email spam filter uses a blacklist
 - Blacklist: Allow by default, deny blacklisted subjects
 - Hard to reason about who can access resource
 - Whitelist: Deny by default, allow whitelisted subjects
 - Hard to deal with adding whitelist entries
 - Both lists can grow quite large

→ How to extend for modern systems with multiple subjects and objects?

Access control matrix (ACM)

- Allow multi-subject multi-object access control
- $\text{access}(\text{subject}, \text{object}) = 1 \text{ or } 0$
 - 1 (true): access granted
 - 0 (false): access denied

		Objects			
		A	B	C	D
Subjects	Alice	1	0	0	1
	Bob	0	1	1	1
	Claire	1	0	0	0
	Dave	0	1	1	1


Access control matrix (ACM)

- Finer-grained control is available with access rights
 - None, Own, Read, Write, Execute
- Problems
 - ACM is a “sparse matrix” by nature
 - High storage overhead
 - Size of ACM grows significantly as the number of subjects and objects increases

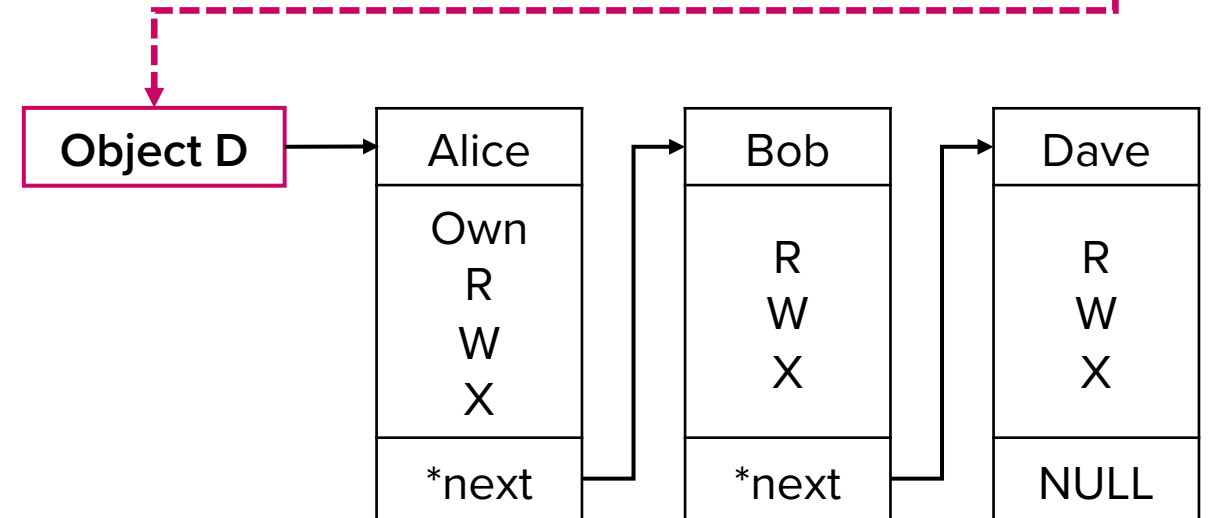
		Objects			
		A	B	C	D
Subjects	Alice	R	-	-	ORWX
	Bob	-	RW	ORW	RWX
	Claire	ORW	-	-	-
	Dave	-	ORW	R	RWX

Access control lists (ACL)

- Slice ACM by columns (objects)



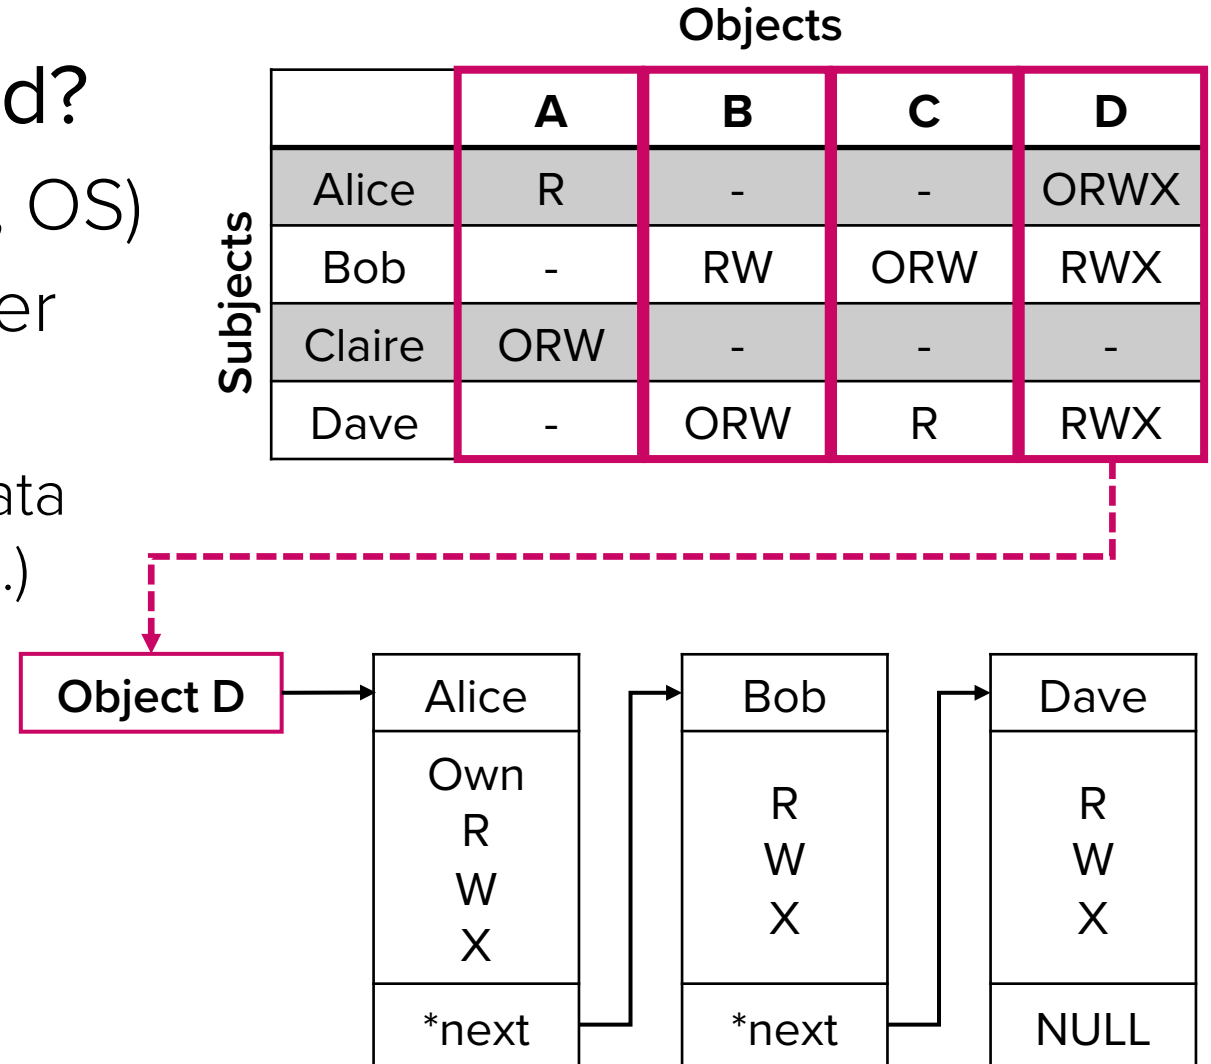
		Objects			
		A	B	C	D
Subjects	Alice	R	-	-	ORWX
	Bob	-	RW	ORW	RWX
	Claire	ORW	-	-	-
	Dave	-	ORW	R	RWX



- + Convenient to determine “who can access this resource?”
- Very inefficient to determine the objects that a specific subject can access


Access control lists (ACL)

- Where should an ACL be stored?
 - In trusted part of the system (e.g., OS)
 - Storing the ACL with object's other metadata would be natural
 - A file object has associated metadata (size, created time, modified time, ...)
 - ACL can also be stored as a metadata

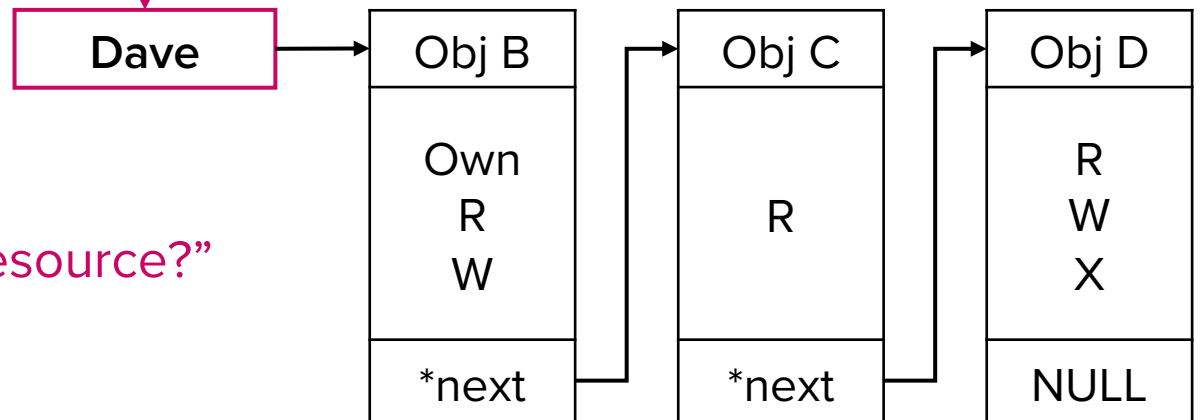


Capability lists (C-list)

- Slice ACM by rows (subjects)



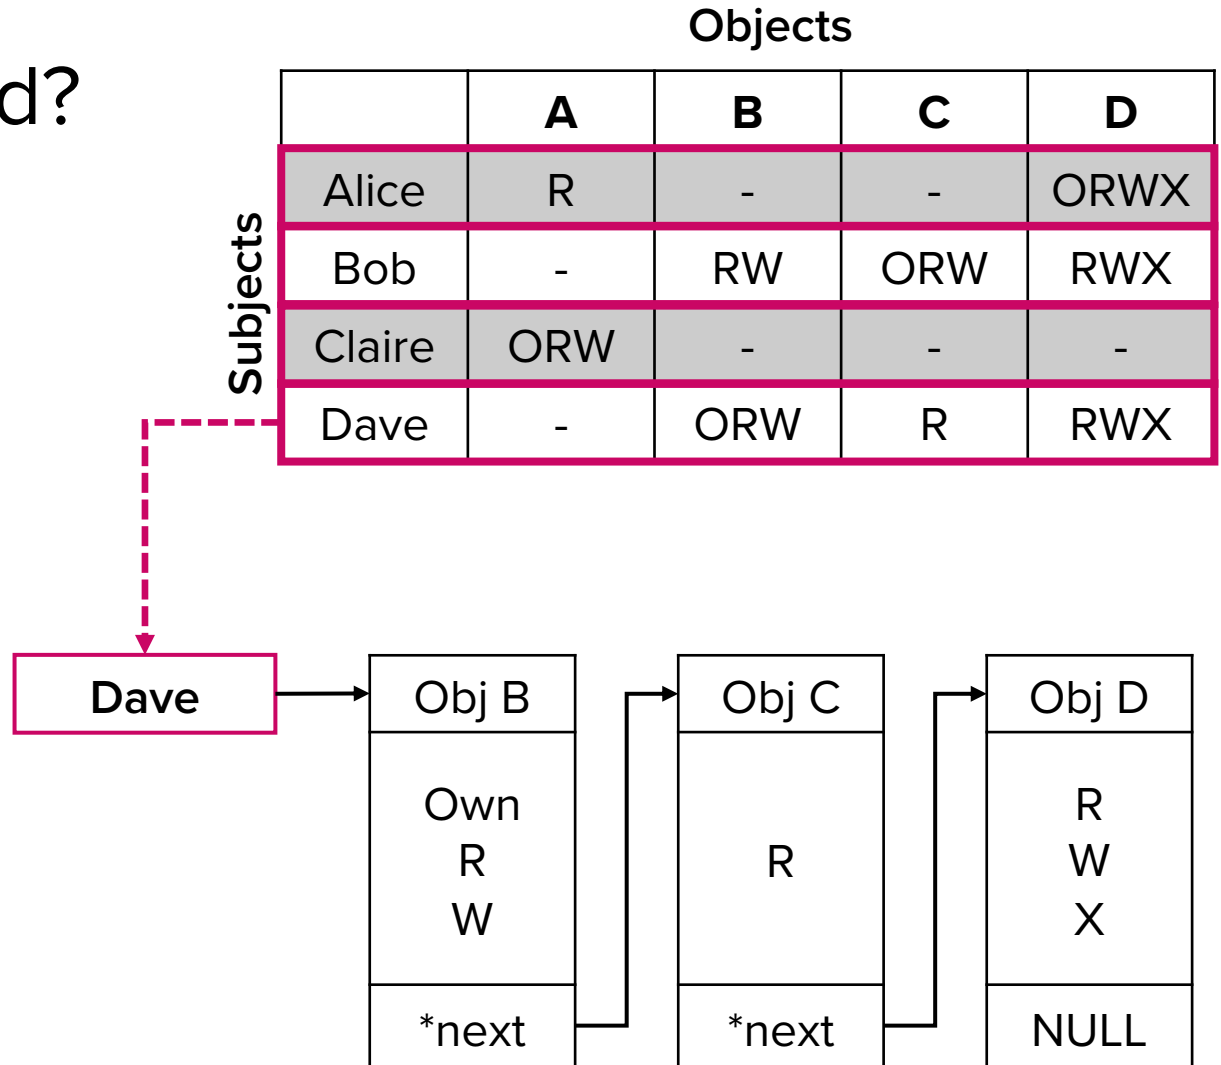
		Objects			
		A	B	C	D
Subjects	Alice	R	-	-	ORWX
	Bob	-	RW	ORW	RWX
	Claire	ORW	-	-	-
	Dave	-	ORW	R	RWX



- + Good for checking what a user can do
- + Provides flexibility for delegation
- Challenging to determine “who can access this resource?”
- Revocation is tricky

Capability lists (C-list)

- Where should a C-list be stored?
 - A “capability” is an unforgeable reference/handle for a resource
 - OS should maintain C-lists of all subjects (users)
 - Object sharing requires propagation of capabilities

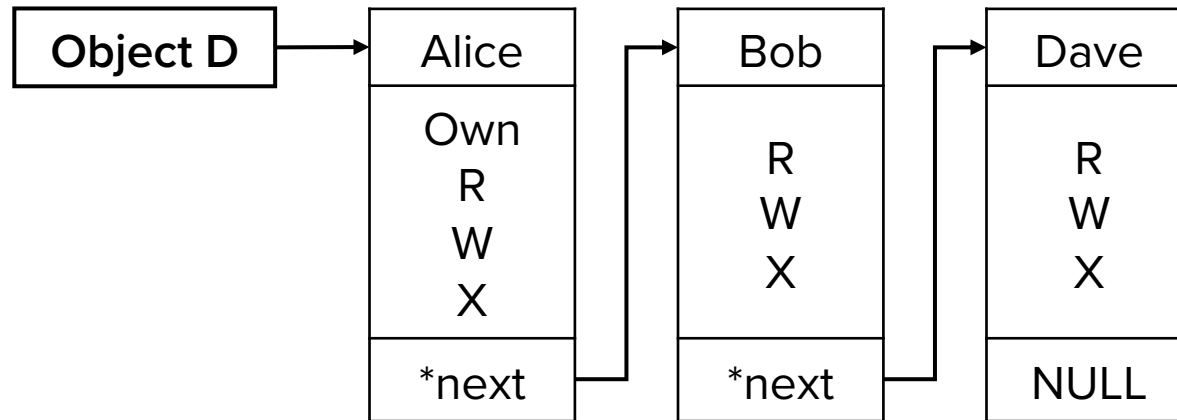


ACL vs C-list

- Which one is better?

1. Checking efficiency (e.g., subject tries to access an object)

- ACL – System needs to traverse object's ACL to find the user ($O(n)$)



If Eve requests to read object D, time complexity is $O(n)$ where n is the length of object D's ACL

- C-list – Subject can present its capability ticket for the object ($O(1)$)

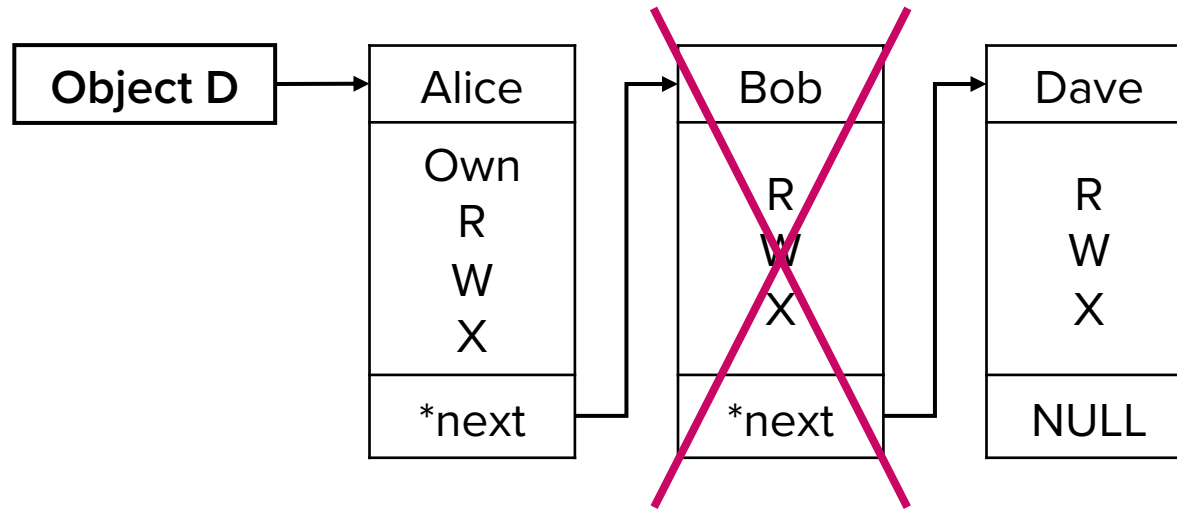
→ C-list wins

ACL vs C-list

- Which one is better?

2. Revocation (removing a subject's access to an object)

- ACL – Alice (owner) can remove Bob's permission from the ACL of object D



- C-list – Alice (owner) cannot control Bob's C-list. System needs to intervene
→ ACL wins

ACL vs C-list

- Which one is better?

3. Accountability (e.g., a sensitive file has been accessed and you want to find potential subject)

- ACL – All information is available in one place, i.e., the ACL of the file
- C-list – All subjects need to be investigated for their capabilities

→ ACL wins

Overall, ACL seems to outperform C-list

But are there any problems that ACL cannot handle?

Confused deputy problem

- Setting
 - Compiler is a pay-per-use service
`$ compiler input_filename output_filename`
 - System wants to charge users when they use the Compiler
 - The compiler updates a Billing file after it is executed
 - Alice wants to use the Compiler

Subjects	Objects	
	Compiler	Billing
	Alice	RX
Compiler	RX	RW

Access Control Matrix

Confused deputy problem

- Malicious behavior of Alice
 - Alice executes the Compiler several times to compile programs
 - Compiler updates Billing file with Alice's records
 - Then, Alice executes:
`$ compile input_filename Billing`
 - Billing file gets corrupted
 - Compiler, a deputy acting on behalf of Alice, is confused!

		Objects	
		Compiler	Billing
Subjects	Alice	RX	-
	Compiler	RX	RW

Access Control Matrix

Confused deputy problem

- What's the issue with ACL?
 - $\text{access}(\text{Alice}, \text{Compiler}, \text{execute}) = 1$
 - $\text{access}(\text{Compiler}, \text{Billing}, \text{write}) = 1$



		Objects	
		Compiler	Billing
Subjects	Alice	RX	-
	Compiler	RX	RW

Access Control Matrix

Confused deputy problem

- C-list can solve this problem through delegation
 - Alice does not have a capability to write to Billing
 - Alice **must delegate** her C-list to the Compiler when executing it
 - The Compiler cannot write to Billing
- Free from the confused deputy problem

		Objects	
		Compiler	Billing
Subjects	Alice	RX	-
	Compiler (on behalf of: Alice)	RX	-

Downside: The system should implement an additional monitor to write to Billing

Access Control Matrix

DAC in Practice

Unix-like systems use ACL

- Background
 - In Unix, every access-controlled resource is represented as a file
 - Memory
 - Device drivers
 - Named pipes
 - ...

Unix-like systems use ACL

- Background

- Each file has an owner (UID), group (GID), and everyone (world)
 - User is someone capable of using files
 - Group is a list of user accounts
 - One user may belong to many groups
 - User's details are in `/etc/passwd`
 - `lab02:x:1012:1012::/home/lab02:/bin/bash`
username uid gid home directory login command
 - Group details are in `/etc/group`
 - `lab02:x:1012:target,lab02,target02`
group name gid member list

Unix-like systems use ACL

- File permissions
 - Available file permissions are read (r), write (w), and execute (x)
 - Original ACL implementation had a 9-bit representation
 - 3 bits for the owner, 3 bits for the group, 3 bits for everyone else
 - e.g., **rwXrw-r--** : Owner can rwx, group members can rw, and others can r

Unix-like systems use ACL

- Directory permissions work differently
 - Read: List contents of directory
 - Write: Create or delete files in directory
 - Execute: Use anything in or change working directory to directory

```
mkdir /tmp/xxx  
cd /tmp/xxx  
mkdir kkk  
stat kkk | grep Access  
cd kkk  
cd ..  
chmod a-x kkk  
cd kkk
```

→ temporary directory for testing

→ shows (0775/**d**rwxrwxr-x)

→ can cd (change directory) to kkk

→ remove x permission from all (user, group, others)

→ access denied

Unix-like systems use ACL

- ACL implementation
 - Original ACL implementation had a 9-bit representation
 - 3 bits for the owner, 3 bits for the group, 3 bits for everyone else
 - e.g., **rw-rw-r--** : Owner can rw, group members can rw, and others can r
 - Modern OSes support full ACL (Linux, BSD, MacOS, ...)
 - 3 additional bits: setuid, setgid, and sticky bit

Unix-like systems use ACL

- ACL implementation

- Numeric representation of permission bits consists of four digits

- User, group, others permissions:

Bit position	2	1	0
Permission	Read	Write	Execute

r: $2^2 = 4$ w: $2^1 = 2$ x: $2^0 = 1$

→ rwx: read + write + execute = $4 + 2 + 1 = 7$

→ rw: read + write = $4 + 2 = 6$


- Special permissions:

Bit position	2	1	0
Permission	setuid	setgid	sticky bit

→ setuid: $2^2 = 4$

Represent full permission with 4 digits:

special owner group others



Q) what does 4750 mean?

Q) what does 3000 mean?

Unix-like systems use ACL

- Extended permissions
 - Sticky bit
 - Originally used to lock file in memory (sticky!)
 - Now used on directories to limit delete operation
 - If sticky bit is set, must own file or directory to delete
 - Other users cannot delete even with write permission
 - Example

```
cd /tmp/xxx  
mkdir mmm  
chmod +t mmm  
stat mmm | grep Access
```

→ temporary directory for testing

→ shows (1775/drwxrwxr-**t**)

Q) Where can you find this permission?

Unix-like systems use ACL

- Extended permissions
 - SetUID: If set, program runs as the owner no matter who executes it
 - SetGID: If set, program runs as the member of the group
 - “Runs as” == Runs with the same privileges as
- Examples:

- Lab target binaries

```
$ stat /home/lab02/target | grep Access
Access: (4750/-rwsr-x---)  Uid: ( 1999/  target)   Gid: ( 1012/  lab02)
```

- sudo

```
$ stat /usr/bin/sudo | grep Access
Access: (4755/-rwsr-xr-x)  Uid: (    0/   root)   Gid: (    0/   root)
```

Unix-like systems use ACL

- Extended permissions
 - SetUID/SetGID binaries are great targets for attackers seeking privilege escalation
 - e.g., CVE-2012-0809 in sudo binary
 - sudo had a debug mode (-D), which invokes sudo_debug function
 - The function printf's the program name (i.e., sudo)
 - The program name can be controlled by an attacker
 - e.g., by creating a symbolic link: `ln -s /usr/bin/sudo arbitrary_name`
 - “arbitrary_name” can be a format string payload, which leads to arbitrary code execution (format string vulnerabilities are not covered in this course)
 - Huge impact - The attacker can become root and execute any command as root

Unix-like systems use ACL

- Interacting with files in Unix-like systems through syscalls
 - Creating
 - `creat(filename, mode);`
 - `open(filename, flags, mode);` // specify `O_CREAT` in flags to create file
 - Opening
 - `int fd = open(filename, flags);`
 - flags: `O_RDONLY`, `O_WRONLY`, or `O_RDWR`
 - OS returns a file descriptor (`fd`) if the file exists
 - ACL check happens at this stage!
 - System traverses the file's ACL and checks whether the permission in the open flags match the subject's access rights
 - Afterwards, the file can be accessed through the file descriptor (`fd`)

Unix-like systems use ACL

- Interacting with files in Unix-like systems through syscalls
 - open's error check example:

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdio.h>
#include <errno.h>
#include <string.h>

int main(void) {
    int fd = open("myfile", O_RDWR);
    printf("fd: %d\n", fd);

    char* buf = strerror(errno);
    printf("Error: %s\n", buf);

    return 0;
}
```

→ Test with varying permissions of myfile

Unix-like systems use ACL

- Interacting with files in Unix-like systems through syscalls
 - Reading
 - `read(fd, buf, count);`
 - read count bytes and store in buf from the open file referred to by the `fd`
 - Writing
 - `write(fd, buf, count);`
 - write count bytes from buf to the open file referred to by the `fd`
 - Closing
 - `close(fd);`
 - Closes a file descriptor (invalidates the reference)

Reading and writing does not involve any permission check → performance!

Attacking Access Control

TOCTOU vulnerability

- Time-of-check-time-of-use (TOCTOU)
 - Access right checking is performed when a file is opened
 - Once checked, the permission remains available until the file is explicitly closed (or the process terminates and implicitly closed)
 - What if file permission is changed during this time?

TOCTOU vulnerability

- Time-of-check-time-of-use (TOCTOU)
 - Example: `vi` (text editor) - predecessor of `vim` (VI iMproved)
 - `vi` keeps a backup of the original file upon save
 - Save: **rename** the original file as a backup and **create** a new file with the original name
 - If we run `vi` as root, modify Alice's file and save,
 - `rename("alice_file", "alice_file.bak");` retains the permissions
 - `open("alice_file", 0_CREAT);` → this is owned by root (since `vi` is running as root)
 - `vi` needs to change the owner to Alice by invoking **chown** syscall

TOCTOU vulnerability

- Time-of-check-time-of-use (TOCTOU)
 - Example: `vi` (text editor) - predecessor of `vim` (VI iMproved)

`vi`

```
rename("alice_file", "alice_file.bak");
```

```
open("alice_file", O_CREAT);
```

TOC

```
chown("alice_file", alice_uid, alice_gid);
```

TOU

Q) How can we prevent such attacks?

Attacker

```
$ ln -s /etc/passwd alice_file
```

Result: `/etc/passwd` is owned by Alice

TOCTOU vulnerability

- Time-of-check-time-of-use (TOCTOU)
 - Example: `vi` (text editor) - predecessor of `vim` (VI iMproved)

`vi`

```
rename("alice_file", "alice_file.bak");
```

```
open("alice_file", O_CREAT);
```

TOC

```
chown("alice_file", alice_uid, alice_gid);
```

TOU

Prevention:

1. Make window as small as possible (although not failproof)
2. Use unpredictable path for `alice_file`
3. Create `alice_file` in a secure directory

Attacker

```
$ ln -s /etc/passwd alice_file
```

Result: `/etc/passwd` is owned by Alice

Summary

- Access control allows us to determine if a request from a subject to access an object can be granted
- ACLs and C-lists are two ways to represent states used for discretionary access control decisions
- Unix-like systems use ACL for access controls

Coming up next

- Information flow control problem
 - With DAC, Alice can never be sure that sensitive data she shares with Bob will not be further shared with others
 - Motivation for Mandatory Access Control (MAC)

Questions?