

# Lec 08: Attacks and Defenses (2)

CSED415: Computer Security  
Spring 2025

Seulbae Kim



# Announcement

---

- No class meeting next Tuesday (March 18)
  - I will be traveling to present a funding proposal
  - The lecture video will be uploaded before Tuesday
    - We will begin a new section (Cryptography)
    - **Important:** Please watch the video before next Thursday's lecture to ensure you can follow the subsequent material

# Recap

- NX is effective at preventing return-to-stack attacks
  - MMU raises page fault exception if fetched instruction belongs to a NX page
  - Stack is marked NX by default
- Return-to-libc attack bypasses NX protection
  - Return to a libc function useful for exploitations
- Return-Oriented Programming (ROP) generalizes code reuse attacks

# Recap

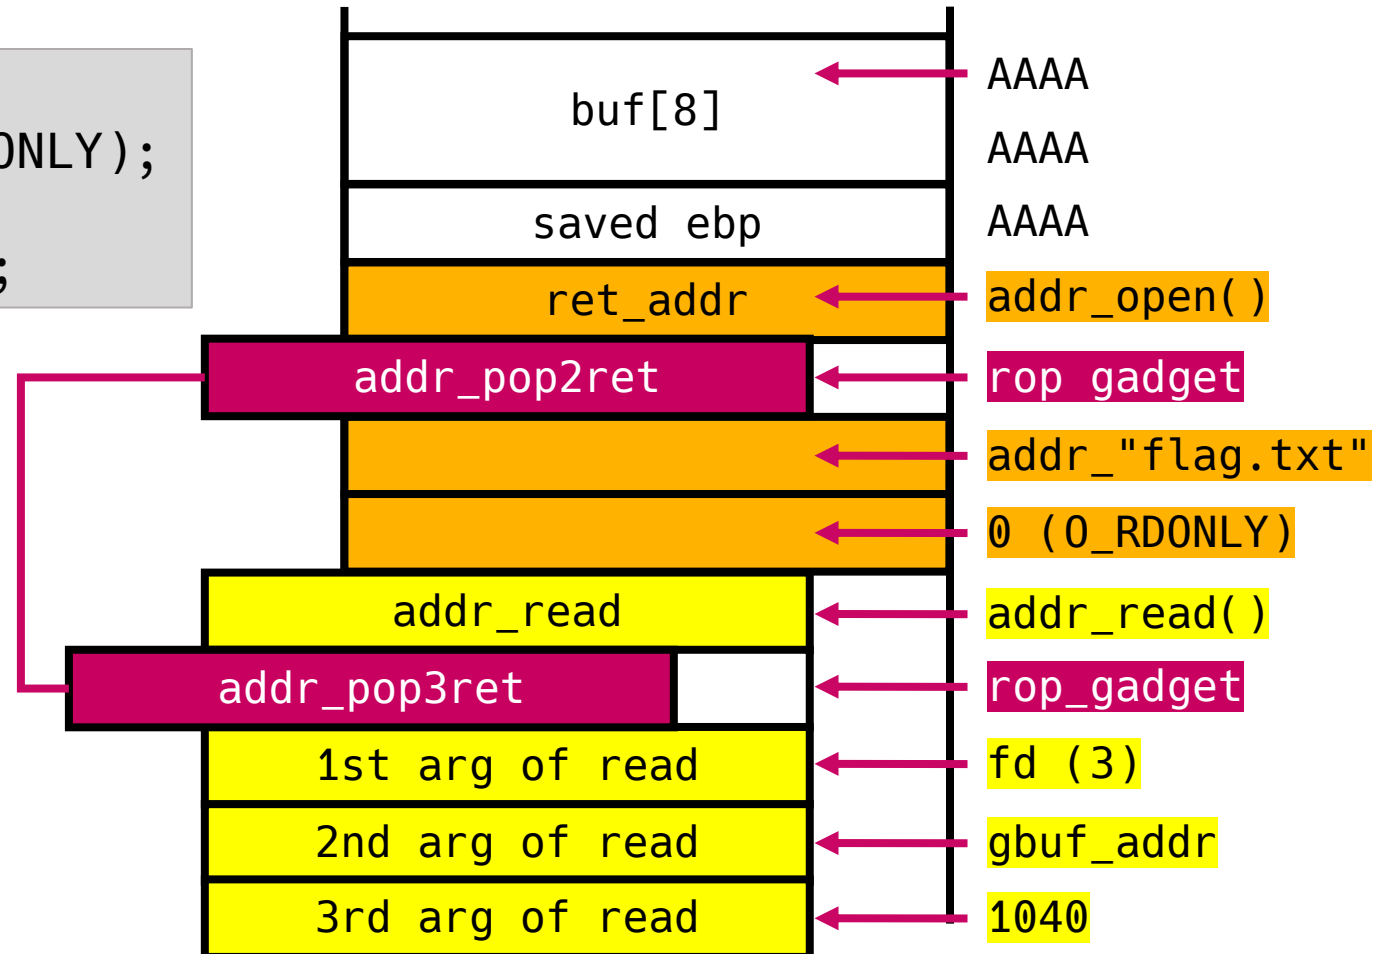
- Chaining multiple func calls

[Goal]

```
1. int fd = open("flag.txt", O_RDONLY);  
2. read(fd, gbuf_addr, 1040);  
3. write(stdout, gbuf_addr, 1040);
```

ROP gadgets are essential for chaining  
returns to multiple functions

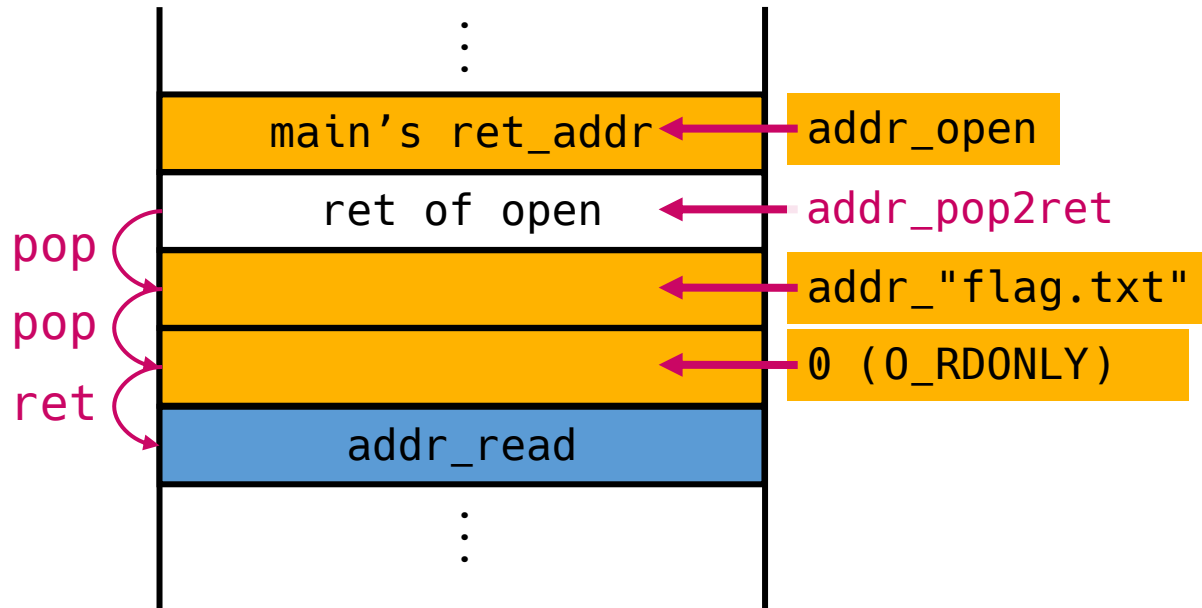
- ROP chain (x86, 32-bit)




# Attack #1-2: ROP (x86\_64)

# ROP for x86\_64 ?

- x86 stores args on the stack
  - Any pop2ret gadget is useful



- x64 puts args in the registers
  - RDI, RSI, RDX, RCX, R8, R9
  - Only the extra args are stored on the stack



Before returning to `addr_open`

`addr_"flag.txt"` should be put in RDI

`0 (0_RDONLY)` should be put in RSI

How?

# ROP for x64

- We should utilize gadgets that POP stack values into registers
  - e.g., `pop {rdi, rsi, rdx, rcx, r8, r9}; ret;`
  - Example gadgets:

```
0x401303: pop rdi
0x401304: ret
```

```
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
```

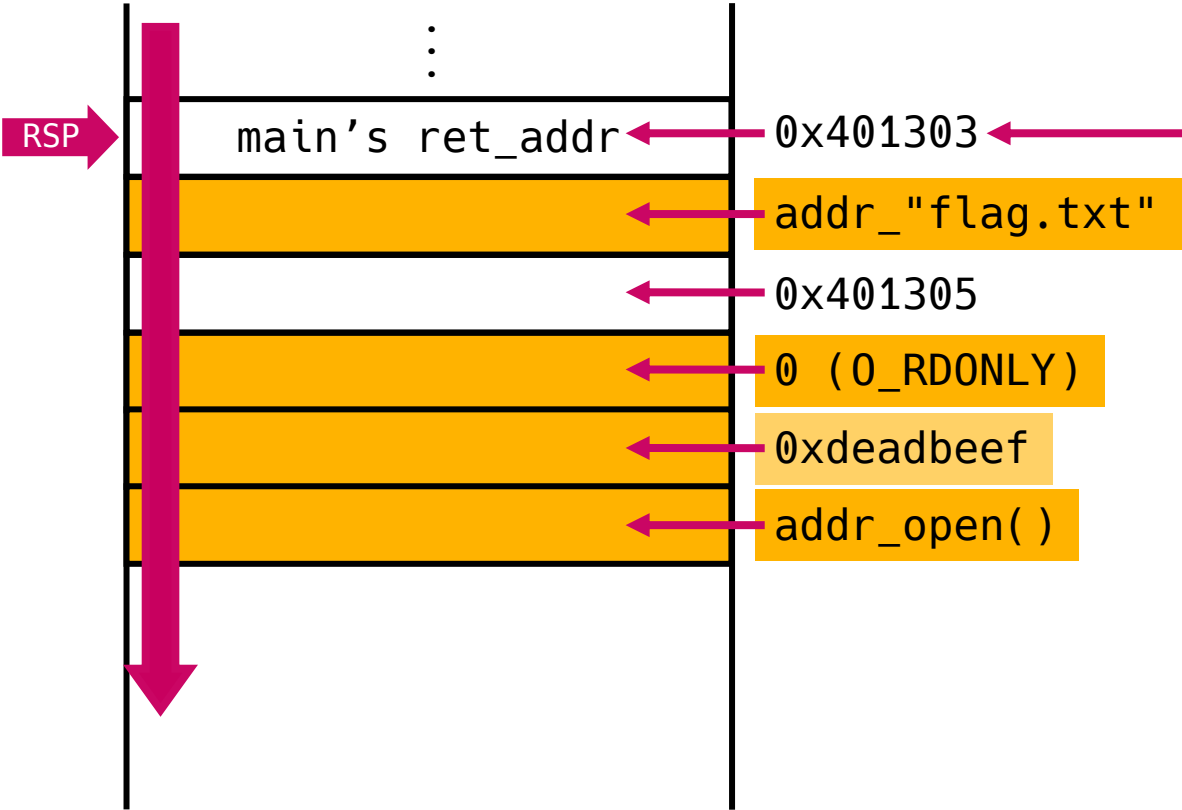
```
0x401309: pop rdx
0x40130a: ret
```

# ROP for x64

- Setting up args using gadgets

(Just before main( ) returns)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
0x401303: pop rdi
0x401304: ret
```

```
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
```

```
0x401309: pop rdx
0x40130a: ret
```

REG	VALUE
RDI	?
RSI	?
RDX	?

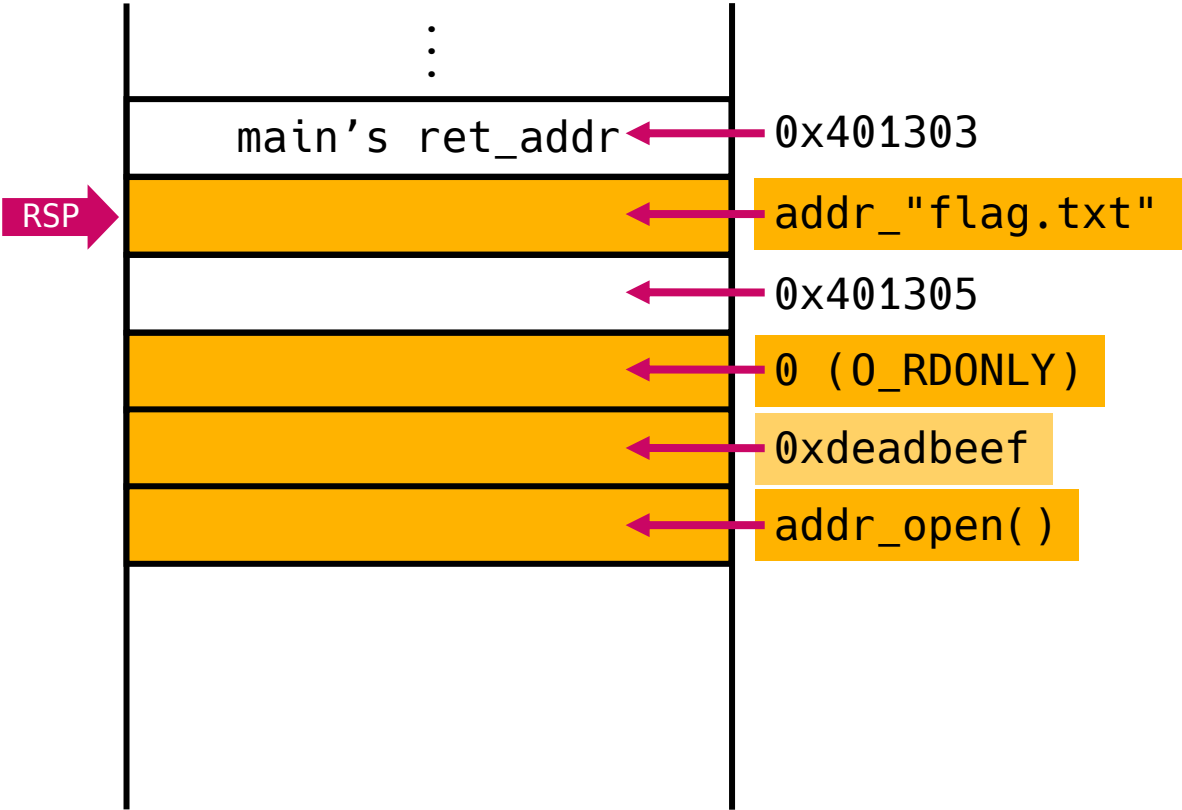


# ROP for x64

- Setting up args using gadgets

(Return to the “pop rdi” gadget)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
RIP → 0x401303: pop rdi
        0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

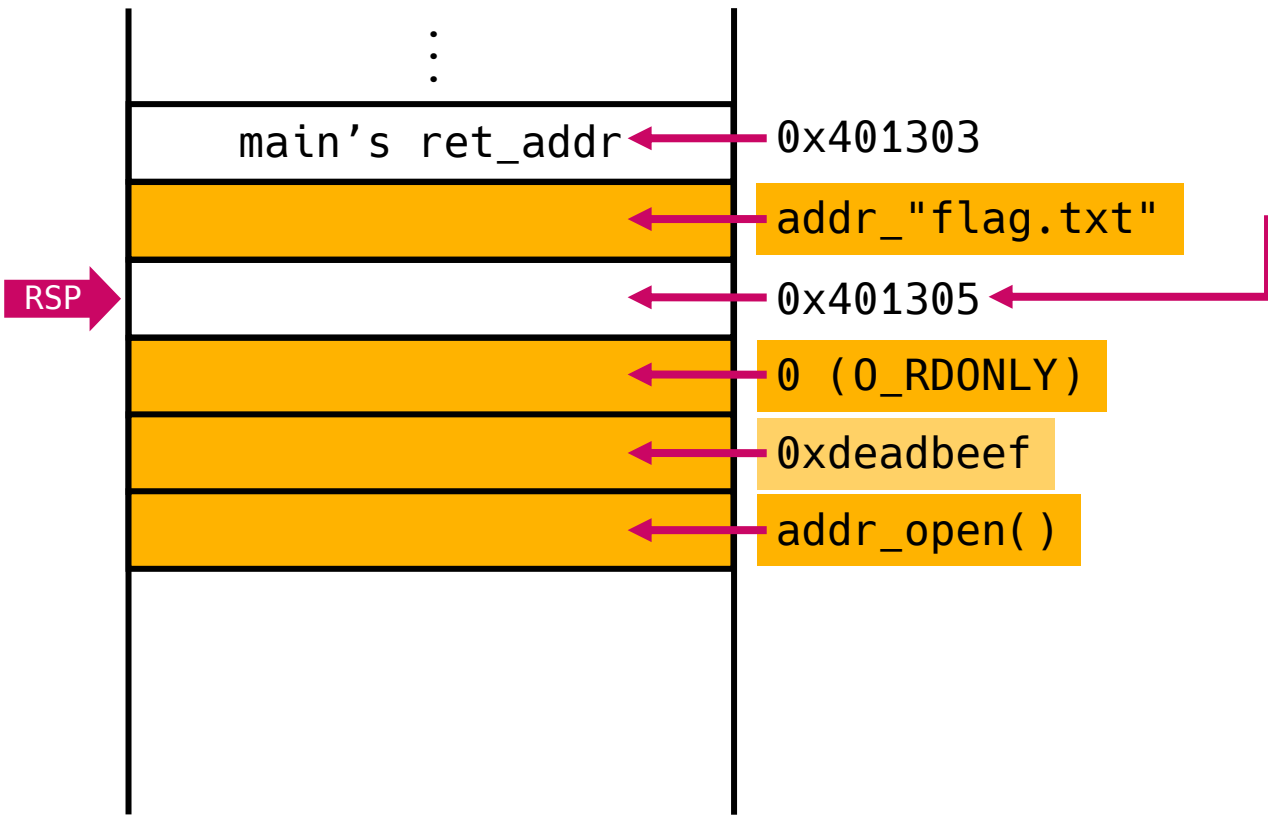
REG	VALUE
RDI	?
RSI	?
RDX	?

# ROP for x64

- Setting up args using gadgets

(addr\_"flag.txt" is popped into rdi)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



RIP

```
0x401303: pop rdi
0x401304: ret
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

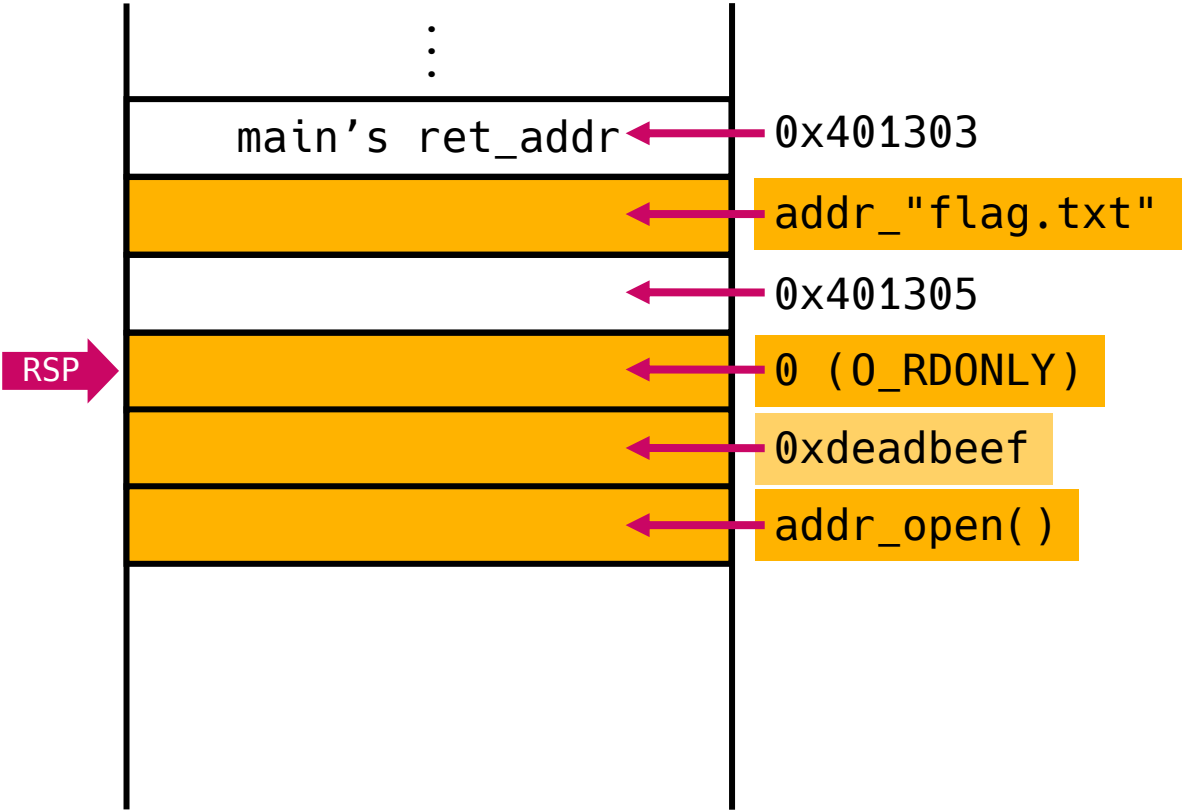
REG	VALUE
RDI	addr_"flag.txt"
RSI	?
RDX	?

# ROP for x64

- Setting up args using gadgets

(return to the “pop rsi” gadget)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



RIP →

```
0x401303: pop rdi
0x401304: ret
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

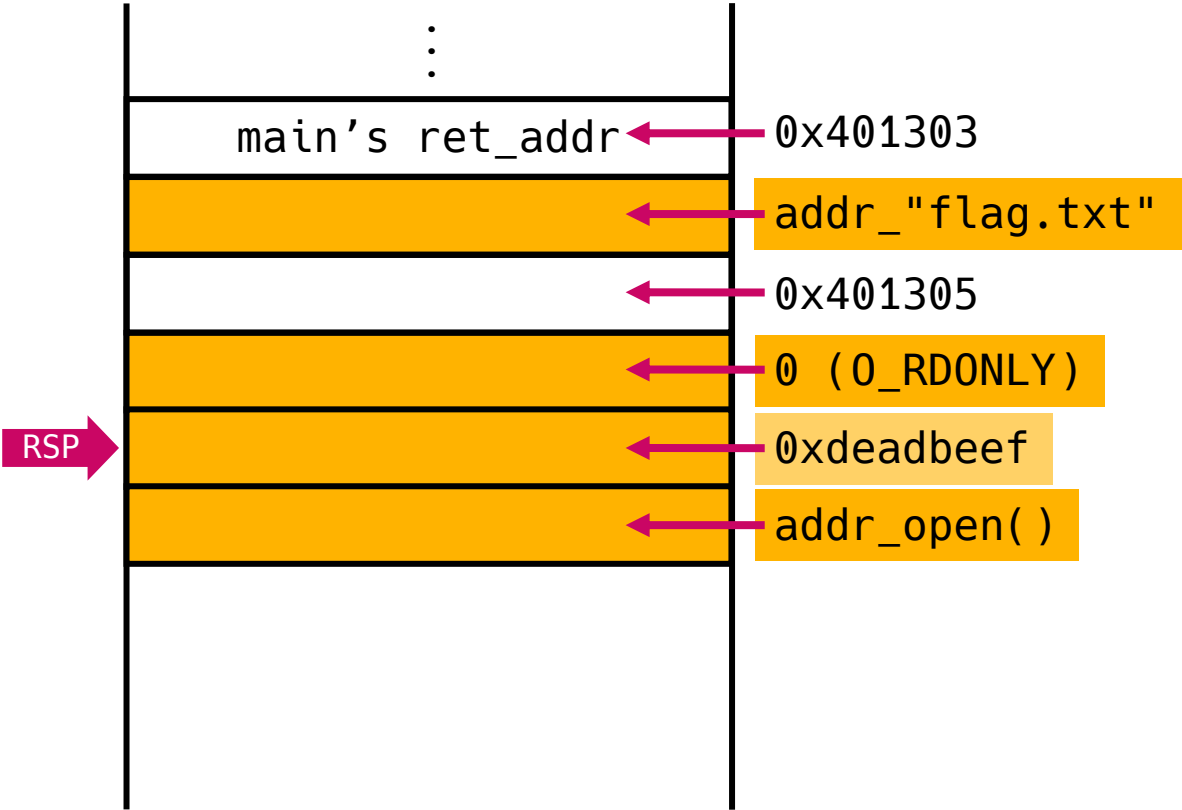
REG	VALUE
RDI	addr_"flag.txt"
RSI	?
RDX	?

# ROP for x64

- Setting up args using gadgets

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

(0 is popped into RSI)



RIP

```
0x401303: pop rdi
0x401304: ret
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

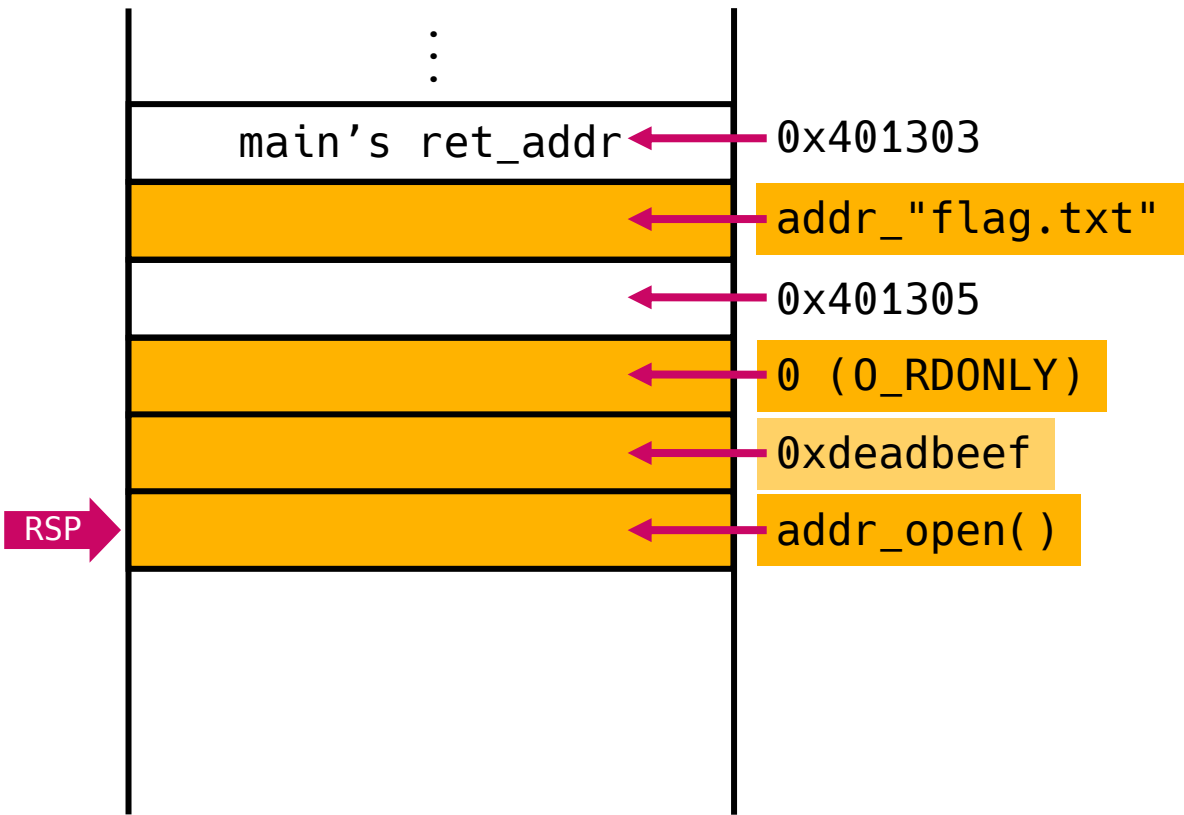
REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

# ROP for x64

- Setting up args using gadgets

Pop a dummy into r15 (irrelevant)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
0x401303: pop rdi
0x401304: ret
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

RIP →

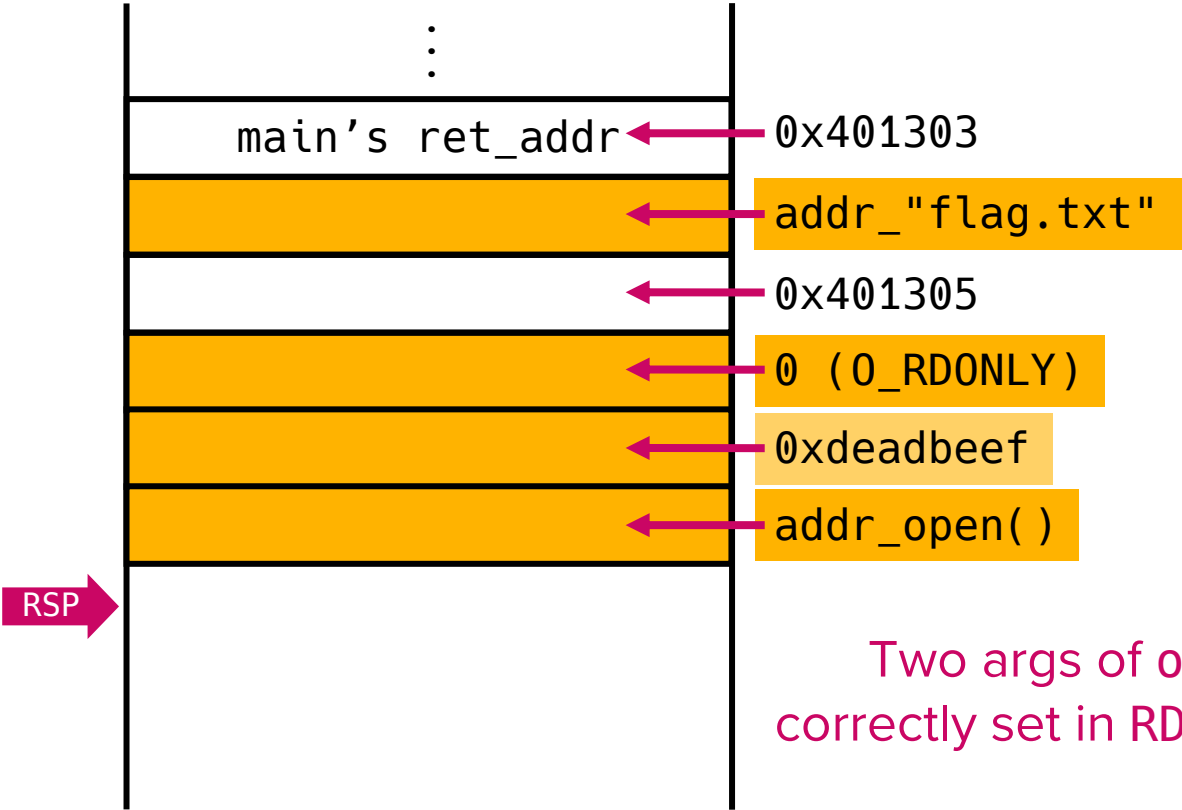
REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

# ROP for x64

- Setting up args using gadgets

Returns to `open(addr_"flag.txt", 0);`

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



RIP →

```
0x401303: pop rdi
0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

Two args of `open()` are correctly set in RDI and RSI

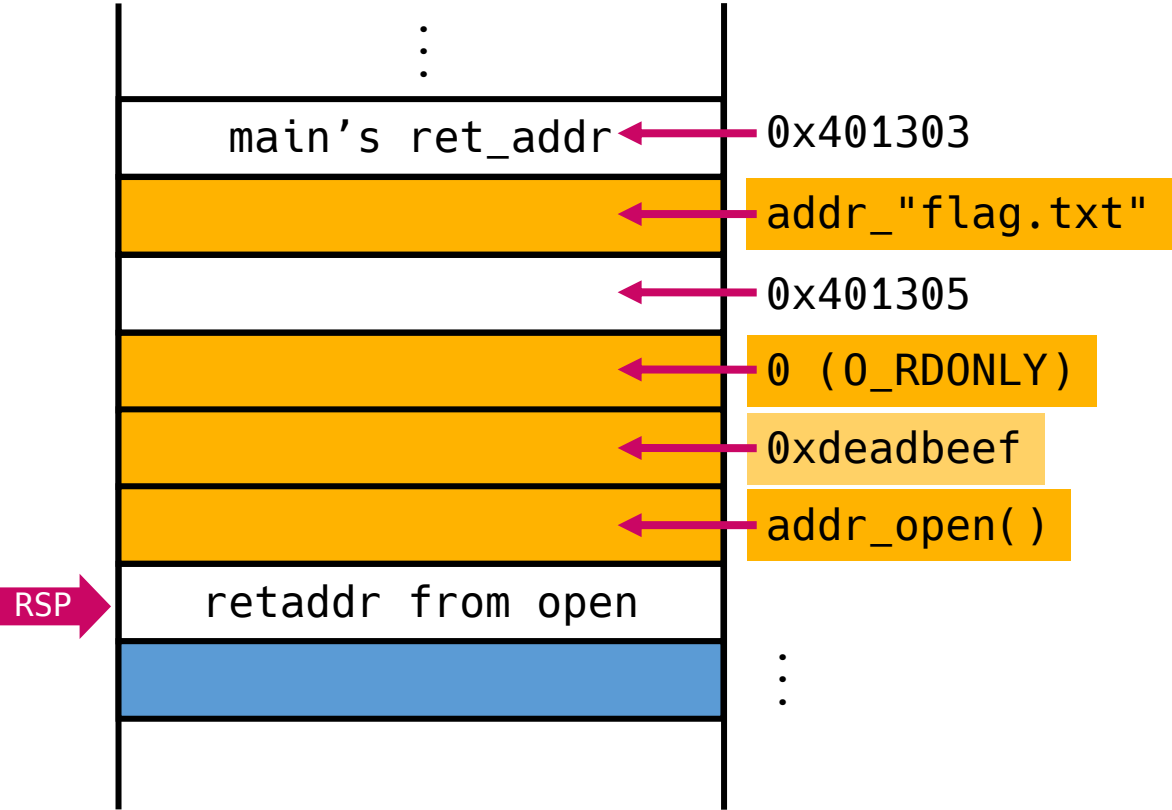
REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

# ROP for x64

- Setting up args using gadgets

We can keep the chain going!

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
0x401303: pop rdi
0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

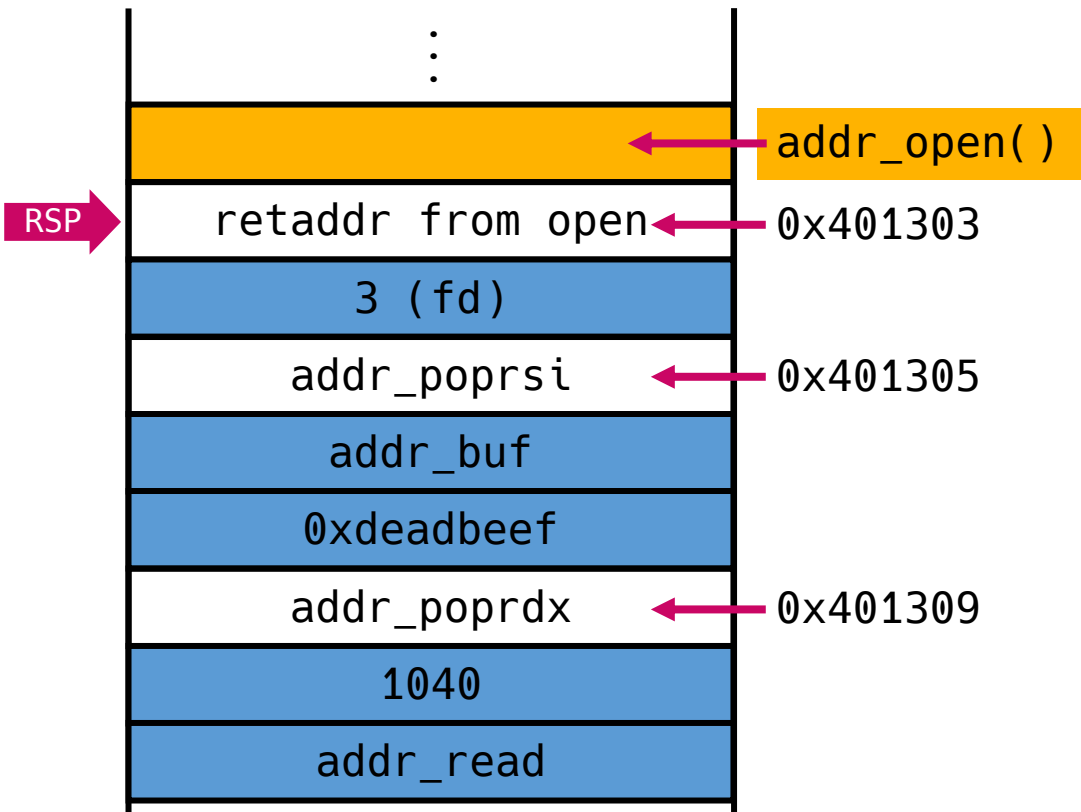
0x401309: pop rdx
0x40130a: ret
```

REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

# ROP for x64

- Setting up args using gadgets

(Scrolled down)



```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

```
0x401303: pop rdi
0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

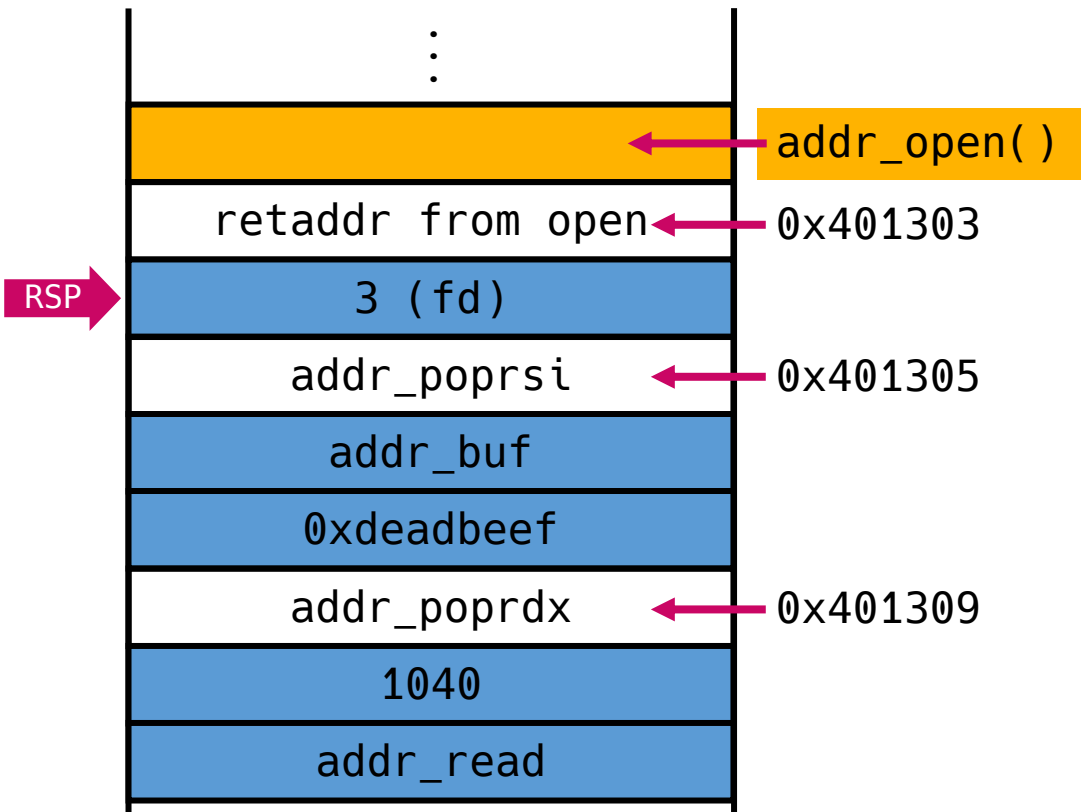


# ROP for x64

- Setting up args using gadgets

(return to the “pop rdi” gadget for arg1)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
RIP → 0x401303: pop rdi
        0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

REG	VALUE
RDI	addr_"flag.txt"
RSI	0
RDX	?

**POSTECH**

- ```
1. int fd = open("/proc/flag", O_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

Stack diagram showing memory layout:

- addr\_open( ) (yellow box)
- retaddr from open (white box, points to 0x401303)
- 3 (fd) (blue box)
- addr\_poprsi (white box, points to 0x401305)
- addr\_buf (blue box)
- 0xdeadbeef (blue box)
- addr\_poprdx (white box, points to 0x401309)
- 1040 (blue box)
- addr\_read (blue box)

A pink arrow labeled RSP points to the stack.

```
0x401303: pop rdi
0x401304: ret
```

```
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
```

```
0x401309: pop rdx
0x40130a: ret
```

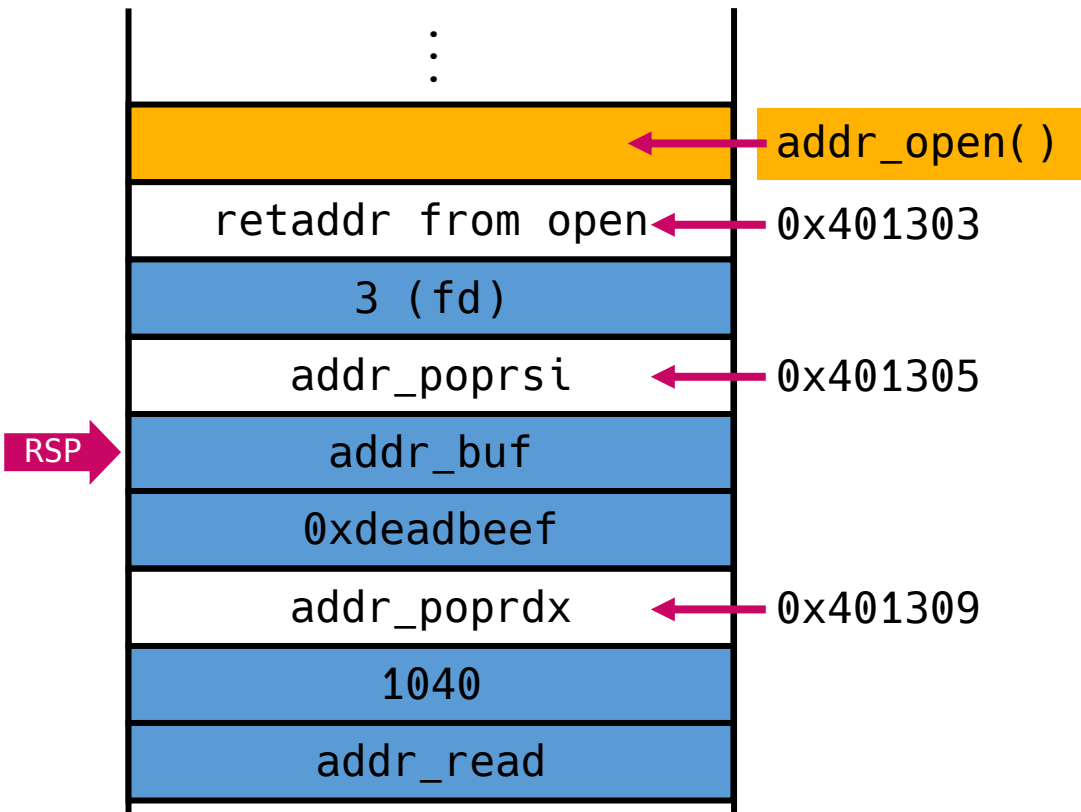
| REG | VALUE |
|-----|-------|
| RDI | 3     |
| RSI | 0     |
| RDX | ?     |

# ROP for x64

- Setting up args using gadgets

(return to the “pop rsi” gadget for arg2)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



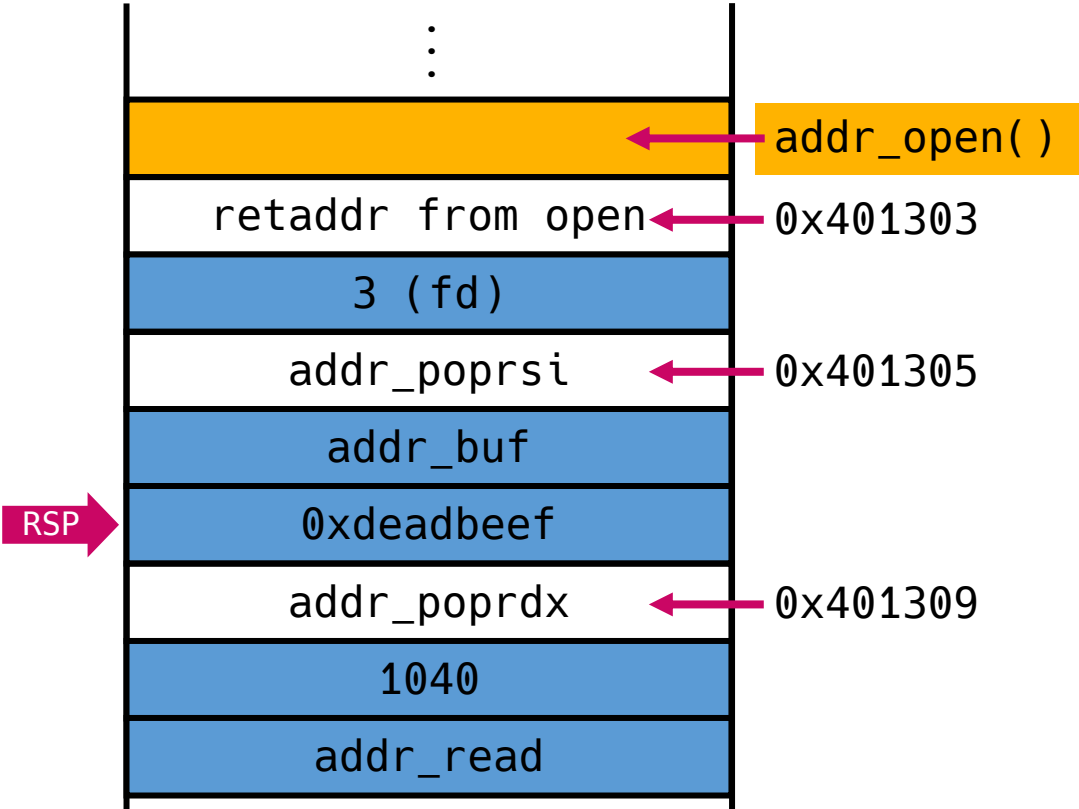
```
0x401303: pop rdi
0x401304: ret
RIP -> 0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

| REG | VALUE |
|-----|-------|
| RDI | 3     |
| RSI | 0     |
| RDX | ?     |

# ROP for x64

- Setting up args using gadgets

(addr\_buf is popped into RSI)



```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

RIP →

|           |         |
|-----------|---------|
| 0x401303: | pop rdi |
| 0x401304: | ret     |
| 0x401305: | pop rsi |
| 0x401306: | pop r15 |
| 0x401307: | ret     |
| 0x401309: | pop rdx |
| 0x40130a: | ret     |

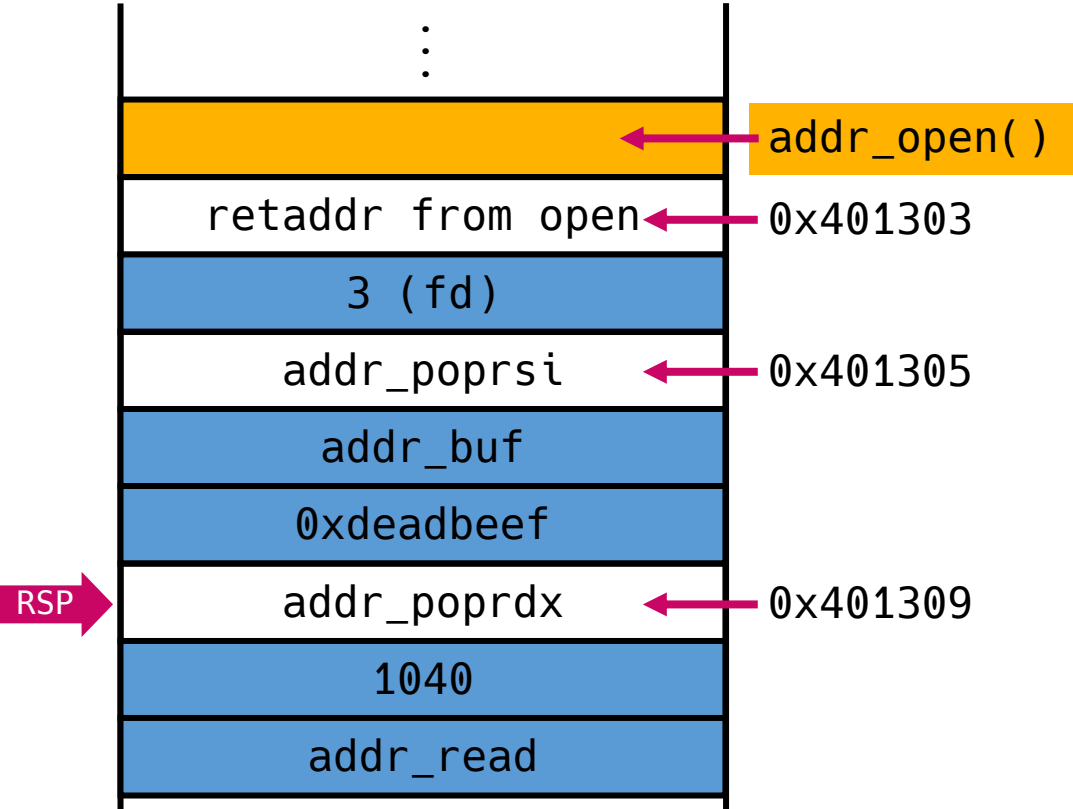
| REG | VALUE    |
|-----|----------|
| RDI | 3        |
| RSI | addr_buf |
| RDX | ?        |

# ROP for x64

- Setting up args using gadgets

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

(A dummy in r15)



0x401303: pop rdi  
0x401304: ret

0x401305: pop rsi  
0x401306: pop r15  
0x401307: ret

0x401309: pop rdx  
0x40130a: ret

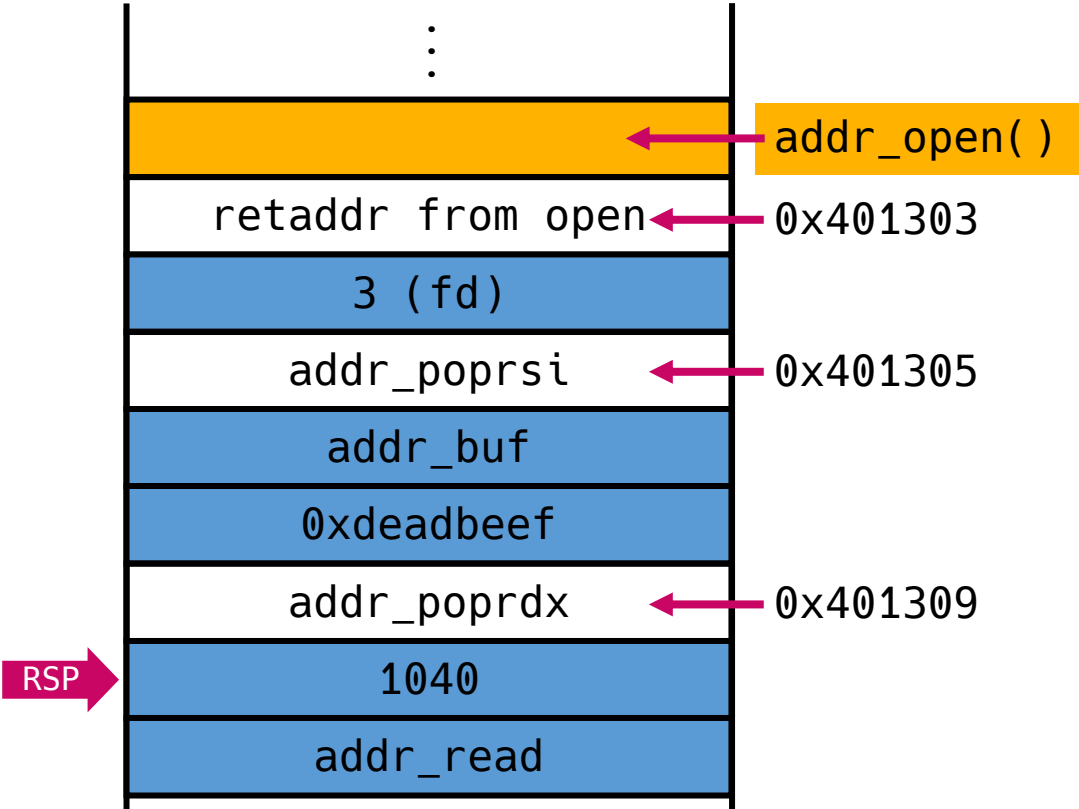
| REG | VALUE    |
|-----|----------|
| RDI | 3        |
| RSI | addr_buf |
| RDX | ?        |

# ROP for x64

- Setting up args using gadgets

(return to the “pop rdx” gadget for arg3)

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
0x401303: pop rdi
0x401304: ret
0x401305: pop rsi
0x401306: pop r15
0x401307: ret
0x401309: pop rdx
0x40130a: ret
```

RIP →

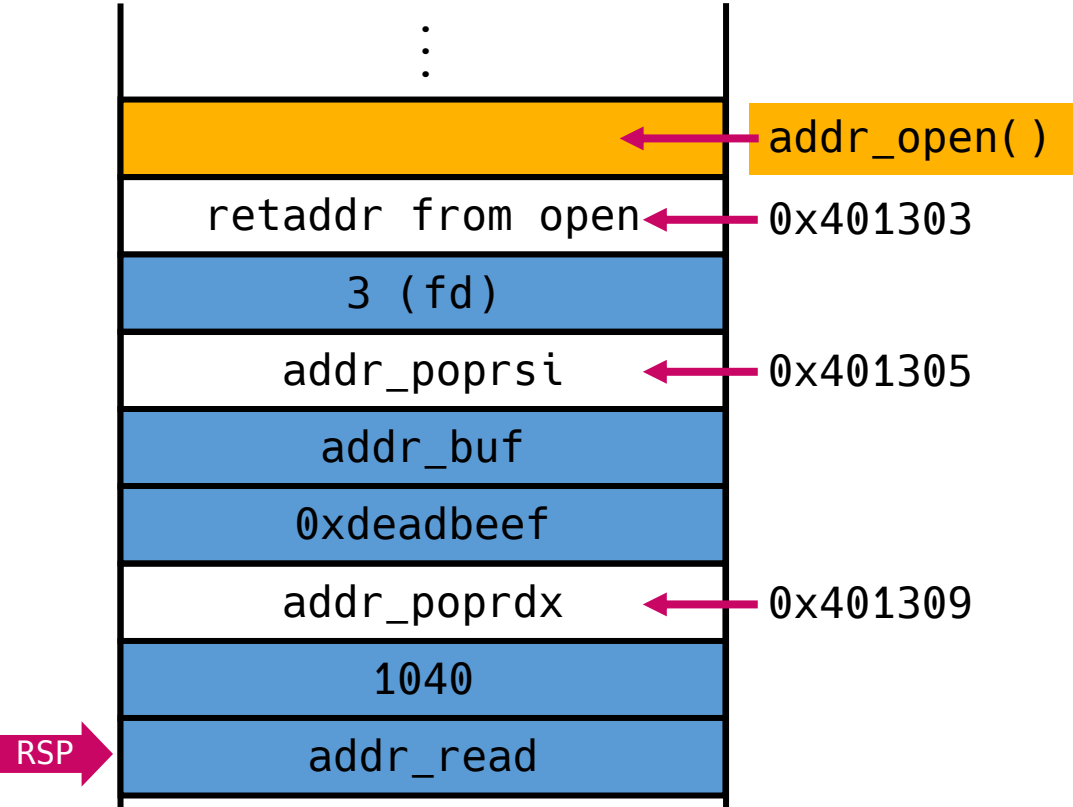
| REG | VALUE    |
|-----|----------|
| RDI | 3        |
| RSI | addr_buf |
| RDX | ?        |

# ROP for x64

- Setting up args using gadgets

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```

(1040 is popped into RDX)



```
0x401303: pop rdi
0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

RIP →

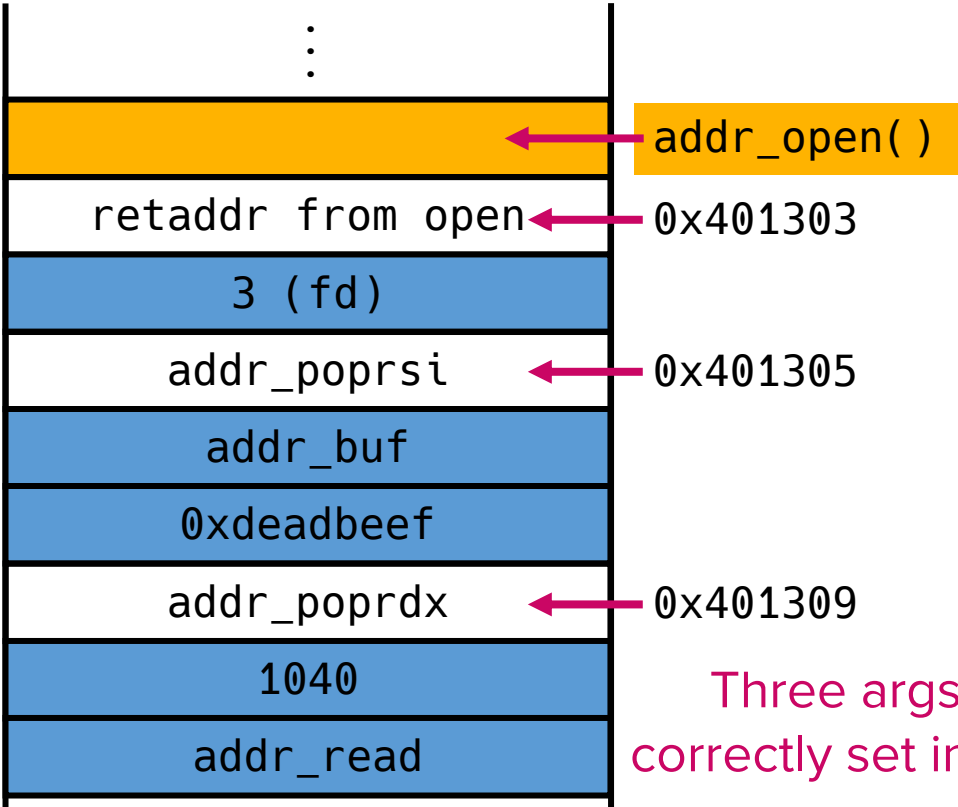
| REG | VALUE    |
|-----|----------|
| RDI | 3        |
| RSI | addr_buf |
| RDX | 1040     |

# ROP for x64

- Setting up args using gadgets

Returns to `read(3, addr_buf, 1040);`

```
1. int fd = open("/proc/flag", 0_RDONLY);
2. read(fd, addr_buf, 1040);
3. write(stdout, addr_buf, 1040);
```



```
0x401303: pop rdi
0x401304: ret

0x401305: pop rsi
0x401306: pop r15
0x401307: ret

0x401309: pop rdx
0x40130a: ret
```

| REG | VALUE    |
|-----|----------|
| RDI | 3        |
| RSI | addr_buf |
| RDX | 1040     |

Three args of `read( )` are correctly set in RDI, RSI, RDX



# More on ROP Gadgets

# ROP gadgets

- Small sequences of instructions found within the existing code of a program or the libraries it links to
- Key property
  - ROP gadget **ends** with an instruction that affects the instruction pointer (**rip** of x86\_64, **eip** of x86)
    - `ret` ; == `pop rip`
    - `call <label>` ; == `push next_rip + jmp <label>`
    - `jmp <label>` ; == `rip = <label>`

# ROP gadgets in Lab 02's target binary

- Finding gadgets through objdump
  - Search for ret instruction from target's disassembly

```
csed415-lab02@csed415:~$ objdump -D ./target -M intel | grep -B3 ret
...
--
1267:    e8 64 ff ff ff      call    11d0 <deregister_tm_clones>
126c:    c6 05 b5 2d 00 00 01  mov     BYTE PTR [rip+0x2db5],0x1
1273:    5d                  pop     rbp
1274:    c3                  ret
--
14b8:    48 83 c4 08          add     rsp,0x8
14bc:    c3                  ret
--
```

No useful gadget found ☹️

# ROP gadgets in Lab 02's target binary

- Finding gadgets through objdump
  - Search for ret instruction from target's disassembly

```
cse415-lab02@cse415:~$ objdump -D ./target -M intel | grep -B3 ret
```

Can we find more gadgets?

```
--  
14b8: 48 83 c4 08          add     rsp,0x8  
14bc: c3                  ret  
--
```

No useful gadget found ☹️

# Background: Variable-length instructions

- Intel's CISC architecture employs variable-length instructions
  - Length of encoded instructions vary

## Asm

## Binary

push ebp

→ 0x55

mov ebp, esp

→ 0x89 0xe5

add esp, 0x10

→ 0x83 0xc4 0x10

endbr32

→ 0xf3 0x0f 0x1e 0xfb

call (near call)

→ 0xe8 0xca 0xfd 0xff 0xff

...

...

## Q) Advantage?

- Smaller binary size  
(We can assign shorter lengths to frequently used instructions)

## Q) Disadvantage?

- Potential security issues  
(See next slides)

# Background: Variable-length instructions

```
csed415-lab02@csed415:~$ objdump -D ./target -M intel
```

```
...
1457:  f3 0f 1e fa      endbr64
145b:  55              push    rbp
145c:  48 89 e5        mov     rbp, rsp
145f:  48 83 ec 10     sub     rsp, 0x10
1463:  89 7d fc        mov     DWORD PTR [rbp-0x4], edi
1466:  48 89 75 f0     mov     QWORD PTR [rbp-0x10], rsi
146a:  e8 1a fe ff ff  call    1289 <init>
146f:  48 8d 05 2a 0c 00 00  lea     rax, [rip+0xc2a]
1476:  48 89 c7        mov     rdi, rax
1479:  e8 72 fc ff ff  call    10f0 <puts@plt>
...
Variable length!
```

# Disassembling x86\_64

- Bytecode

```
0x45 0x08 0x5d 0x48 0x01 0xc8 0xc3 ...
```



- Disassembled code from ↑

```
or      BYTE PTR [r13+0x48],r11b  
add     eax,ecx  
ret  
mov     rdi,rbp
```

# Disassembling x86\_64

- Bytecode

```
0x45 0x08 0x5d 0x48 0x01 0xc8 0xc3 ...
```



- Disassembled code from ↑

```
or      BYTE PTR [r13+0x48],bl  
add     eax,ecx  
ret  
mov     rdi,rbp
```

Completely different results, but instructions are still valid



# Disassembling x86\_64

- Bytecode

```
0x45 0x08 0x5d 0x48 0x01 0xc8 0xc3 ...
```



- Disassembled code from ↑

```
pop    rbp
add    rax,rcx
ret
mov    rdi,rbp
```

Completely different results, but instructions are still valid

# Disassembling x86\_64

- Bytecode

```
0x45 0x08 0x5d 0x48 0x01 0xc8 0xc3 ...
```



- Disassembled code from ↑

```
add    ecx,ecx  
ret  
mov    rdi,rbp  
mov    rdx,rax
```

Completely different results, but instructions are still valid

# Disassembling x86\_64

- Bytecode

```
0x45 0x08 0x5d 0x48 0x01 0xc8 0xc3 ...
```



- Disassembled code from ↑

```
enter 0x48c3,0x89  
out dx,eax  
mov rdx,rax  
pop rbp
```

Completely different results, but instructions are still valid

# Variable-length instructions

- Variable-length instructions can be disassembled (decoded) from any address
  - Q) Is it legal to jump to the middle of variable-length instructions?
    - Why not? It's perfectly legal in terms of program execution
    - Recall Lec 07: “The CPU is agnostic of the execution semantics”
  - Security problem:
    - We can find **unintended** ROP gadgets

# Variable-length instructions

- Aligned instructions (i.e., what objdump sees):

|                   |                |
|-------------------|----------------|
| e8 05 ff ff ff    | call 8048330   |
| 81 c3 59 12 00 00 | add ebx,0x1259 |

- If disassembled from the 2nd byte:

|                |                    |
|----------------|--------------------|
| 05 ff ff ff 81 | add eax,0x81ffffff |
| c3             | ret                |

Unintended ret instructions may appear

# ropper: A tool to find ROP gadgets

- ropper scans a binary for all ROP gadgets
  - Including unintended `ret`, `call`, and `jmp` instructions

```
csed415-lab02@csed415:~$ ropper -f ./target
...
0x000000000000001017: add esp, 8; ret;
0x000000000000001016: add rsp, 8; ret;
0x000000000000001014: call rax;
0x0000000000000011ef: jmp rax;
0x000000000000001273: pop rbp; ret;
...
```

(Search for ROP gadgets in the target's bytecode)

# ropper: A tool to find ROP gadgets

- ropper scans a binary for all ROP gadgets
  - Including unintended `ret`, `call`, and `jmp` instructions

```
csed415-lab02@csed415:~$ ropper -f ./target
```

```
...  
0x000000000000001017: add esp, 8; ret;  
0x000000000000001016: add rsp, 8; ret;  
0x000000000000001014: call rax;  
0x0000000000000011ef: jmp rax;  
0x000000000000001273: pop rbp; ret;  
...
```

But.. What if the gadgets required for x86\_64 rop attack do not exist in the binary?

e.g.,

```
pop rdi  
ret
```

```
pop rsi  
ret
```

```
pop rdx  
ret
```

# ropper: A tool to find ROP gadgets

- Finding more gadgets:
  - LIBC inevitably contains many gadgets as it is compiled from a huge code base

```
$ ldd ./target | grep libc (Find the LIBC's shared object file that target links to)
libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6
```

```
$ ropper -f /lib/x86_64-linux-gnu/libc.so.6 --search "pop rdi; ret"
0x0000000000002a3e5: pop rdi; ret; (Search for specific gadgets)
```

```
$ ropper -f /lib/x86_64-linux-gnu/libc.so.6 --search "pop rsi; ret"
0x000000000000163f88: pop rsi; ret; (Search for specific gadgets)
```

```
$ ropper -f /lib/x86_64-linux-gnu/libc.so.6 --search "pop % ret"
(a lot of pop*ret gadgets) (% matches any character)
```



# Defense #2: ASLR


# Address Space Layout Randomization

- World without ASLR
  - An executable is always loaded at the same virtual address region
  - Therefore, all addresses of code and data are invariant
    - They never change across multiple executions
    - All function addresses, data (e.g., string) addresses, and ROP gadget addresses can be easily identified
  - Attackers can easily launch ret-to-libc or ROP attacks

Mitigation idea: What if we “randomize” memory segments every time a program is loaded?


# Recall: Process memory layout

- In theory: Packed

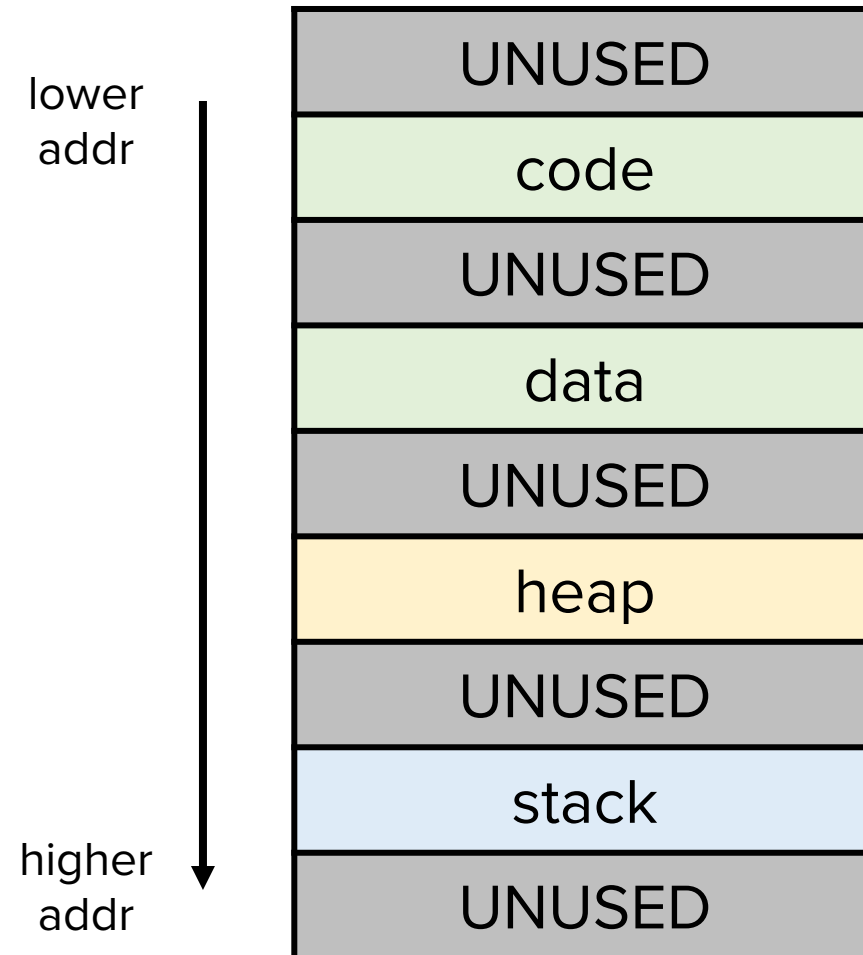


# Recall: Process memory layout

- In theory: Packed



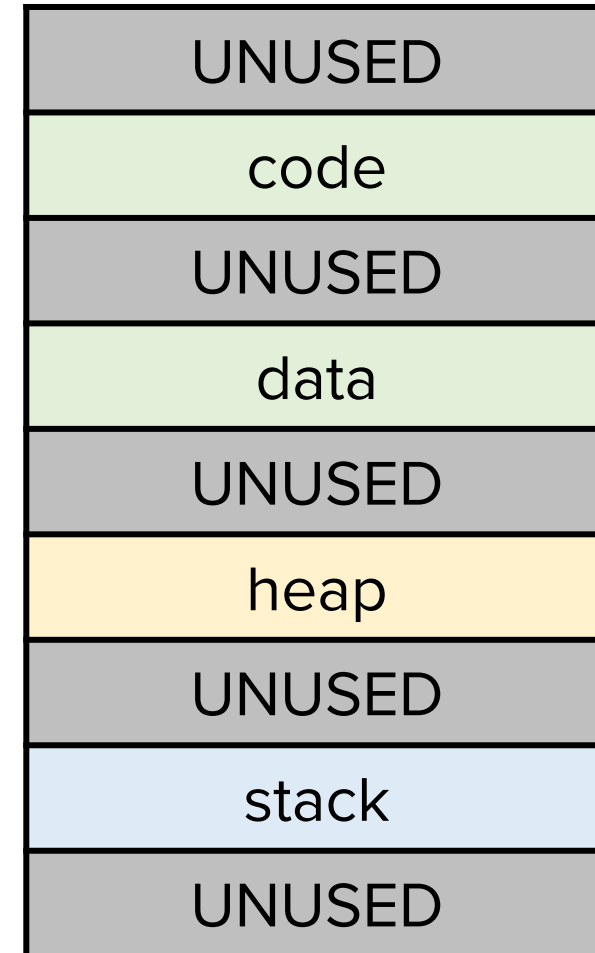
- In practice: Mostly empty



# Recall: Process memory layout

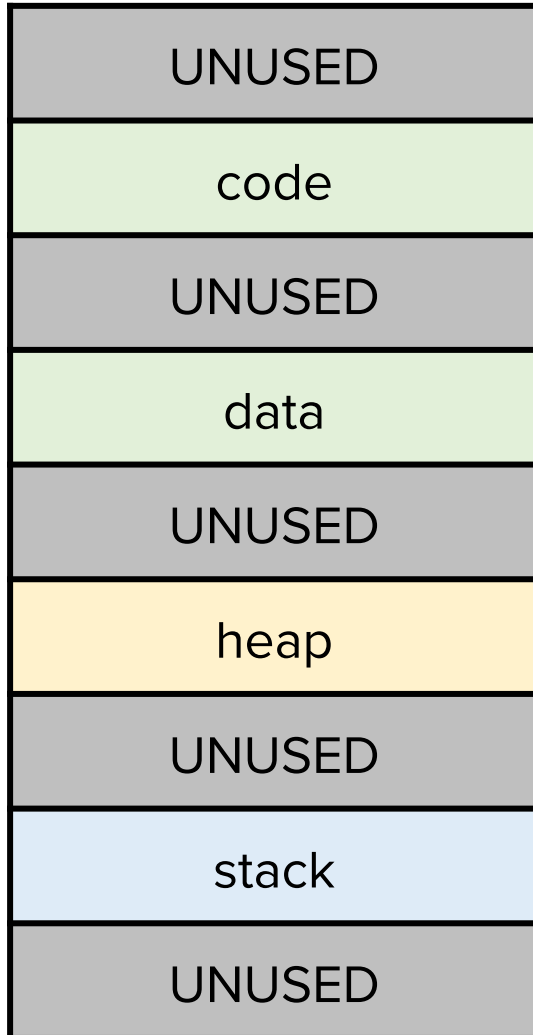
- In practice: Mostly empty

“Wiggle room” exists

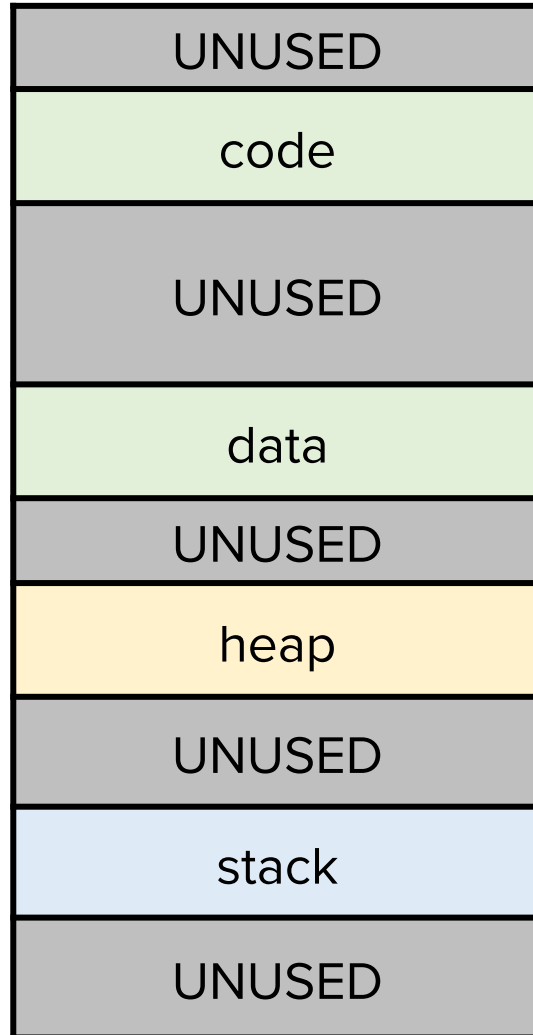


# Idea: Load each segment at different address

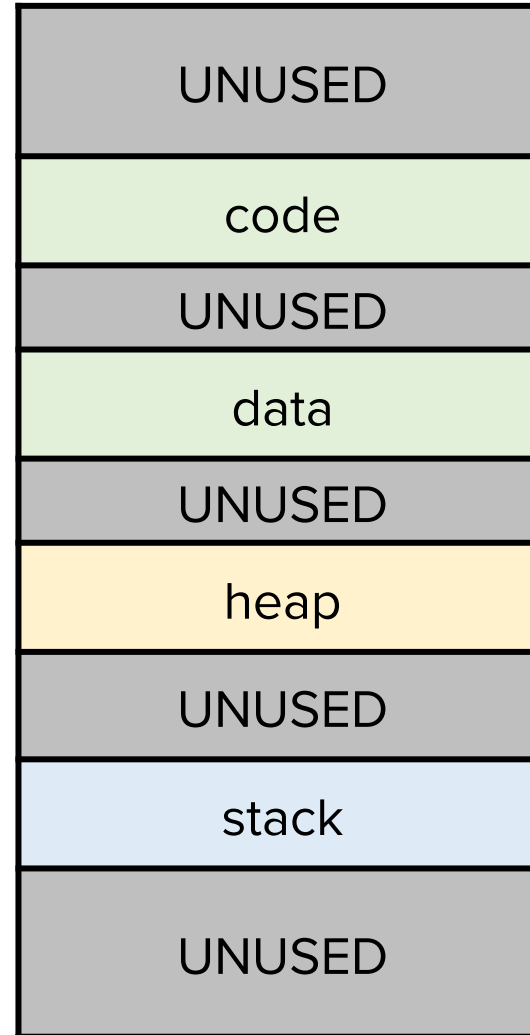
Run #1



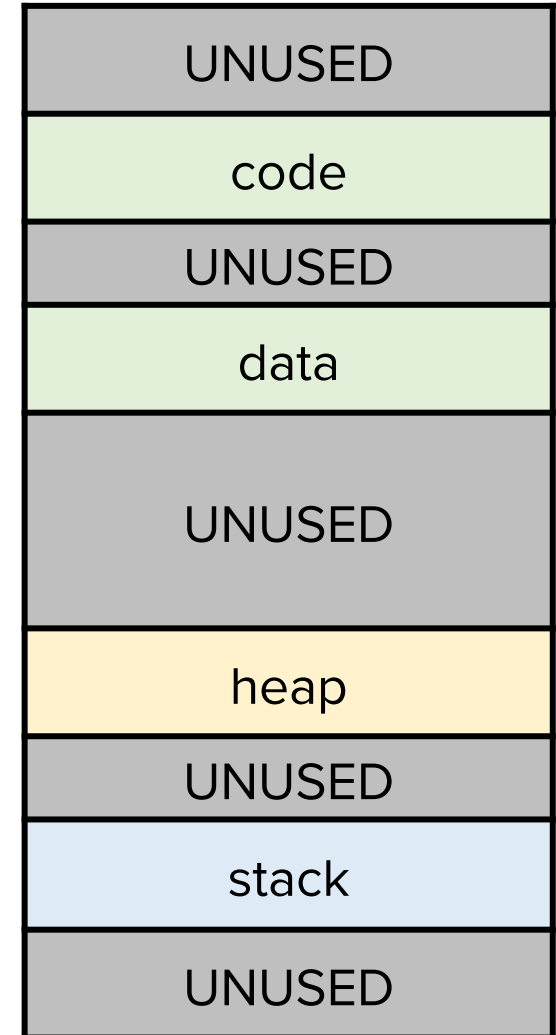
Run #2



Run #3



Run #4



# ASLR can shuffle all four segments

---

- Randomized stack:
  - Can't return to shellcode on stack without knowing the address of the stack buffer or environment variable
- Randomized heap:
  - Can't return to shellcode on the heap without knowing chunk's addr
- Randomized code:
  - Can't ret-to-libc without knowing libc function addresses (e.g., system)
- Randomized data:
  - Can't reuse existing data (e.g., do not know the address of /bin/sh)

# Checking ASLR – Try it yourself

```
#include <stdio.h>

int main(void) {
    int x = 0xdeadbeef;
    return printf("%p\n", &x);
}
```

Code to print stack variable's address

```
$ gcc aslr.c -o aslr
```

Compilation

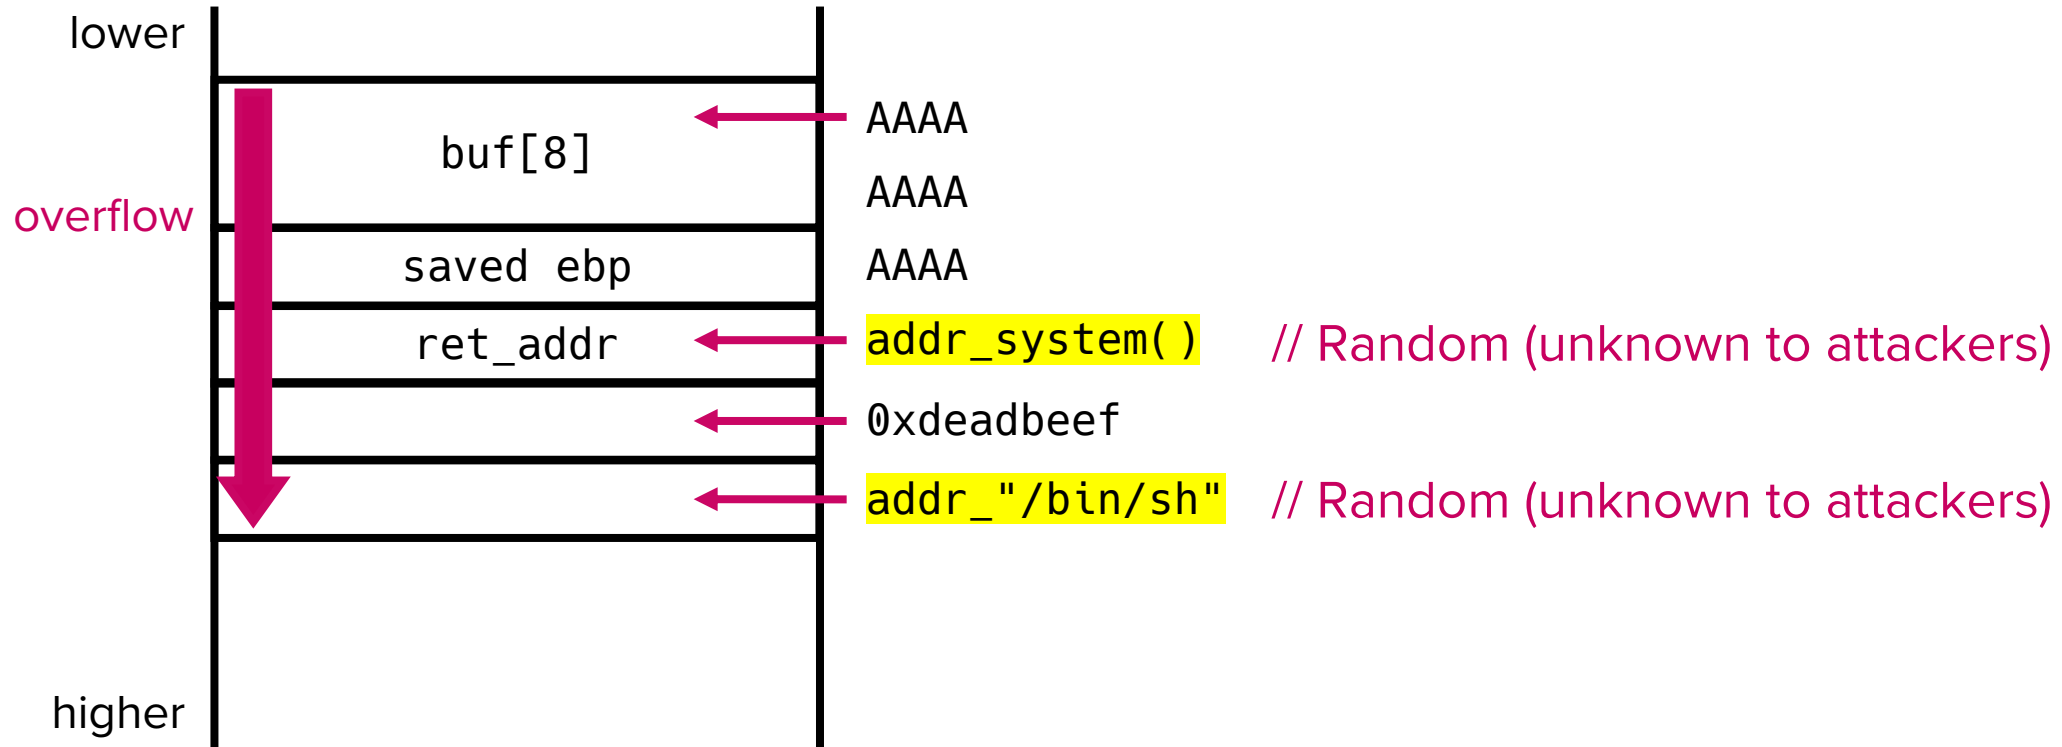
```
$ ./aslr
0x7ffc438c1e84
$ ./aslr
0x7ffd2c1698a4
$ ./aslr
0x7ffdf1b1ca94
```

Stack address is different  
across executions



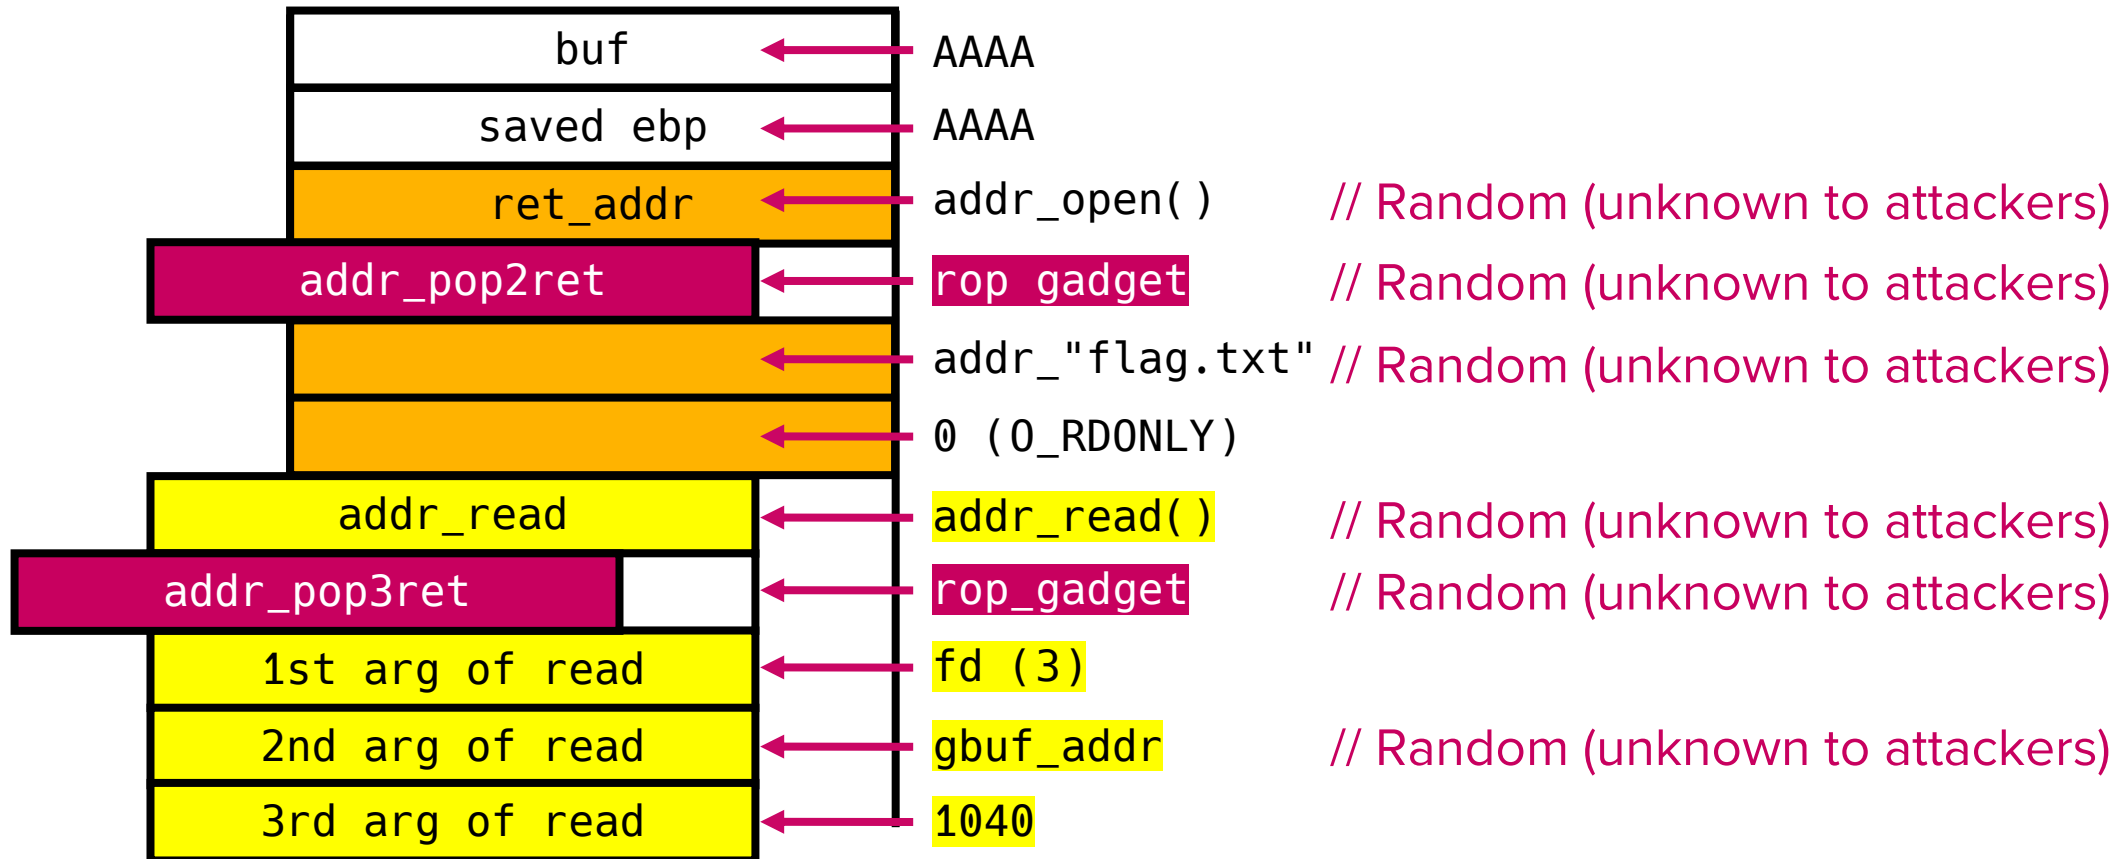
# The effectiveness of ASLR

- ASLR prevents ret-to-libc attacks
  - Cannot return to libc without knowing function and data addresses



# The effectiveness of ASLR

- ASLR prevents ROP attacks
  - Similarly, cannot do ROP without knowing addresses

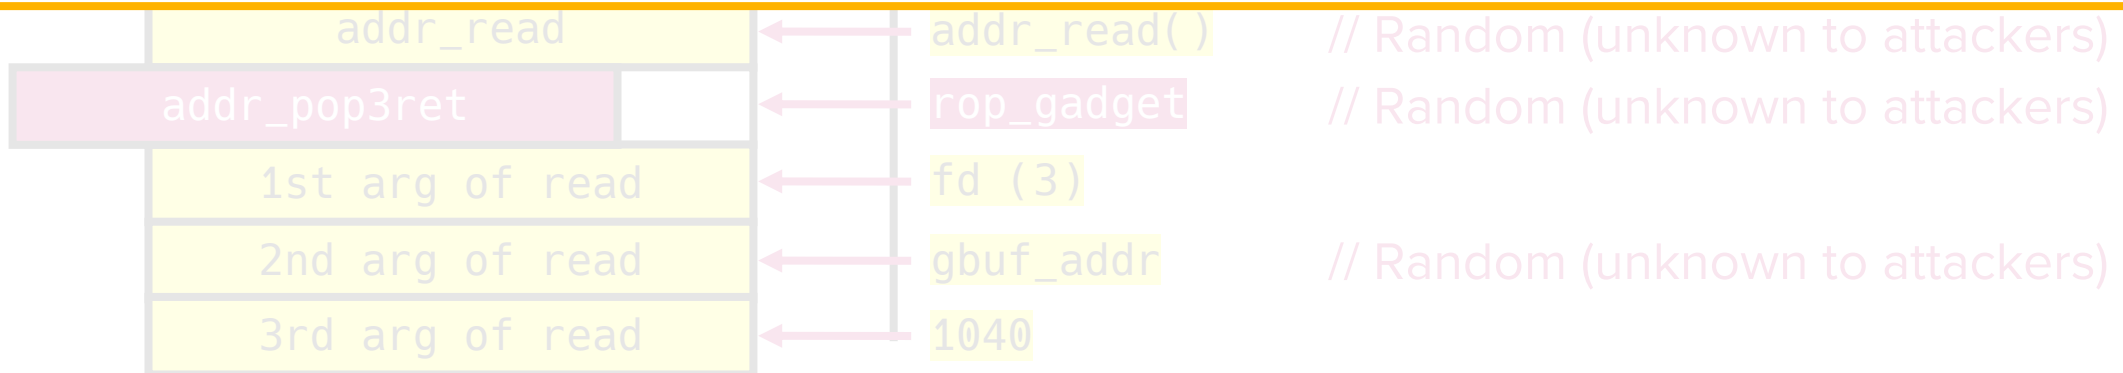


# The effectiveness of ASLR

- ASLR prevents ROP attacks
  - Similarly, cannot do ROP without knowing addresses


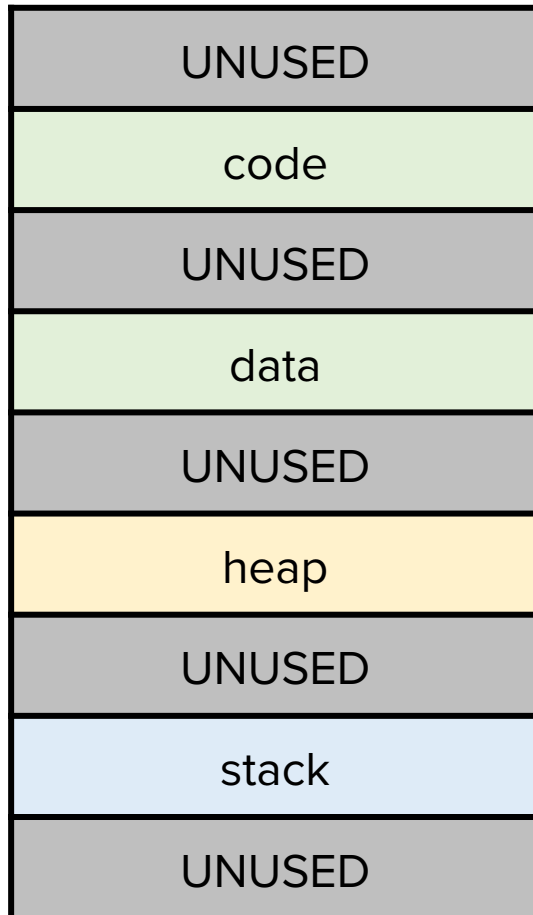


So, are we safe now?



# Subverting ASLR (1)

- ASLR only randomizes the base address of segments

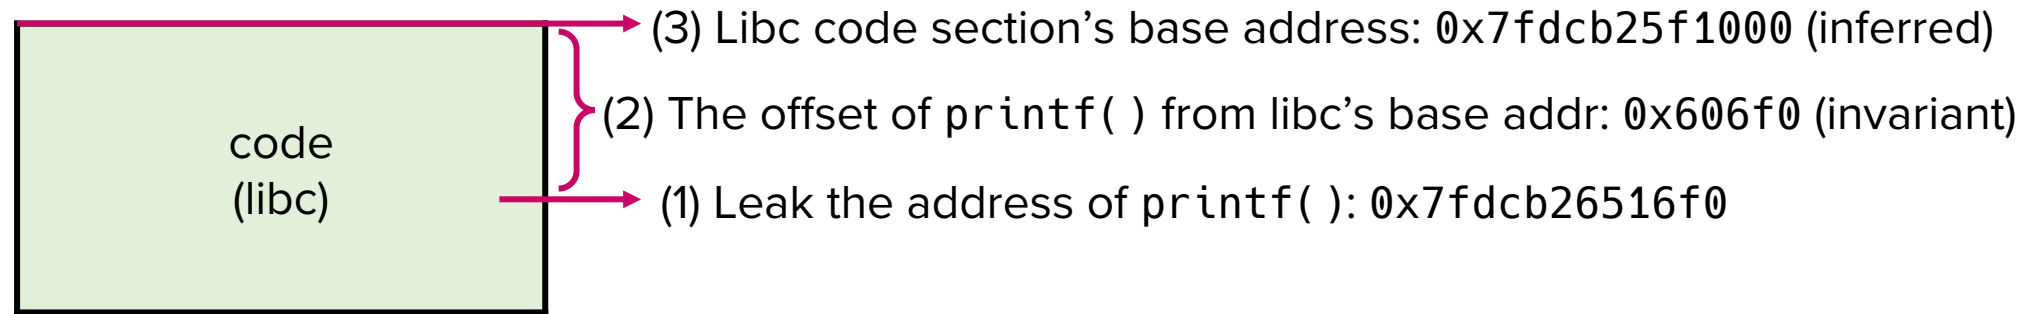


```
<main>:  
base + 11ad: lea    ecx,[esp+0x4]  
base + 11b1: and    esp,0xffffffff0  
base + 11b4: push  DWORD PTR [ecx-0x4]  
base + 11b7: push  ebp  
base + 11b8: mov   ebp,esp  
...
```

Relative addresses (offsets) are fixed!

# Subverting ASLR (1)

- If the address of a pointer within a segment is leaked, the base address can be inferred



- Then, the runtime addresses of other functions within the same segment can be inferred
  - e.g., Offset of `"/bin/sh"` in x86\_64 libc: 0x1d8678  
→  $\text{inferred\_libc\_base\_addr} + 0x1d8678 = \text{addr\_binsh}$

# Subverting ASLR (1)

Q) How can we leak code pointers?  
→ Advanced topic. Take CSED702C!



Q) Why does ASLR only randomize segments' base address?  
→ To ensure reasonable performance!

→  $\text{inferred\_libc\_base\_addr} + 0x1d8678 = \text{addr\_binsh}$

# Subverting ASLR (2)

- Entropy (randomness) of ASLR on x86 Linux is small
  - x86 only randomizes 16 bits (out of 32 bits) of base address
    - There are only  $2^{16} = 65536$  possible addresses for a function
    - Brute-forcible (Pick an address and repeat until the attack works)

# Defense #3: Stack Canary



# Canary



img: Rio Wiki

# Canary in coal mines (Late 1800's-1986)

- Canaries are sensitive to toxic gas
  - e.g., CO (odorless, colorless)
- Coal miners brought canaries into coal mines
  - Canaries die if toxic gases build up
- Miners bail out if a canary dies
  - Sacrificed canaries for miners' lives




img: Times Higher Education

# Canaries as stack protectors

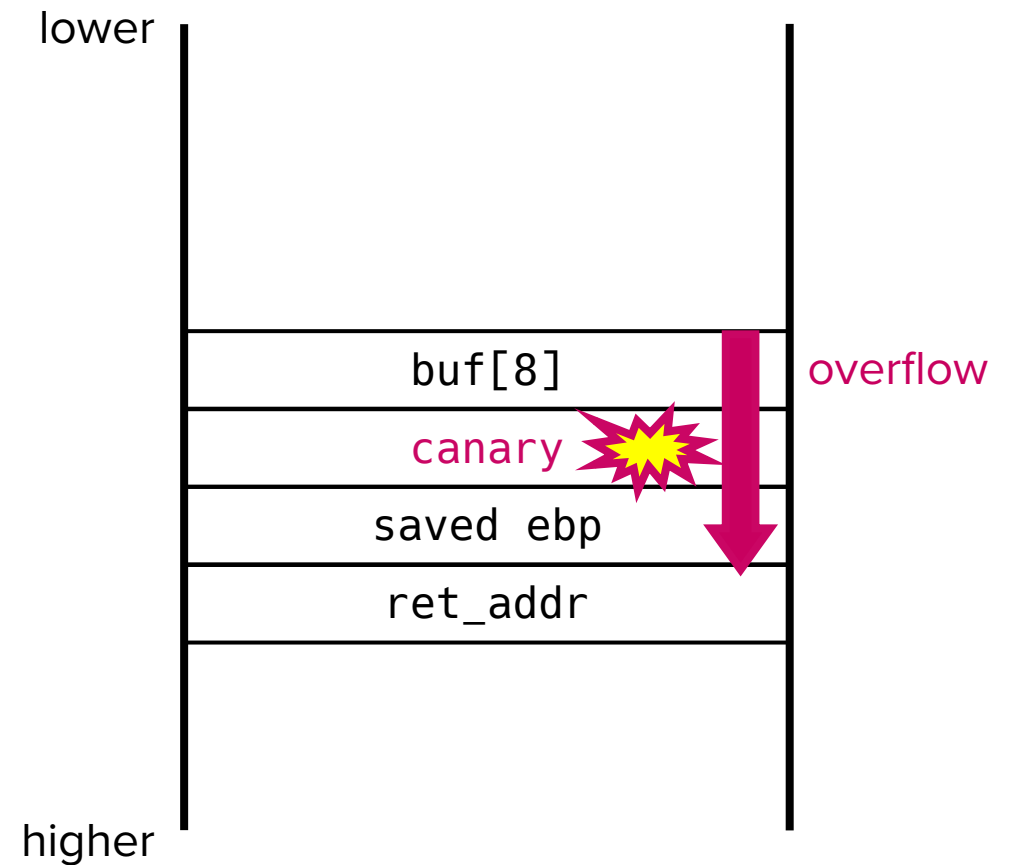
- Workflow of stack smashing attacks:
  - Overwrite the return address of a function
  - Wait for the function to complete and execute **RET** instruction
- Idea for protection:
  - Insert a **canary value** to the stack between the local variables and return address
    - To overwrite the return address, the canary value should be modified
  - Check if the canary is “still alive” before the function returns
    - i.e., check if the value has been modified
  - Terminate execution if the canary is dead

# Stack diagram (x86\_64)

- Without a canary



- With a canary

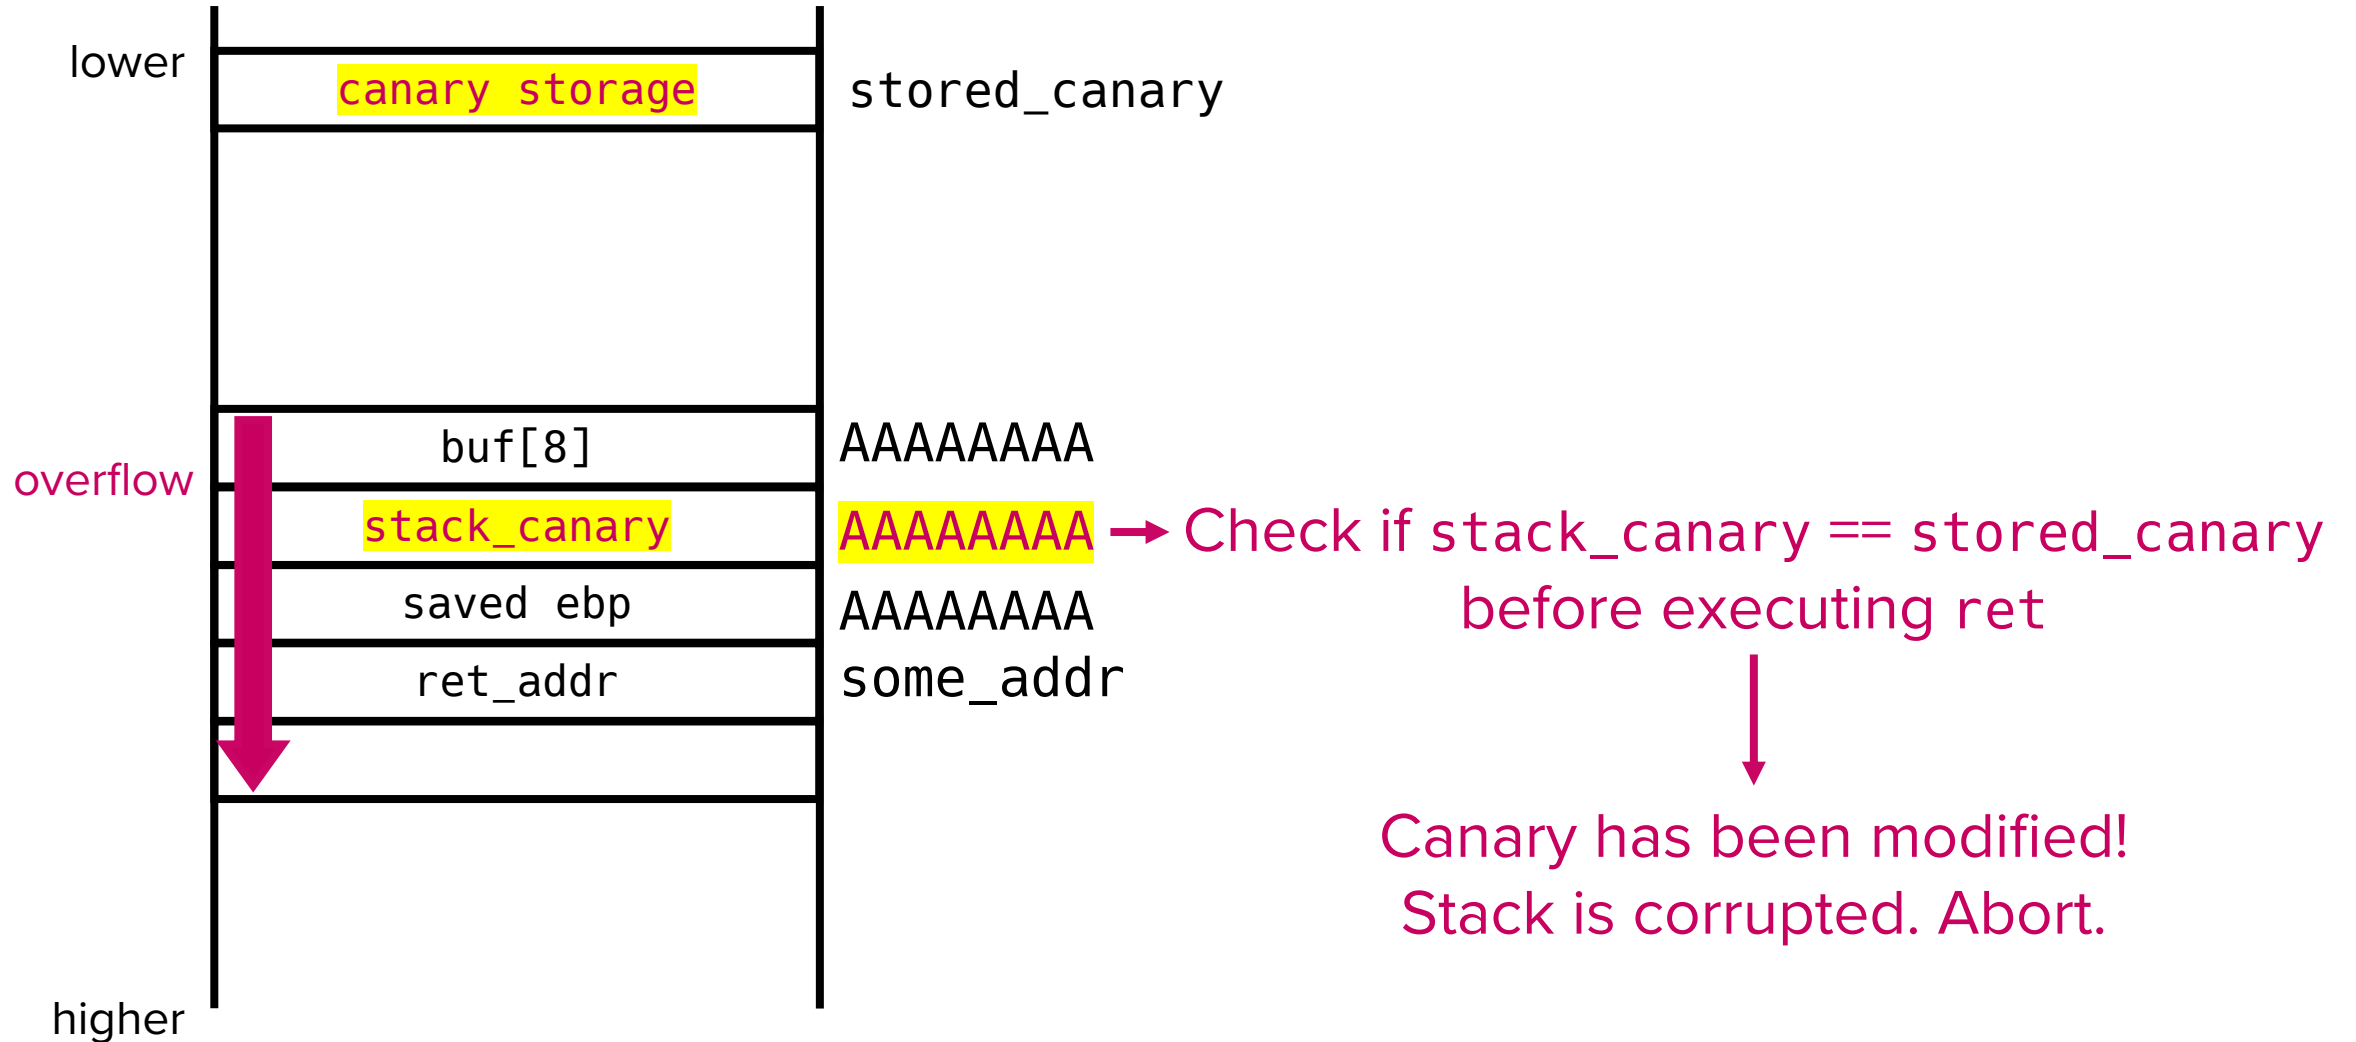


# Canary mechanism

- Binary is executed
- Generate a random secret value and store it in the canary storage
- In the function prologue, place the canary right before the saved frame pointer and return address
- In the function epilogue, check the value on the stack and compare it against the value in canary storage

If the stack canary changes, it is (most likely) due to an attack

# Canary mechanism



# Stack canary in practice

- Implementing canary

```
/* test.c */
#include <string.h>

int main(int argc, char* argv[]) {
    char buf[8];
    memcpy(buf, argv[1], 64);
    return 0;
}
```

PATCH

```
/* test-canary.c */
#include <string.h>
long CANARY = gen_random_canary();

int main(int argc, char* argv[]) {
    long canary = CANARY;
    char buf[8];
    memcpy(buf, argv[1], 64);
    if (canary != CANARY)
        exit(-1);
    return 0;
}
```

This is not what we do in practice.

It is not desirable to manually do this for every function of a software

# Stack canary in practice

- Compilers, during compilation, automatically generate and insert code for loading and checking canary

```
/* test.c */
#include <string.h>

int main(int argc, char* argv[]) {
    char buf[8];
    memcpy(buf, argv[1], 64);
    return 0;
}
```

→ No need to manually patch the code

Compilation (canary is enabled by default)

```
$ gcc test.c -o with-canary
```

Compilation (disable canary)

```
$ gcc test.c -fno-stack-protector -o without-canary
```



# Stack canary in practice

```
; disasm of without-canary
endbr64
push rbp
mov rbp, rsp
sub rsp, 0x20
mov DWORD PTR [rbp-0x14], edi
mov QWORD PTR [rbp-0x20], rsi
```

```
mov rax, QWORD PTR [rbp-0x20]
add rax, 0x8
mov rcx, QWORD PTR [rax]
lea rax, [rbp-0x8]
mov edx, 0x40
mov rsi, rcx
mov rdi, rax
call 0x1050 <memcpy@plt>
mov eax, 0x0
```

```
leave
ret
```

```
; disasm of with-canary
endbr64
push rbp
mov rbp, rsp
sub rsp, 0x20
mov DWORD PTR [rbp-0x14], edi
mov QWORD PTR [rbp-0x20], rsi
mov rax, QWORD PTR fs:0x28
mov QWORD PTR [rbp-0x8], rax
xor eax, eax
mov rax, QWORD PTR [rbp-0x20]
add rax, 0x8
mov rcx, QWORD PTR [rax]
lea rax, [rbp-0x10]
mov edx, 0x40
mov rsi, rcx
mov rdi, rax
call 0x55555555070 <memcpy@plt>
mov eax, 0x0
mov rdx, QWORD PTR [rbp-0x8]
sub rdx, QWORD PTR fs:0x28
je 0x555555551c3 <main+90>
call 0x55555555060 <__stack_chk_fail@plt>
leave
ret
```

# Stack canary in practice

(canary storage)  
Fetches a canary value from fs:0x28 →  
and stores it at [rbp-8]  
(between local vars and return addr)

Compares the stored canary at [rbp-8]  
with the value in the canary storage at fs:0x28 →

Calls \_\_stack\_chk\_fail( ) if they differ →

```
; disasm of with-canary
endbr64
push rbp
mov rbp, rsp
sub rsp, 0x20
mov DWORD PTR [rbp-0x14], edi
mov QWORD PTR [rbp-0x20], rsi
mov rax, QWORD PTR fs:0x28
mov QWORD PTR [rbp-0x8], rax
xor eax, eax
mov rax, QWORD PTR [rbp-0x20]
add rax, 0x8
mov rcx, QWORD PTR [rax]
lea rax, [rbp-0x10]
mov edx, 0x40
mov rsi, rcx
mov rdi, rax
call 0x55555555070 <memcpy@plt>
mov eax, 0x0
mov rdx, QWORD PTR [rbp-0x8]
sub rdx, QWORD PTR fs:0x28
je 0x555555551c3 <main+90>
call 0x55555555060 <__stack_chk_fail@plt>
leave
ret
```

# Stack canary in practice

- Result

```
$ ./without-canary aaaa  
[1] 2304931 segmentation fault ./without-canary aaaa
```

→ We can redirect the control flow by overwriting the return address with a valid address)

```
$ ./without-canary aaaa  
*** stack smashing detected ***: terminated  
[1] 2304723 IOT instruction ./without-canary aaaa
```

→ `__stack_chk_fail()` displays the error message and terminates

# Subverting stack canary (1)

---

- Leak the canary's value
  - Exploit any vulnerability that can reveal stack memory
    - Examples:
      - The HeartBleed vulnerability allows attackers to read arbitrary memory
      - Format string bugs allows printing stack values at chosen addresses
  - Once the canary is leaked, overwrite the canary with the leaked value

# Subverting stack canary (2)

- Guess the canary's value
  - The least significant byte (LSB) of canary is always **0x00** // Q) Why?
    - On x86 (32-bit) systems, this leaves 24 bits (out of 32 bits) to guess
    - There are  $2^{24}$  (~16 million) possibilities, which is feasible for brute force
  - Once successfully guessed, overwrite the canary with the correct value
    - Q) Doesn't canary value change each time a process is run?
    - A) In many server applications, the process is forked per request, so the canary (inherited from the parent process) remains the same across child processes

# Subverting stack canary (3)

- Bypass the canary
  - Stack canary can only mitigate sequential writes (e.g., via `strcpy( )`)
  - Some vulnerabilities (e.g., format string bugs) allow you to do “arbitrary writes”
    - An attacker can write to any memory address directly
  - You can overwrite the return address without altering the stack canary!

# Section 1 of CSED415 is over

- Attacks typically begin with a BOF vulnerability
  - Attack surface analysis (Lecture 02)
  - Buggy code is the root of evil (Lecture 03)
- Mitigations exist, but they are not perfect
  - They are designed in a way that their impact on CIA is minimal
  - There is a trade-off between cost and effectiveness

Primary focus so far: System Integrity (ref: Lecture 02)

*“A system performs its intended function in an unimpaired manner.”*

Question: How can we preserve Confidentiality?

# Coming up next

- Cryptographic primitives and their applications

```
/U$P"$&%js9Kn+0q[D%KvVyum%xDG{Re+LeeItEhBXG&|(: (VxDM3YV`2)`geadp&t  
: )VV7oJ)gk}h7>09)b(t+a!ES@3m4e|.j1Z.7J>IiZ+M=|F%Bo^c{RLSjYq"\]aTp@  
o%,<rTf=ExIY?)FRPFhpfPgT'tUNtRW5A}kFR?3qVX=R[f/B8aVGhu"q;8TLh+DlW?  
Q{Of4y;!v3McuZN~gV9Dyi/4yG80)sBV[nWKWm}ZN3Bza8bNkqNU6Cv+9pISyW'x4Y  
ZvPEW8elyAraa7A=0qV&gi+EpmS0p!1QImoPCy78M}bWeFV(tS@q8NbgOE8;]0EP+i  
pL*x8AF*AyMt&q(V9QsWi9Gwt4akvhPb817bjk,q0i3RLkXc*Sq%z46vCthON6hQu(  
28613jY67i*h4d&iMIPEMK*5Uy/8W2)K5Dt8Qr^ZwAQQ6JAp)9akcJ8pRr%?mp9[9&  
JP6]rZ7GxPLJ74W8gv9Y*2RYC)P8l?m)XWrpSeMB);mNJe/07Rxw7P53PU9aHWSuc2  
*E=zVR@`Wn*4dWSxr9IZp4.HQy$AcIZfAWUMQD@h|8PwJYcn2V%&AR$KnHZK\yS=wQ  
day'I/lW[YM9H6aQWhV)D/2eNi6G1k8F>4Ka9&E4kFg118R8JXH[0#8QVjVMEC,mRx  
}gNqp[MD#$S0d0);u~a^]j@8?RCX('$tk8Ypj:"wnwVUP&#8YKc}D0HYvBvjrwWc,  
Na@Ma)T3Nw0024FbRJtKxgs6l3BcTXn4)oRlsc]*R=ol840mW%NJC6&mom!vhJJ8j?  
@/zw5`Cd,d90&Y70\EP!UQya$Q`VWGr0eQ00YgT.rKH#5KB0C2!&HW3Urm%c&Gboni  
)<<l\dQ:aB^[lhq*&7LQeP"qvCz0k'C|x%ekHe^!xI>hLVuPjmlrajeEne"iSAGoHx  
{@b!HL=\y0!VTI"4q\UfMoxl;shQ8hzc?%s!wGT(3gL`b"W(=AlUA;W[=vKzZbg;qT  
L:'Lby+Wn{%AJ@Lux=nAk7V~Irfw%jIV%kOc&,go}+vDv5F7h!J~4FGWy.awxWLd"M  
IhvlTv^1uGq/nmv&*?!F8Q"EhI0`yt5:IwYZIF,XR:`Q+hnlXwEVnwzRdC#eq%E*9a  
VxDM3YV`2)`ge1dp&t0WquwXb:`P~u%gD{sAr^<lu080Z:)VV7oJ)gk}h7;/Ig;yfA
```



# Questions?